Binary segmentation

Toby Dylan Hocking

Introduction

Analysis of time complexity

Comparing empirical timings between implementations

Motivation for changepoint detection in time series data

- Detecting changes/abnormalities important in medicine.

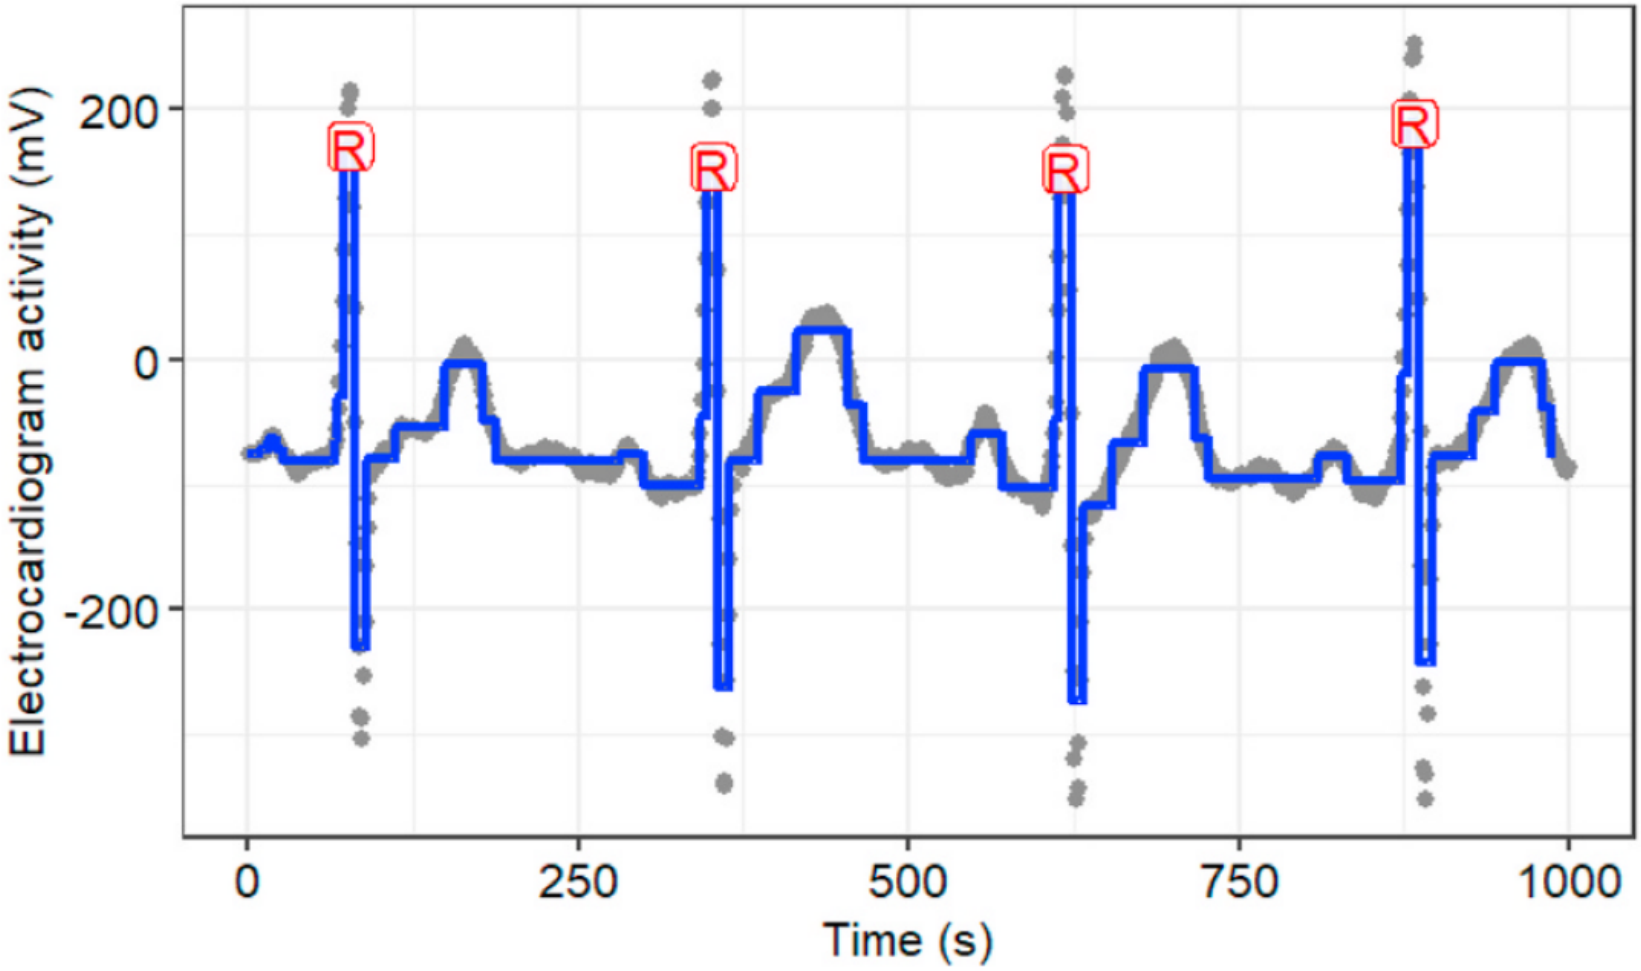

Figure 1: Electrocardiograms (heart monitoring), Fotoohinasab et al, Asilomar conference 2020.

Motivation for changepoint detection in time series data

- Detecting the time when a spike occurs is important in neuroscience.

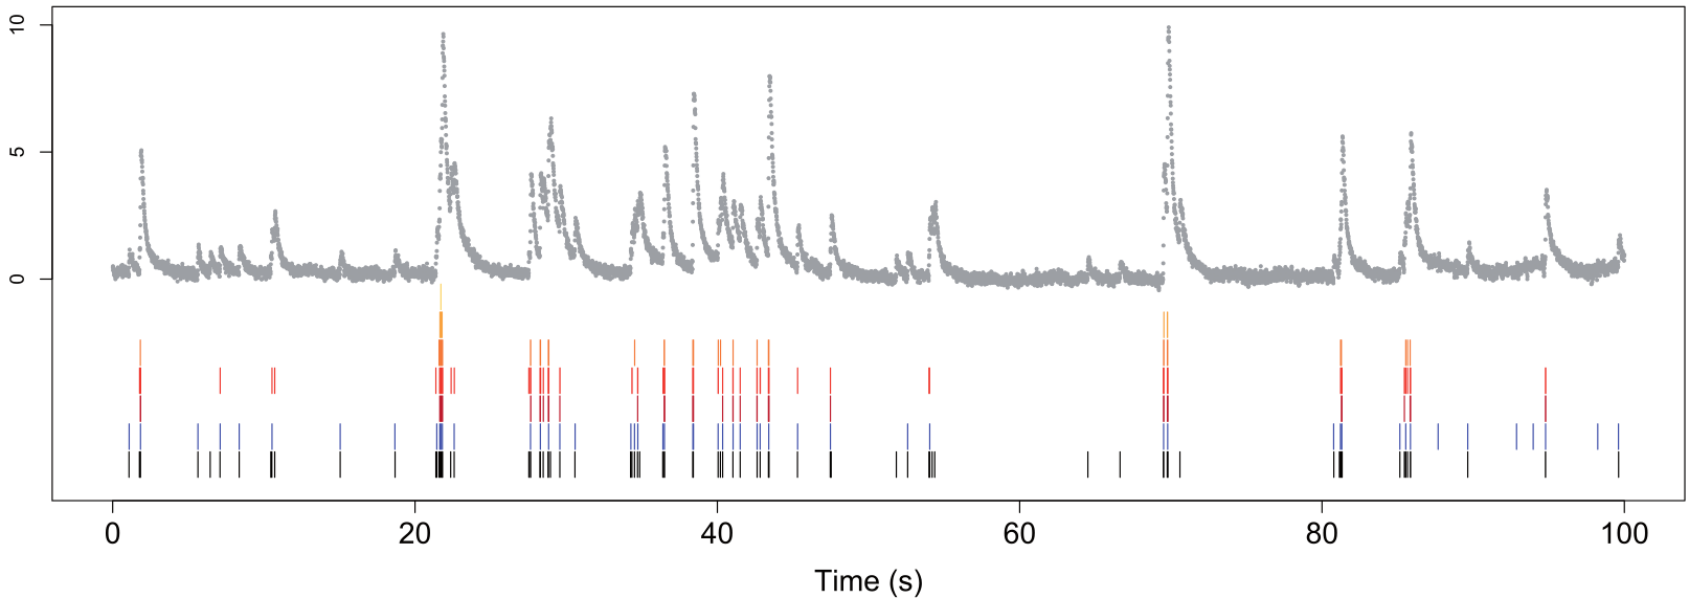

Figure 2: Neural spikes in calcium imaging data, Jewell et al, Biostatistics 2019.

Motivation for changepoint detection in genomic data sequences

- Detecting breakpoints is important in diagnosis of some types of cancer, such as neuroblastoma.

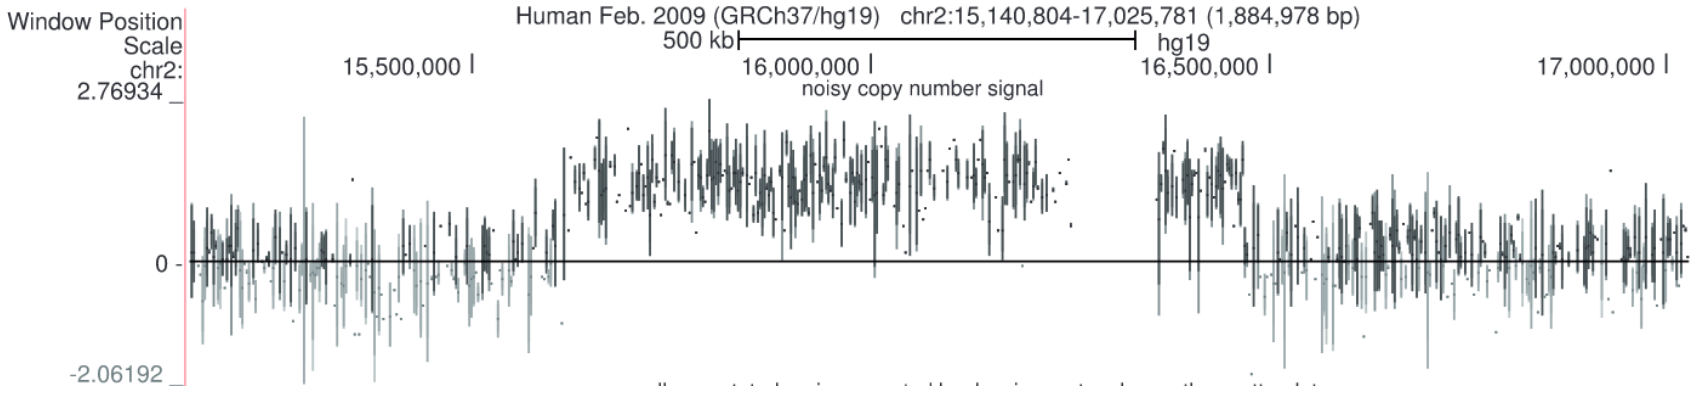

Figure 3: DNA copy number data, breakpoints associated with aggressive cancer, Hocking et al, Bioinformatics 2014.

Motivation for changepoint detection in genomic data sequences

- Detecting peaks (up/down changes) in genomic data is important in order to understand which genes are active or inactive.

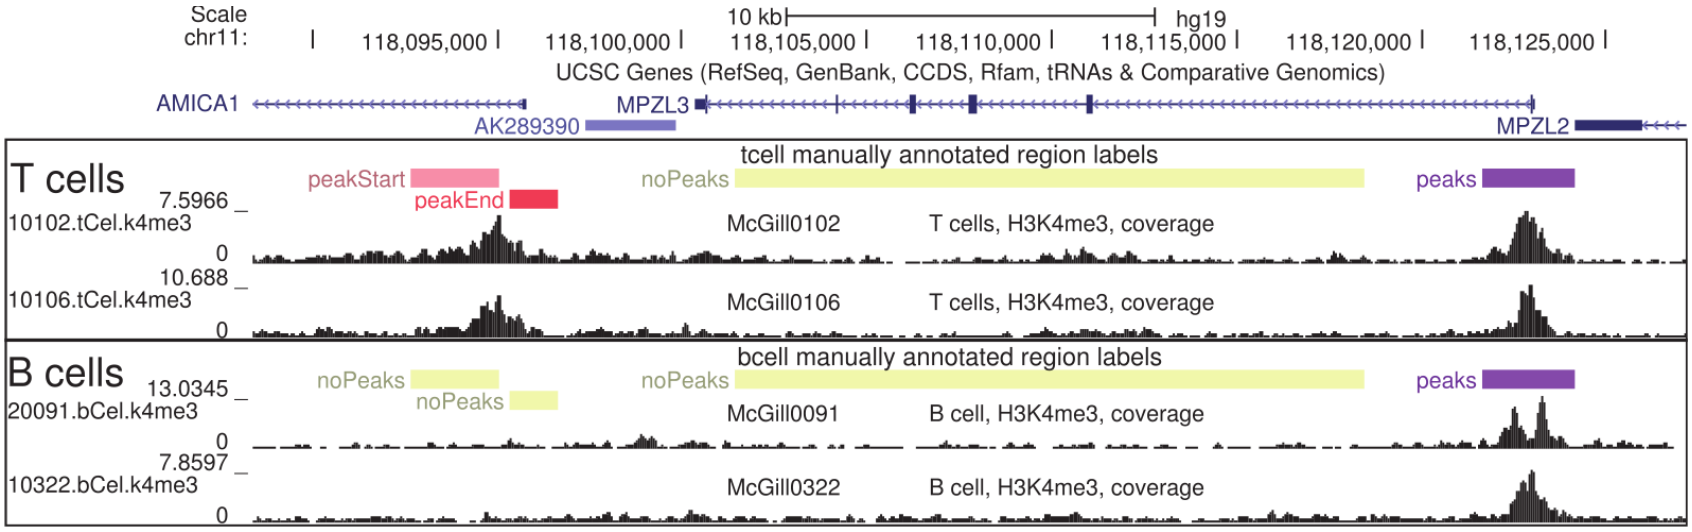

Figure 4: ChIP-seq data for characterizing active regions in the human genome, Hocking et al, Bioinformatics 2017.

Segmentation / changepoint detection framework

- Let $x_1, \ldots, x_n \in \mathbb{R}$ be a data sequence over space or time (logratio column in DNA copy number data below).
- Where are the abrupt changes in the data sequence?

```
##       profile.id chromosome  position    logratio
##           <fctr>     <fctr>     <int>       <num>
##    1:          4          2   1472476  0.44042072
##    2:          4          2   2063049  0.45943162
##    3:          4          2   3098882  0.34141652
##    4:          4          2   7177474  0.33571191
##    5:          4          2   8179390  0.31730407
##   ---
## 230:          4          2 239227603  0.01863417
## 231:          4          2 239471307  0.01720929
## 232:          4          2 240618997 -0.10935876
## 233:          4          2 242024751 -0.13764780
## 234:          4          2 242801018  0.17248752
```

Segmentation / changepoint data visualization

- Let $x_1, \ldots, x_n \in \mathbb{R}$ be a data sequence over space or time.
- Where are the abrupt changes in the data sequence?

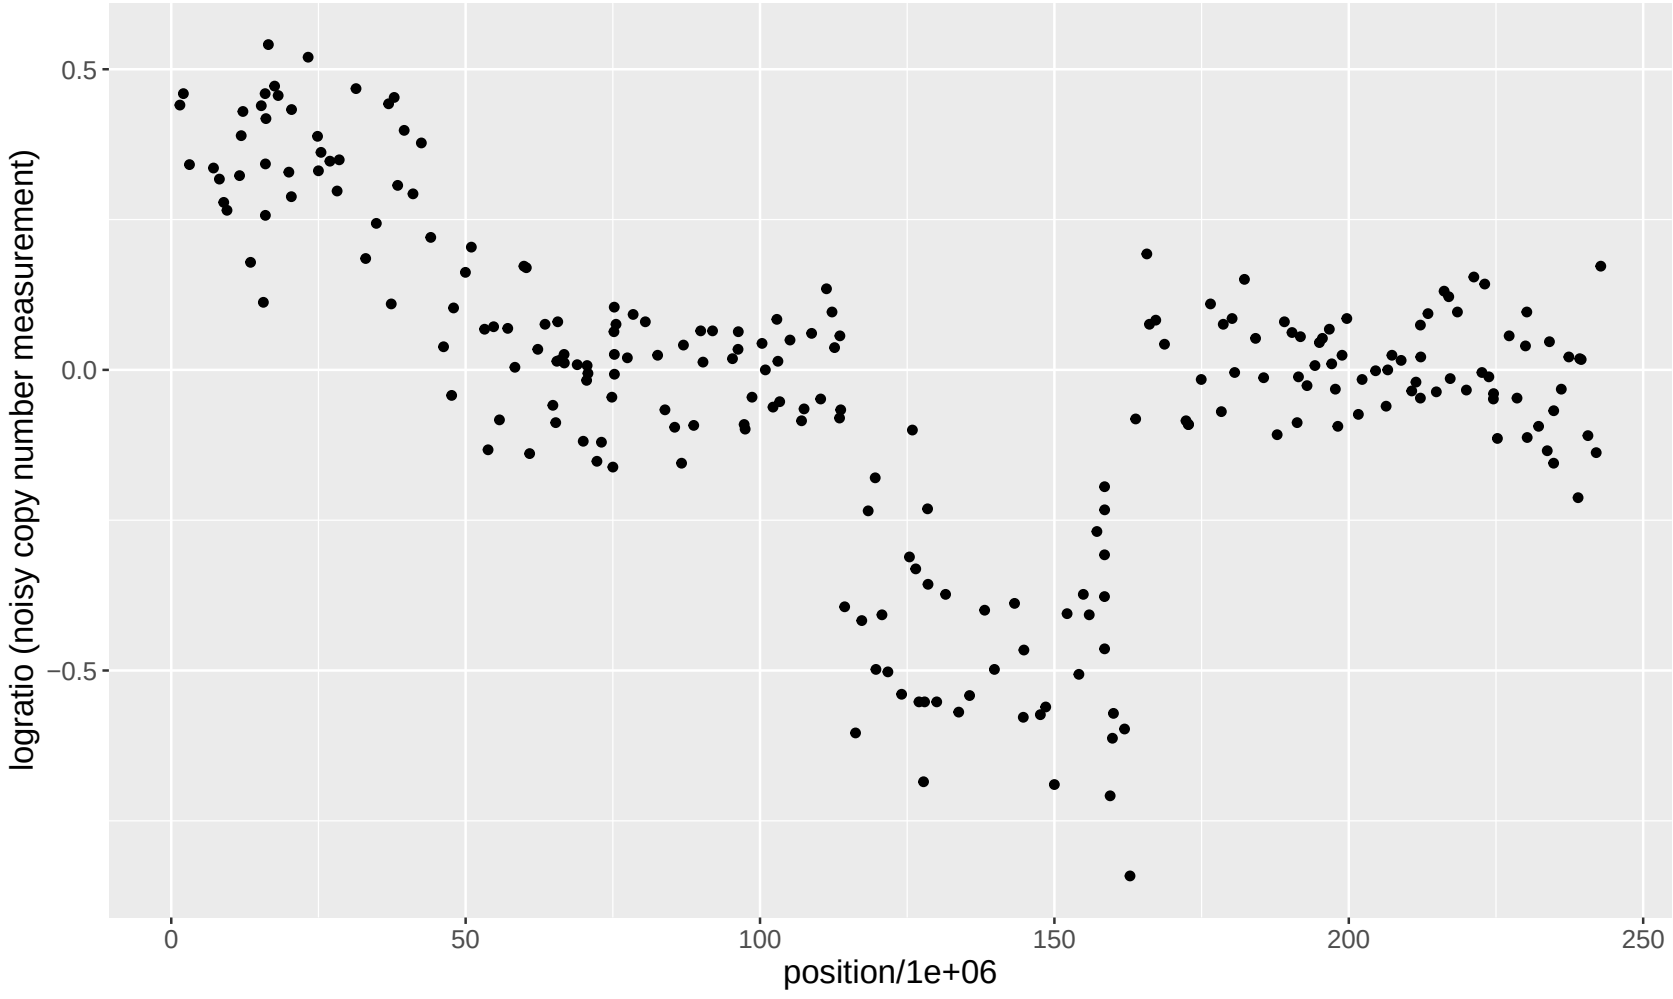

Assume normal distribution with change in mean, constant variance

- There are a certain number of clusters/segments $K \in \{1, \ldots, n\}$.
- Each segment $k \in \{1, \ldots, K\}$ has its own mean parameter $\mu_k \in \mathbb{R}$.
- There is some constant variance parameter $\sigma^2 > 0$ which is common to all segments.
- For each data point i on segment $k \in \{1, \ldots, K\}$ we have $x_i \sim N(\mu_k, \sigma^2)$ – normal distribution.
- This normal distribution assumption means that we want to find segments/changepoints with mean m that minimize the square loss, $(x - m)^2$.
- Other distributional assumptions / loss functions are possible.

Visualize data sequence

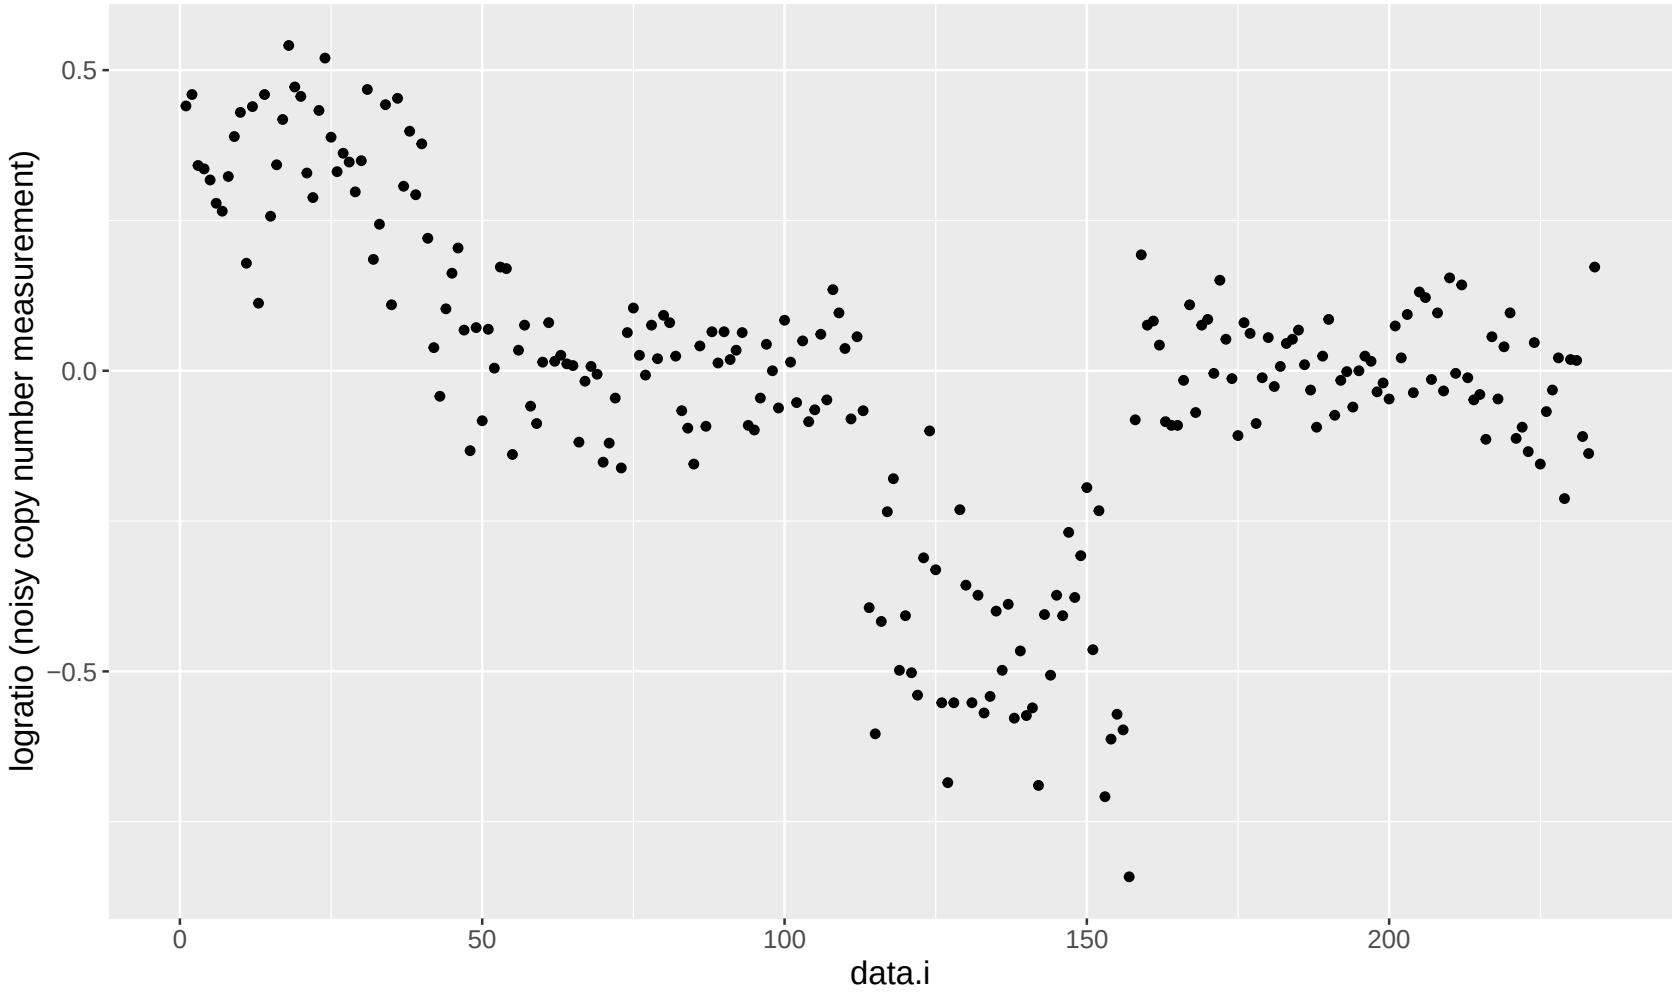

Simplest model, 1 segment, 0 changepoints

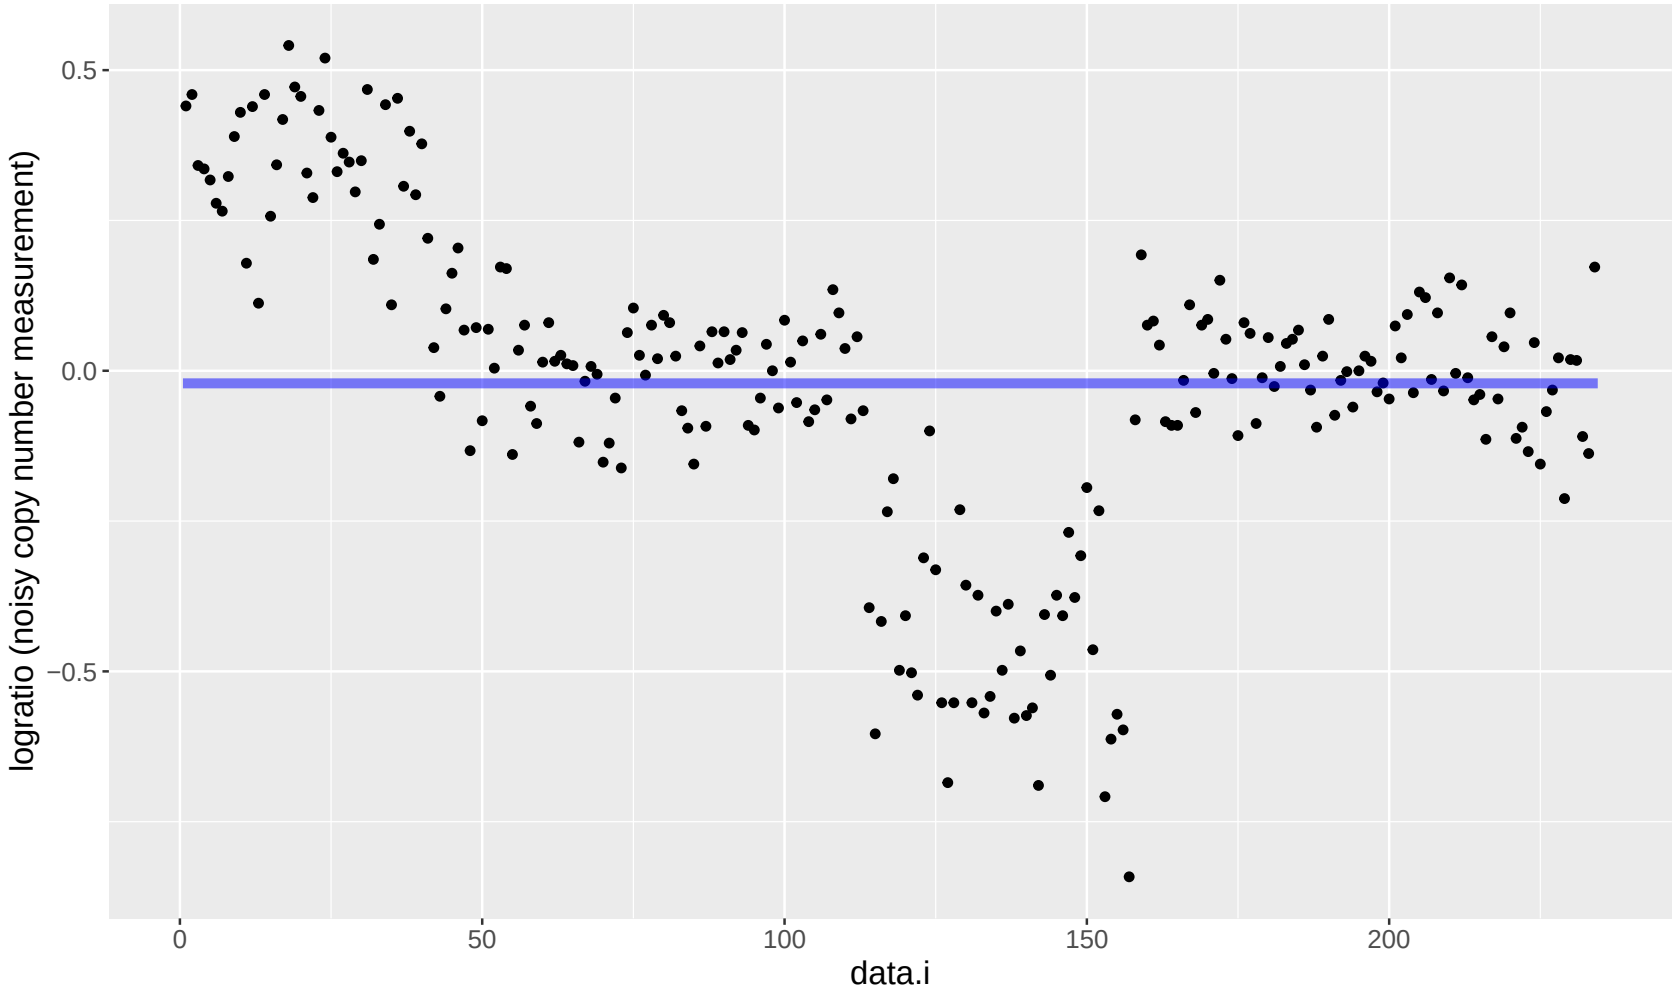

Find best single changepoint (two segments)

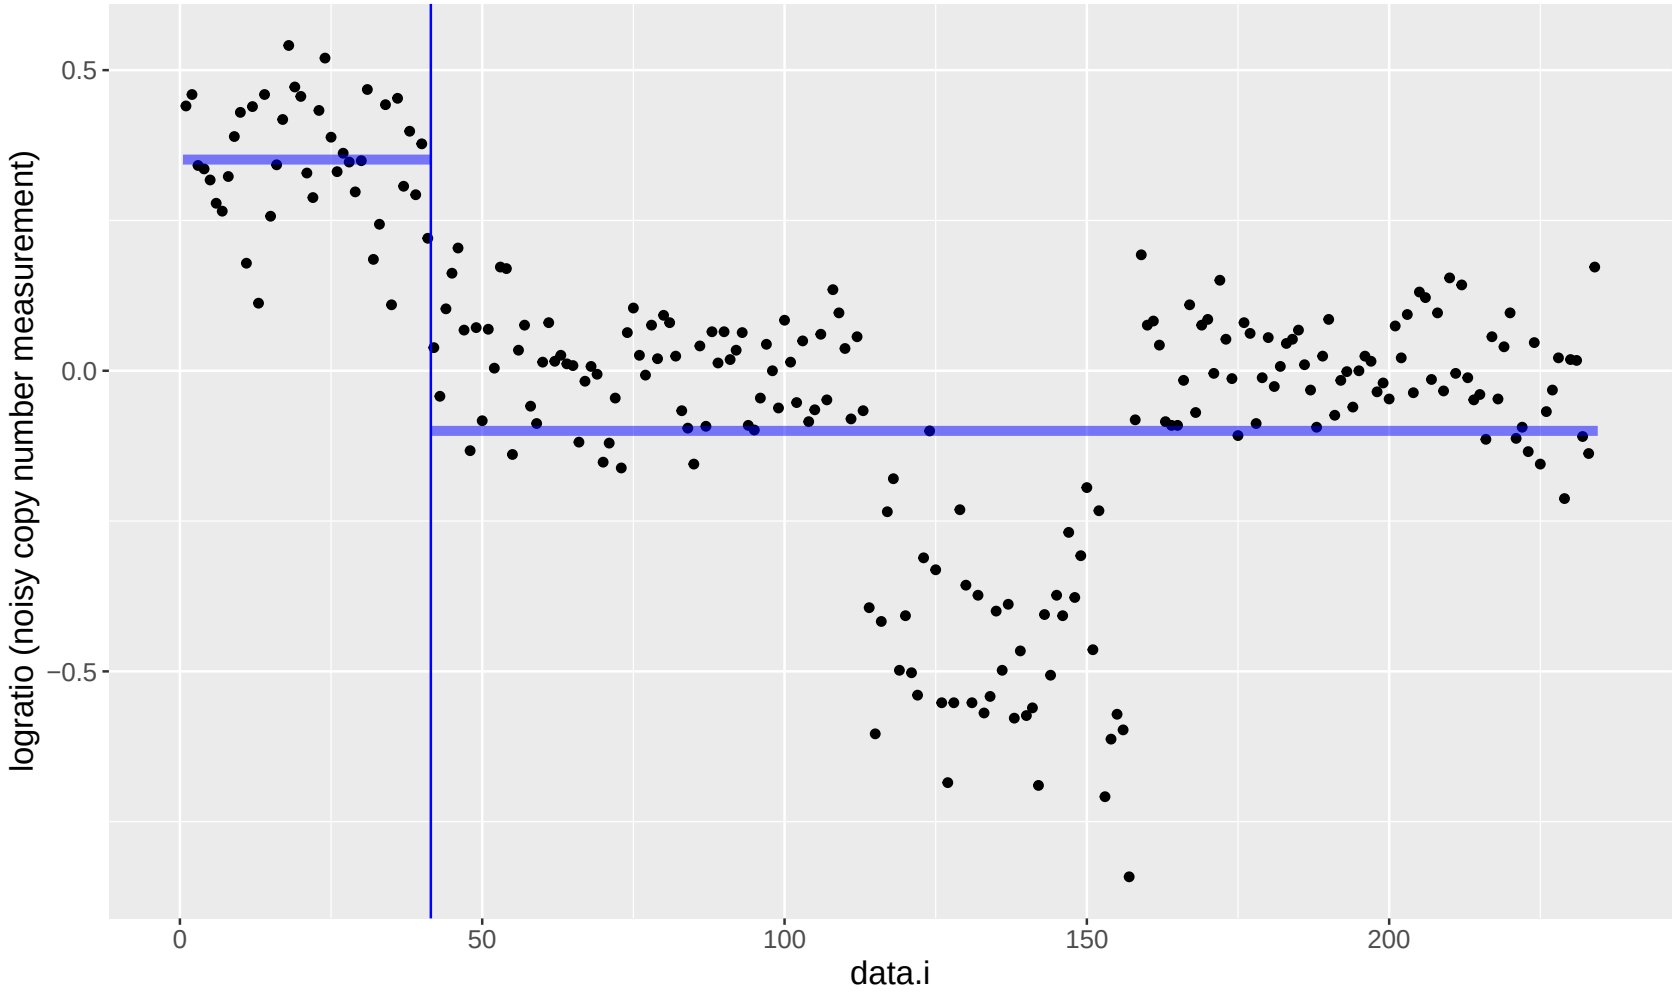

Find two changepoints (three segments)

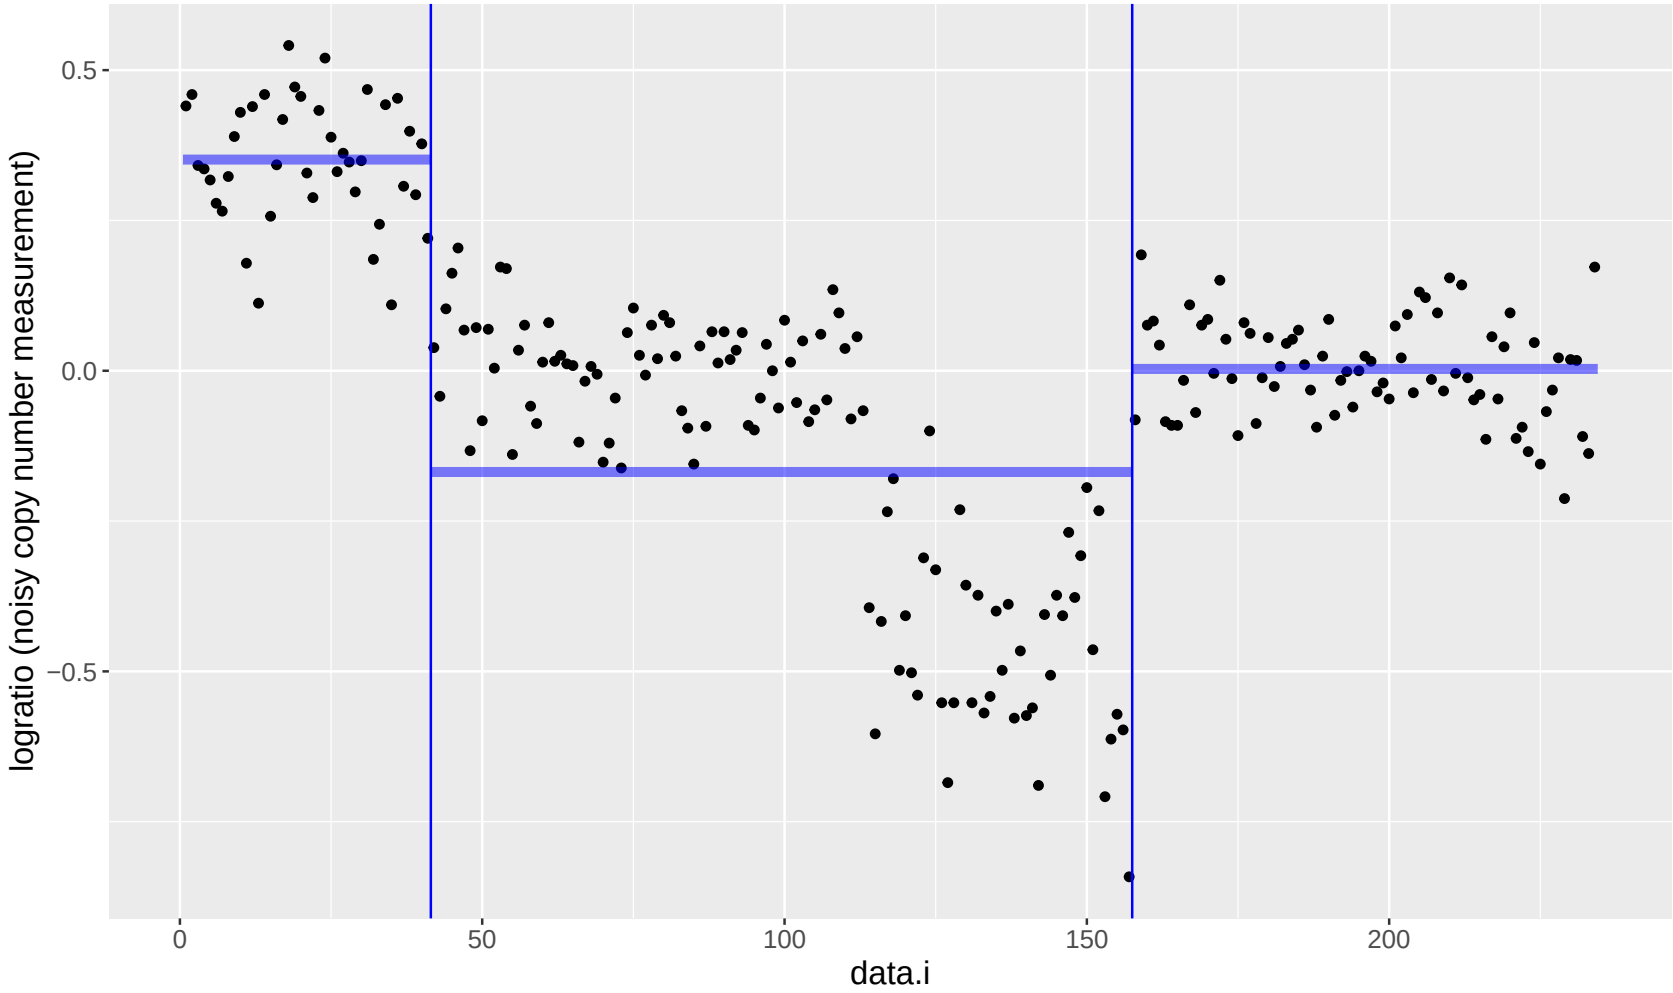

Find four segments

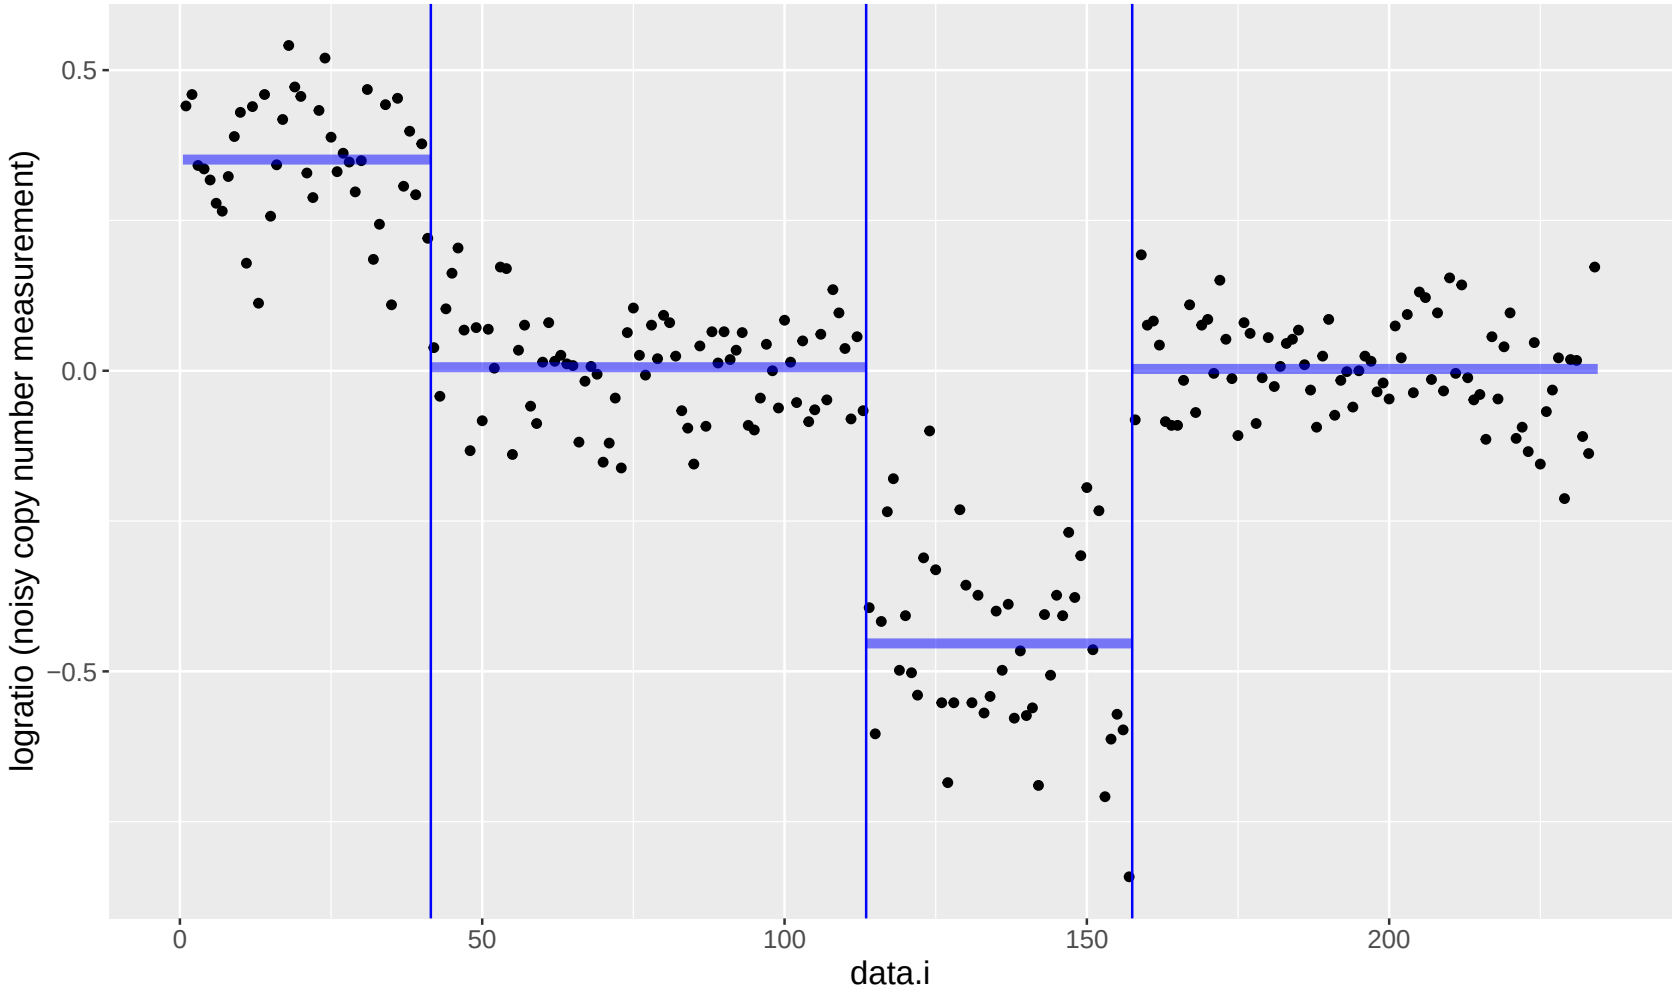

Find five segments

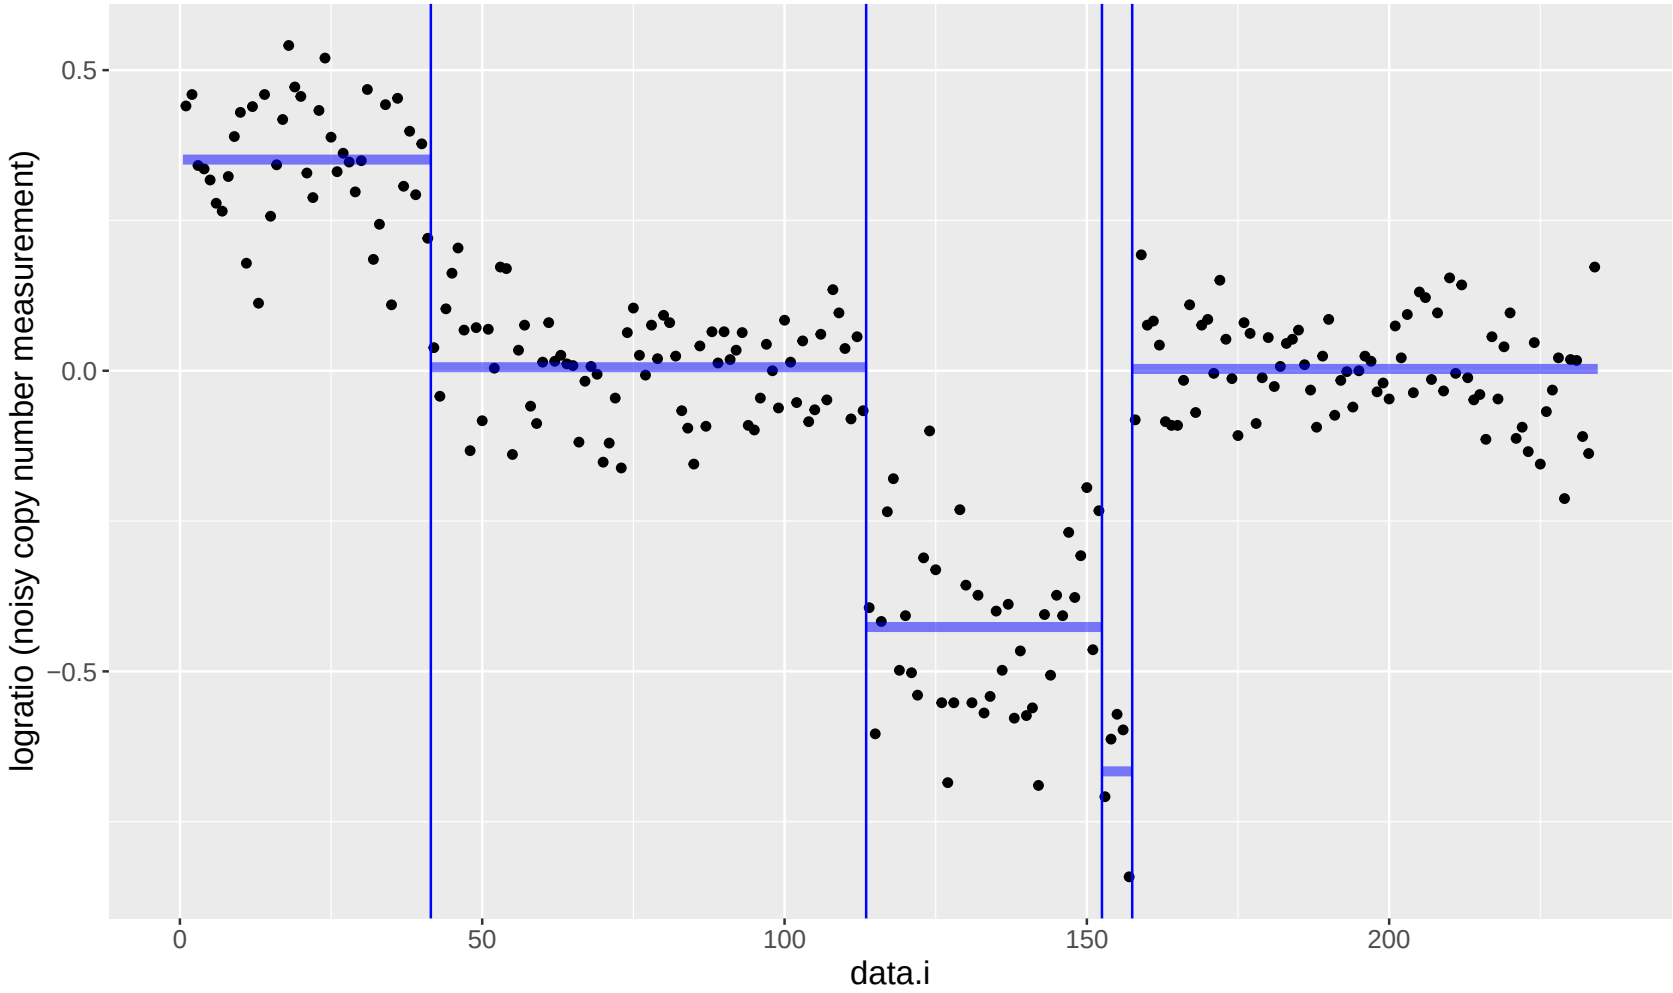

Find six segments

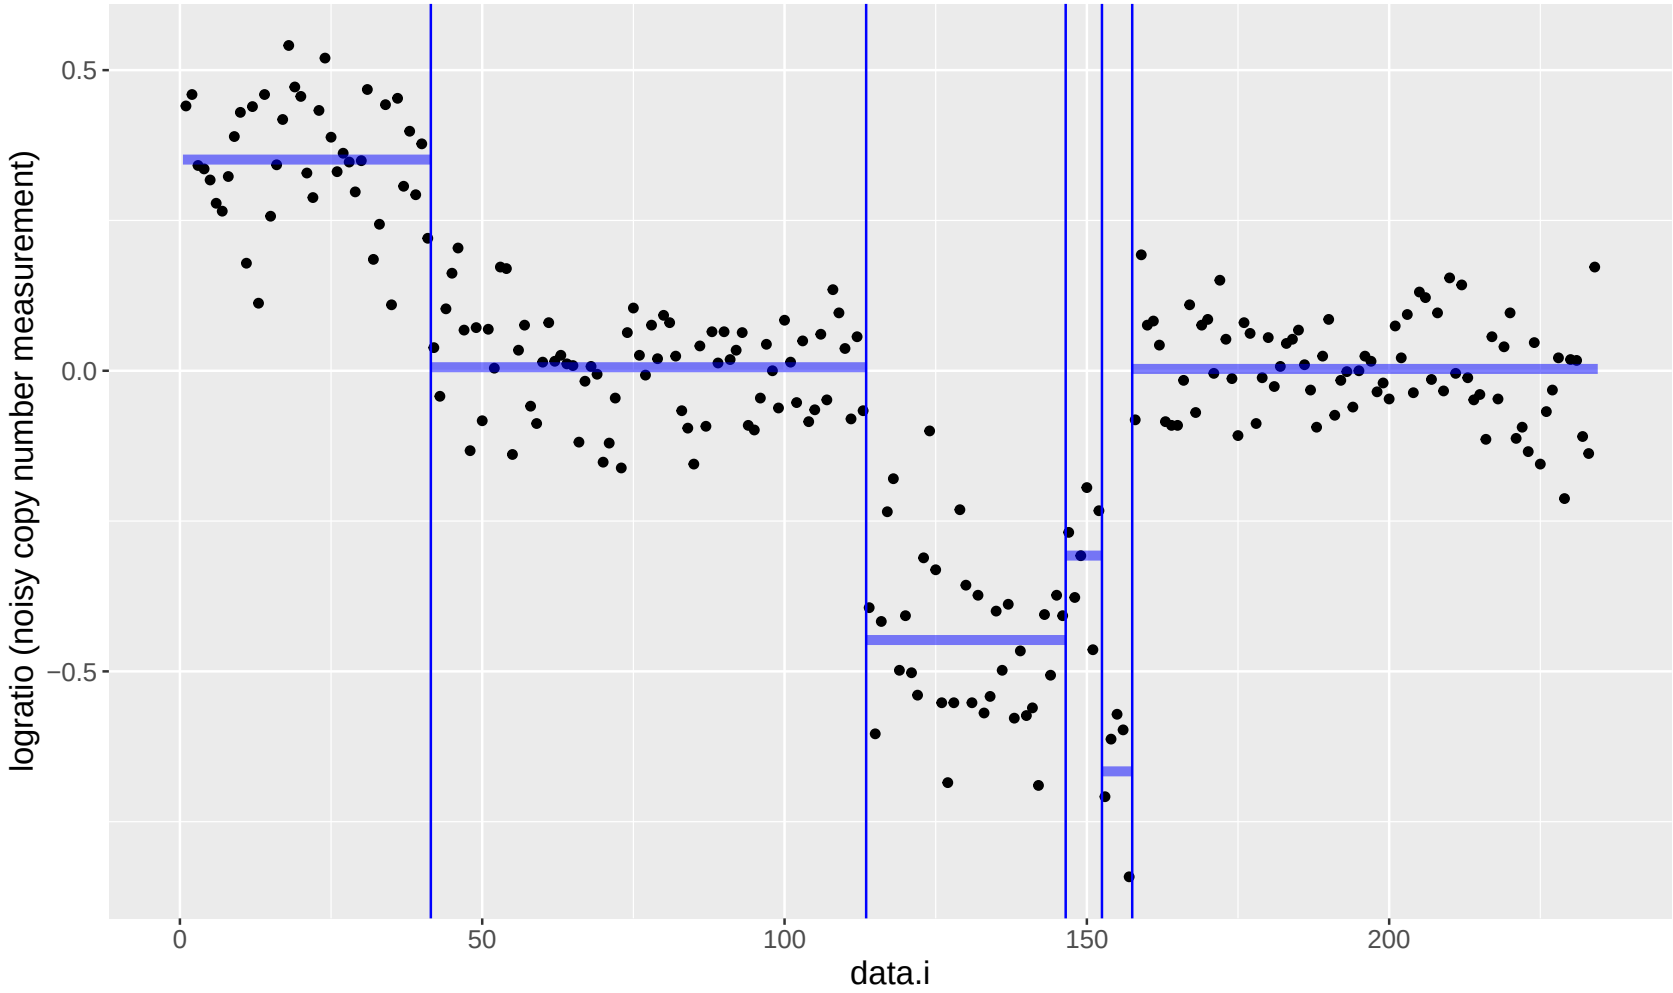

Find seven segments

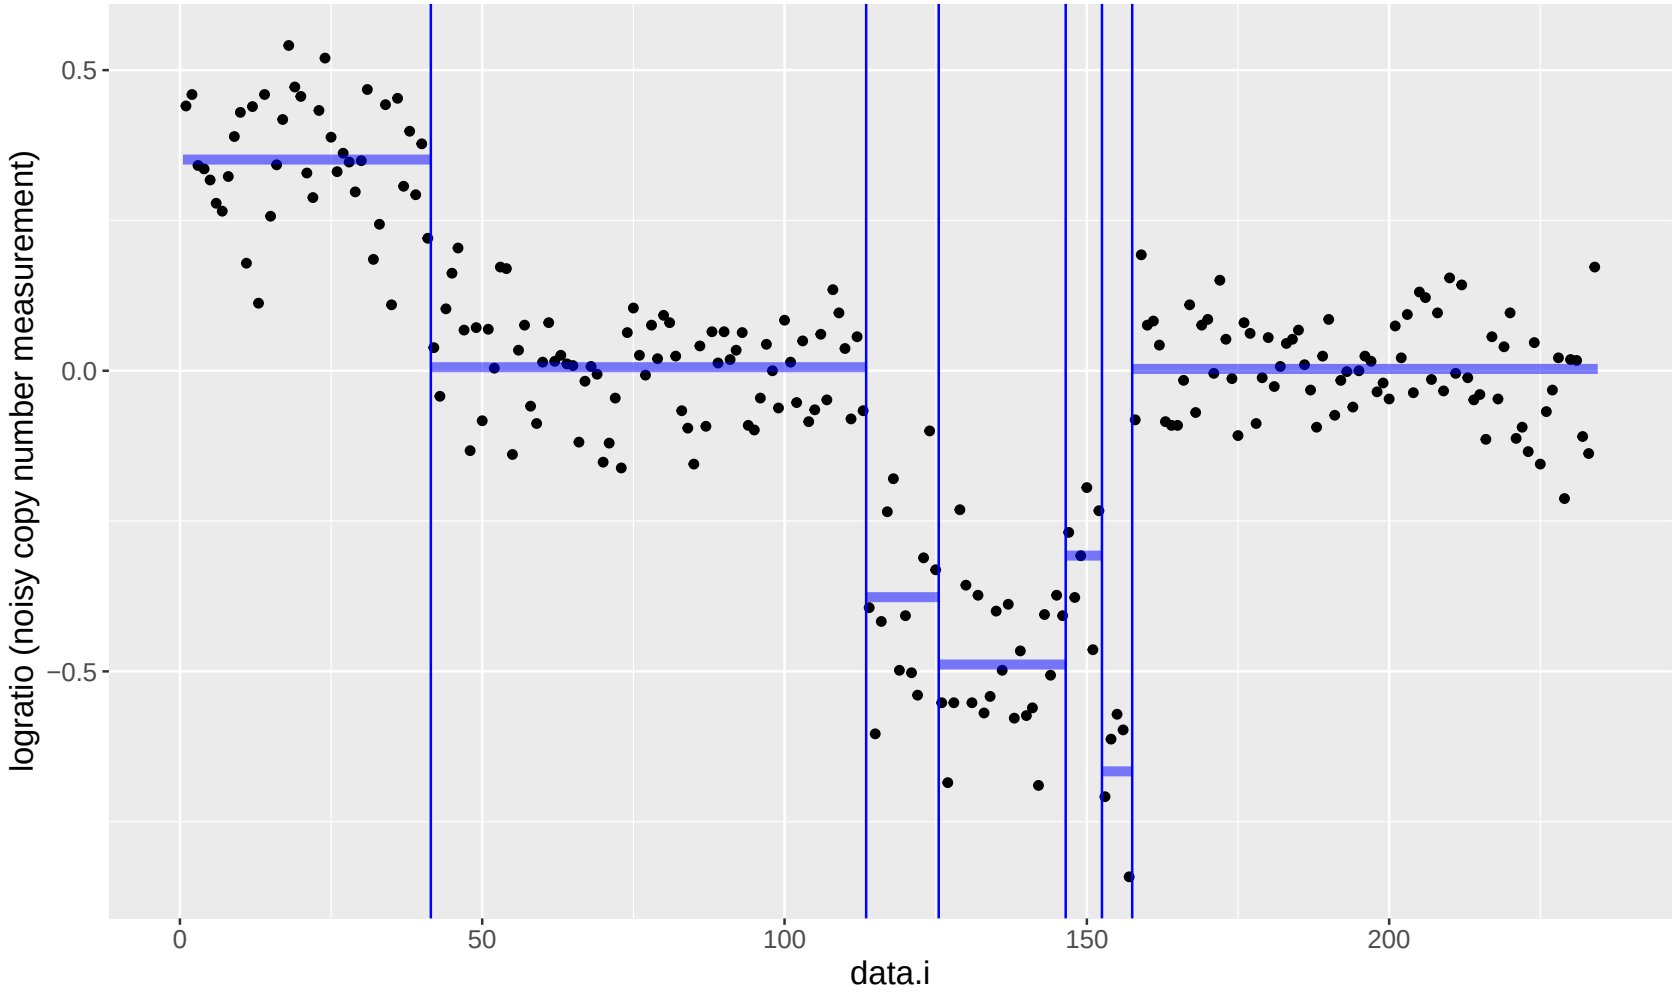

Find 50 segments

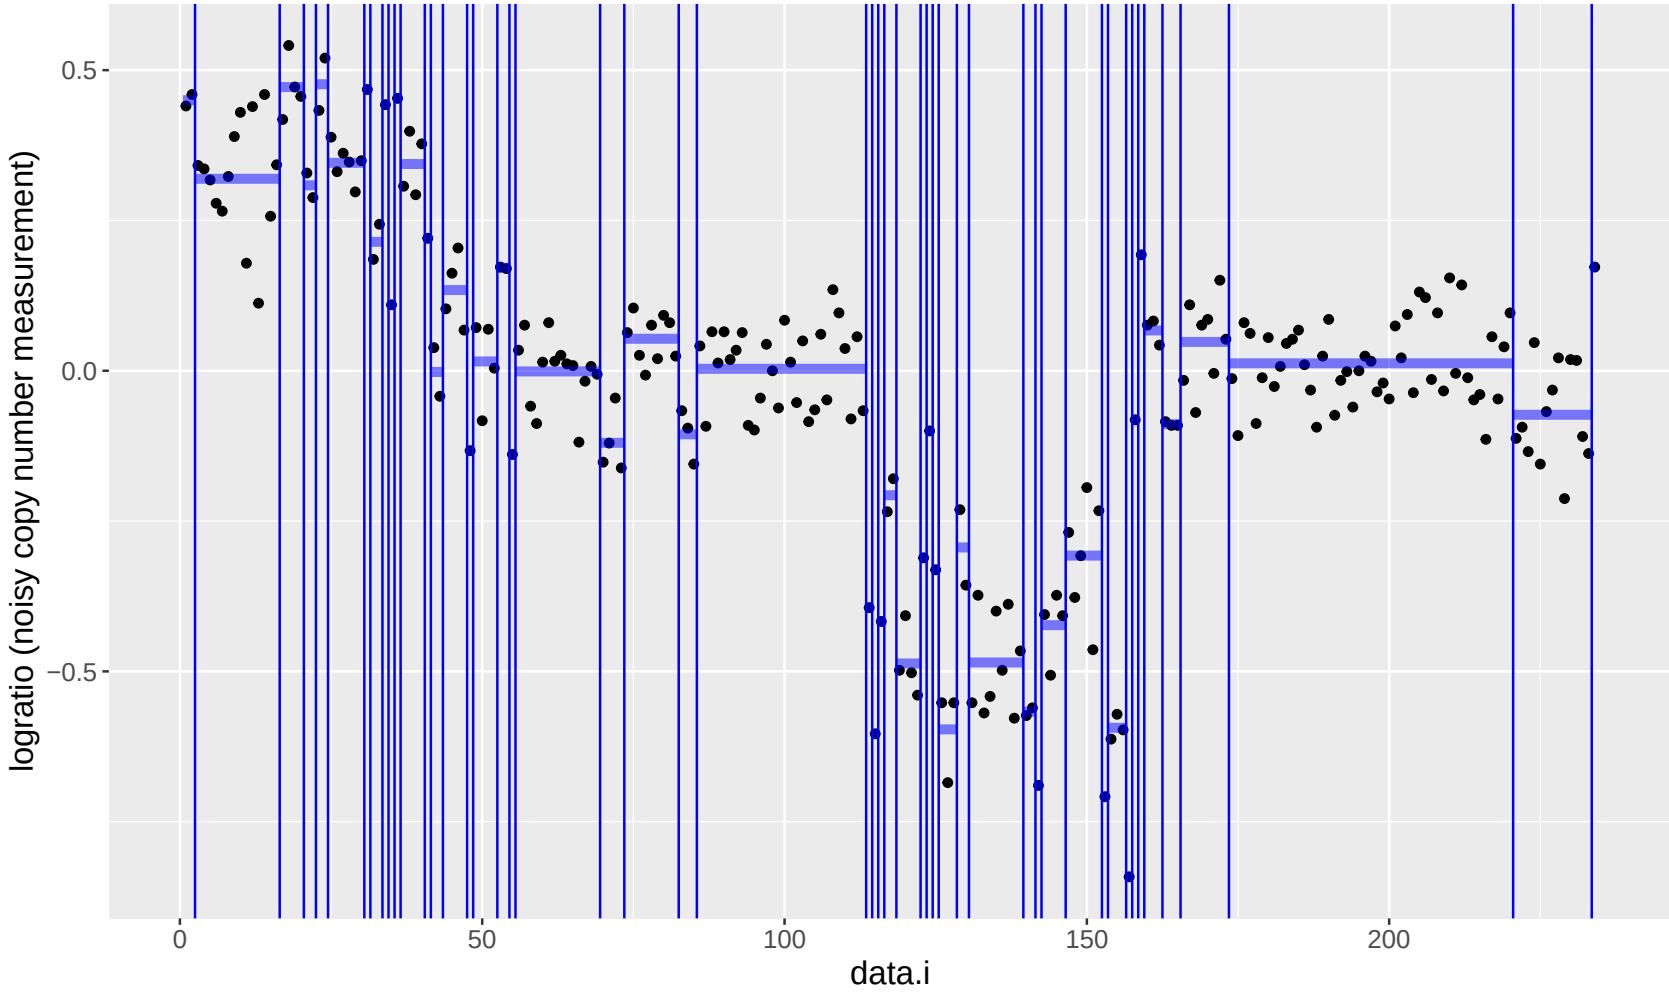

Find 100 segments

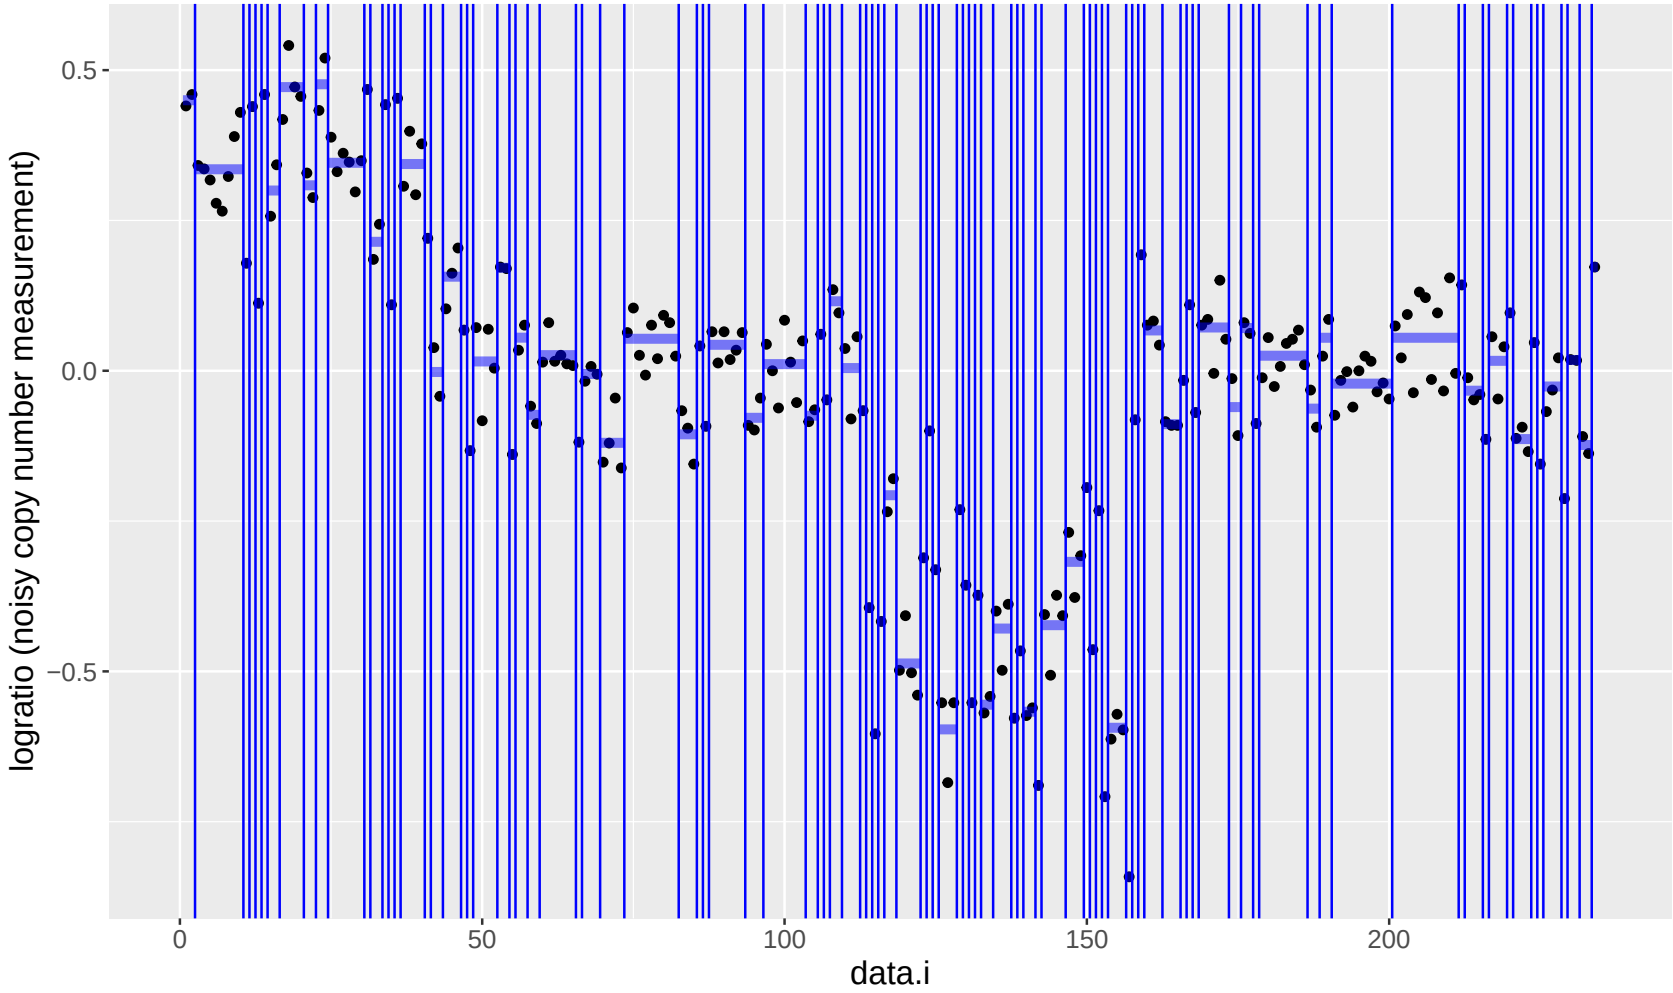

Largest model: 234 segments (changes everywhere)

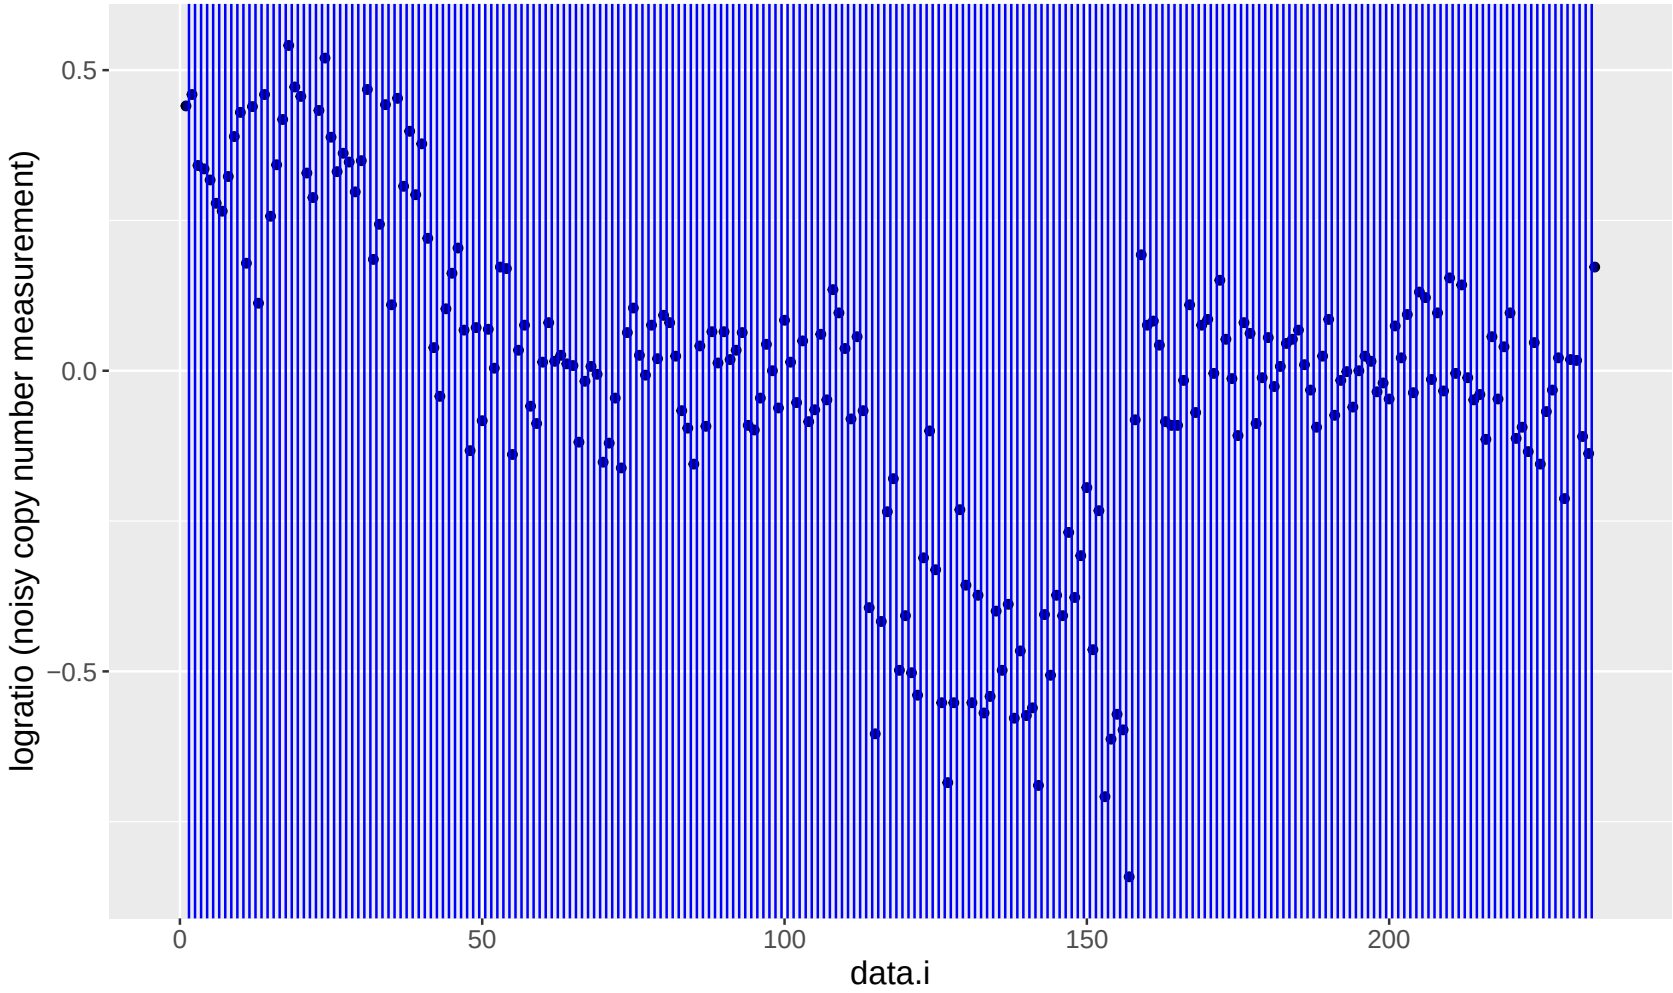

Error/loss function visualization

- Let $m^{(k)} \in \mathbb{R}^n$ be the mean vector with k segments.
- Error for k segments is defined as sum of squared difference between data x and mean $m^{(k)}$ vectors, $E_k = \sum_{i=1}^{n} (x_i - m_i^{(k)})^2$
- As in previous clustering models, kink in the error curve can be used as a simple model selection criterion.

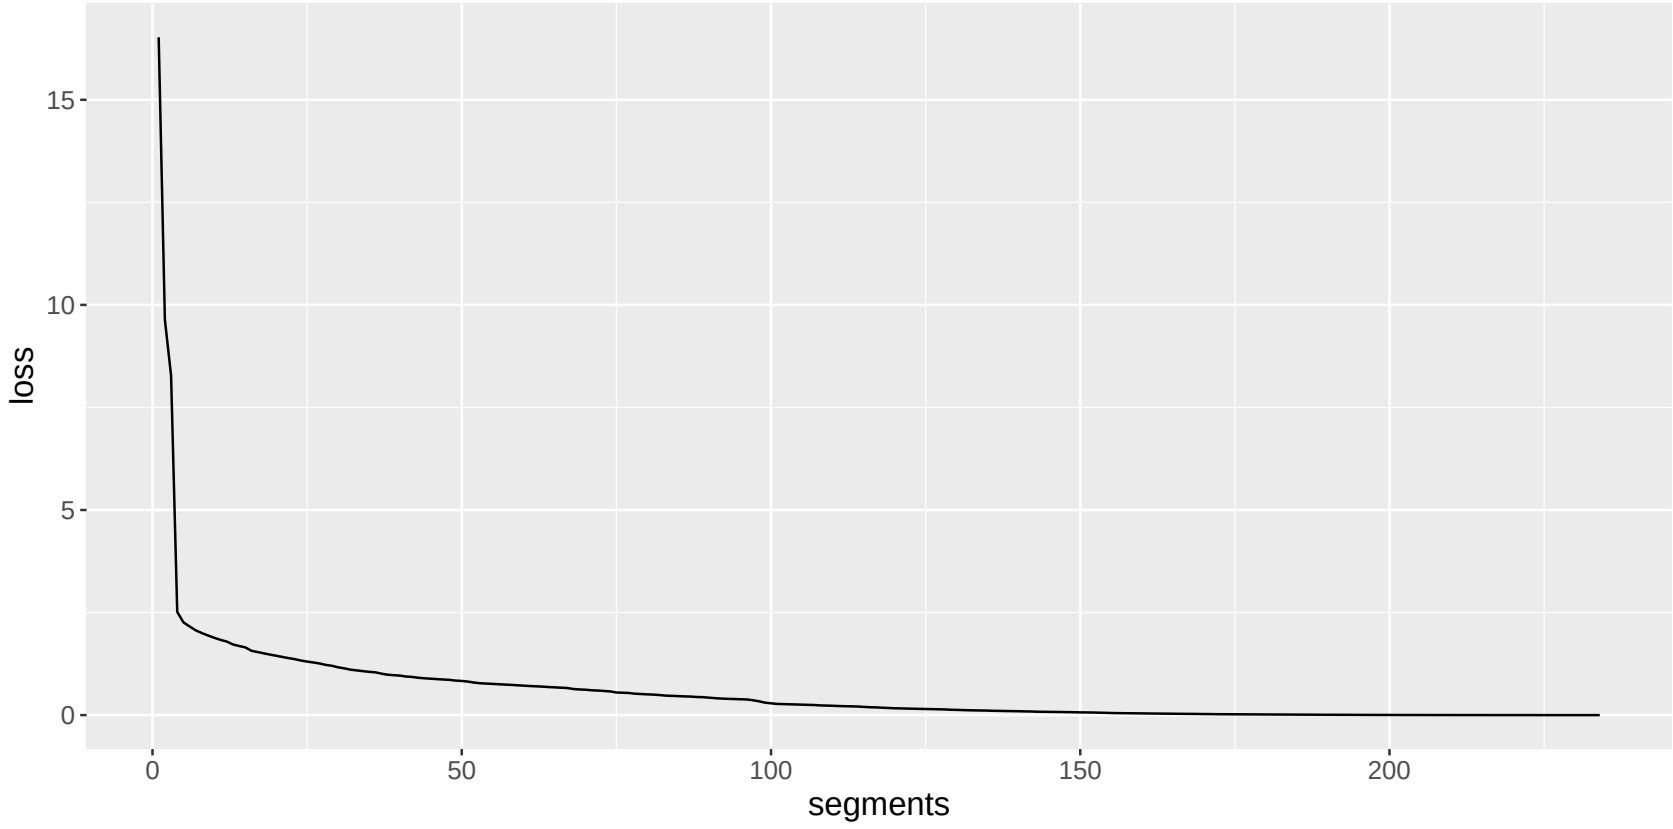

Error/loss function zoom

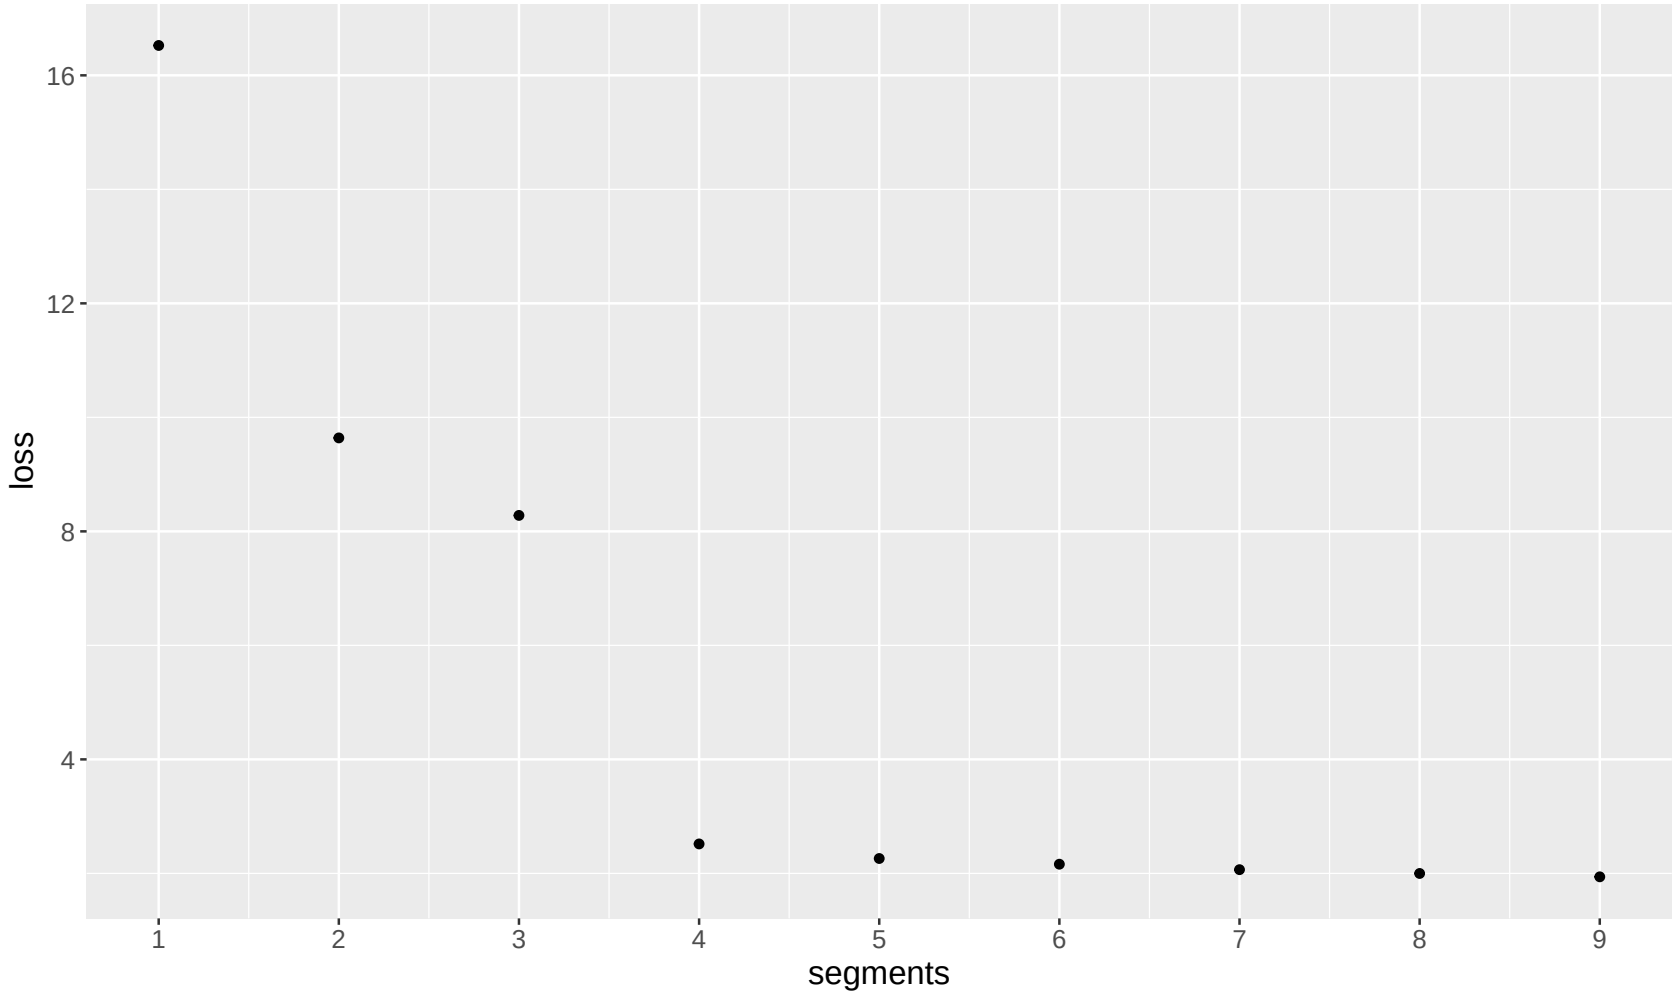

Learning algorithm

- Start with one segment, then repeat:
- Compute loss of each possible split.
- Choose split which results in largest loss decrease.
- If $s = \sum_{i=1}^{n} x_i$ is the sum over n data points, then the mean is s/n and the square loss (from 1 to n) is

$$L_{1,n} = \sum_{i=1}^{n}(x_i - s/n)^2 = \sum_{i=1}^{n}[x_i^2] - 2(s/n)s + n(s/n)^2$$

- Given cumulative sums, and a split point t, we can compute square loss from 1 to t, $L_{1,t}$, and from $t+1$ to n, $L_{t+1,n}$, in constant $O(1)$ time.
- We can minimize over all changepoints t in linear $O(n)$ time,

$$\min_{t \in \{1,\ldots,n-1\}} L_{1,t} + L_{t+1,n}$$

First step of binary segmentation

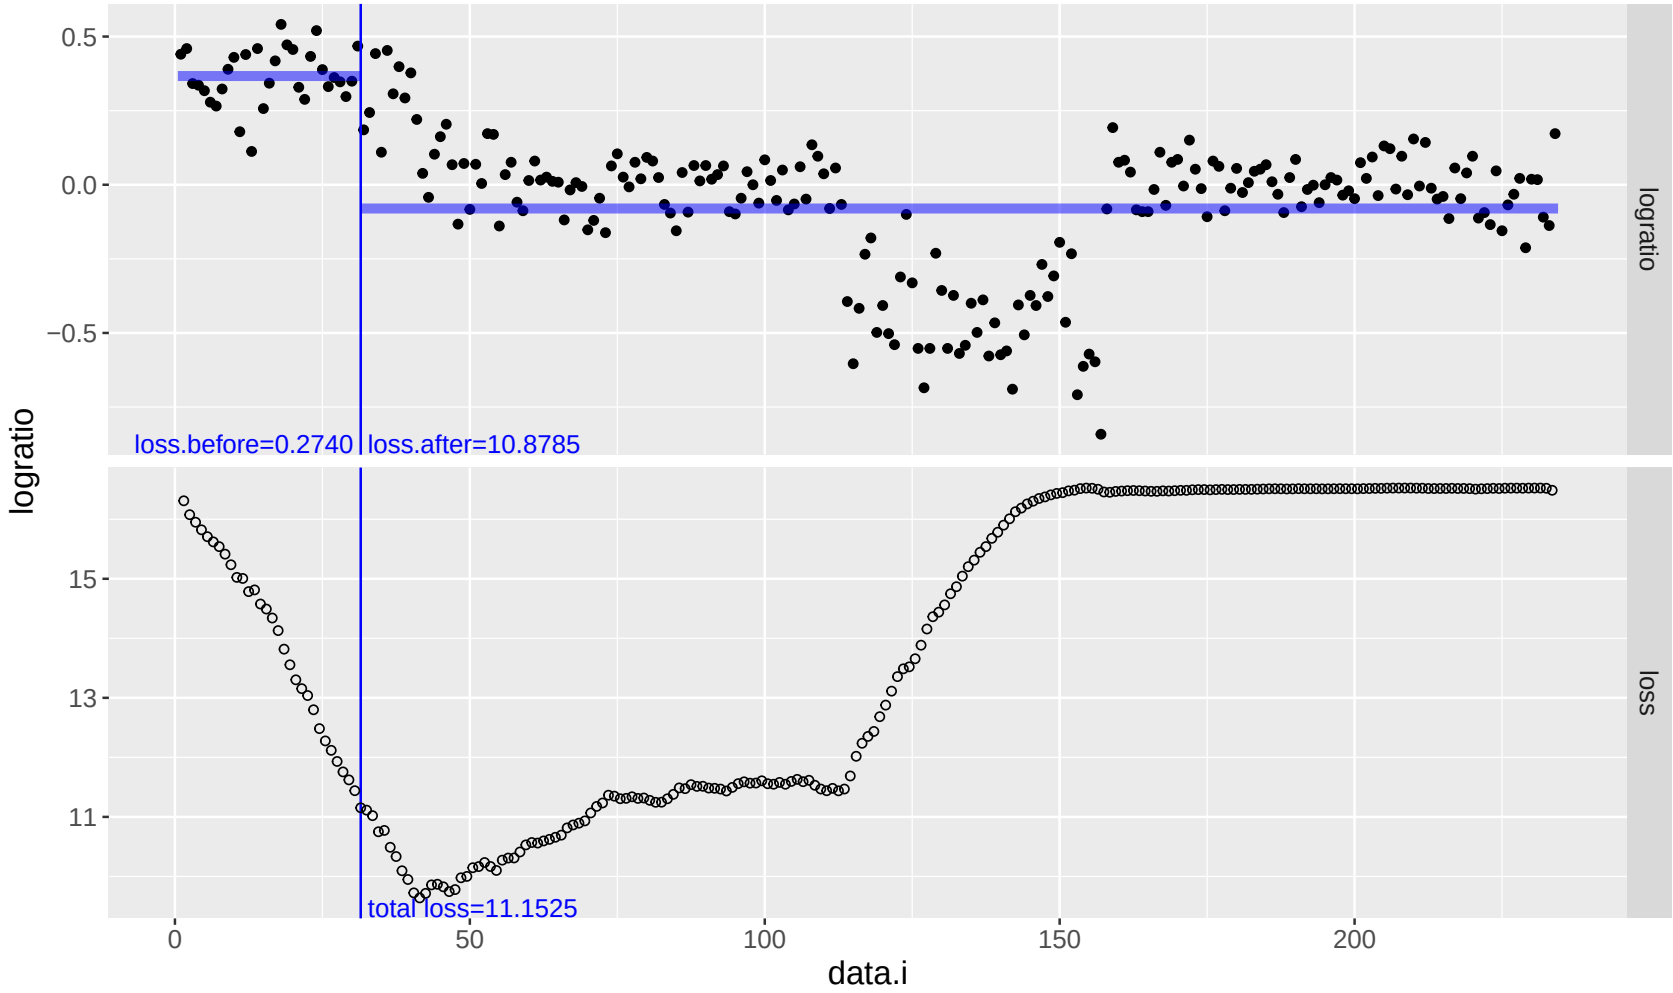

First step of binary segmentation

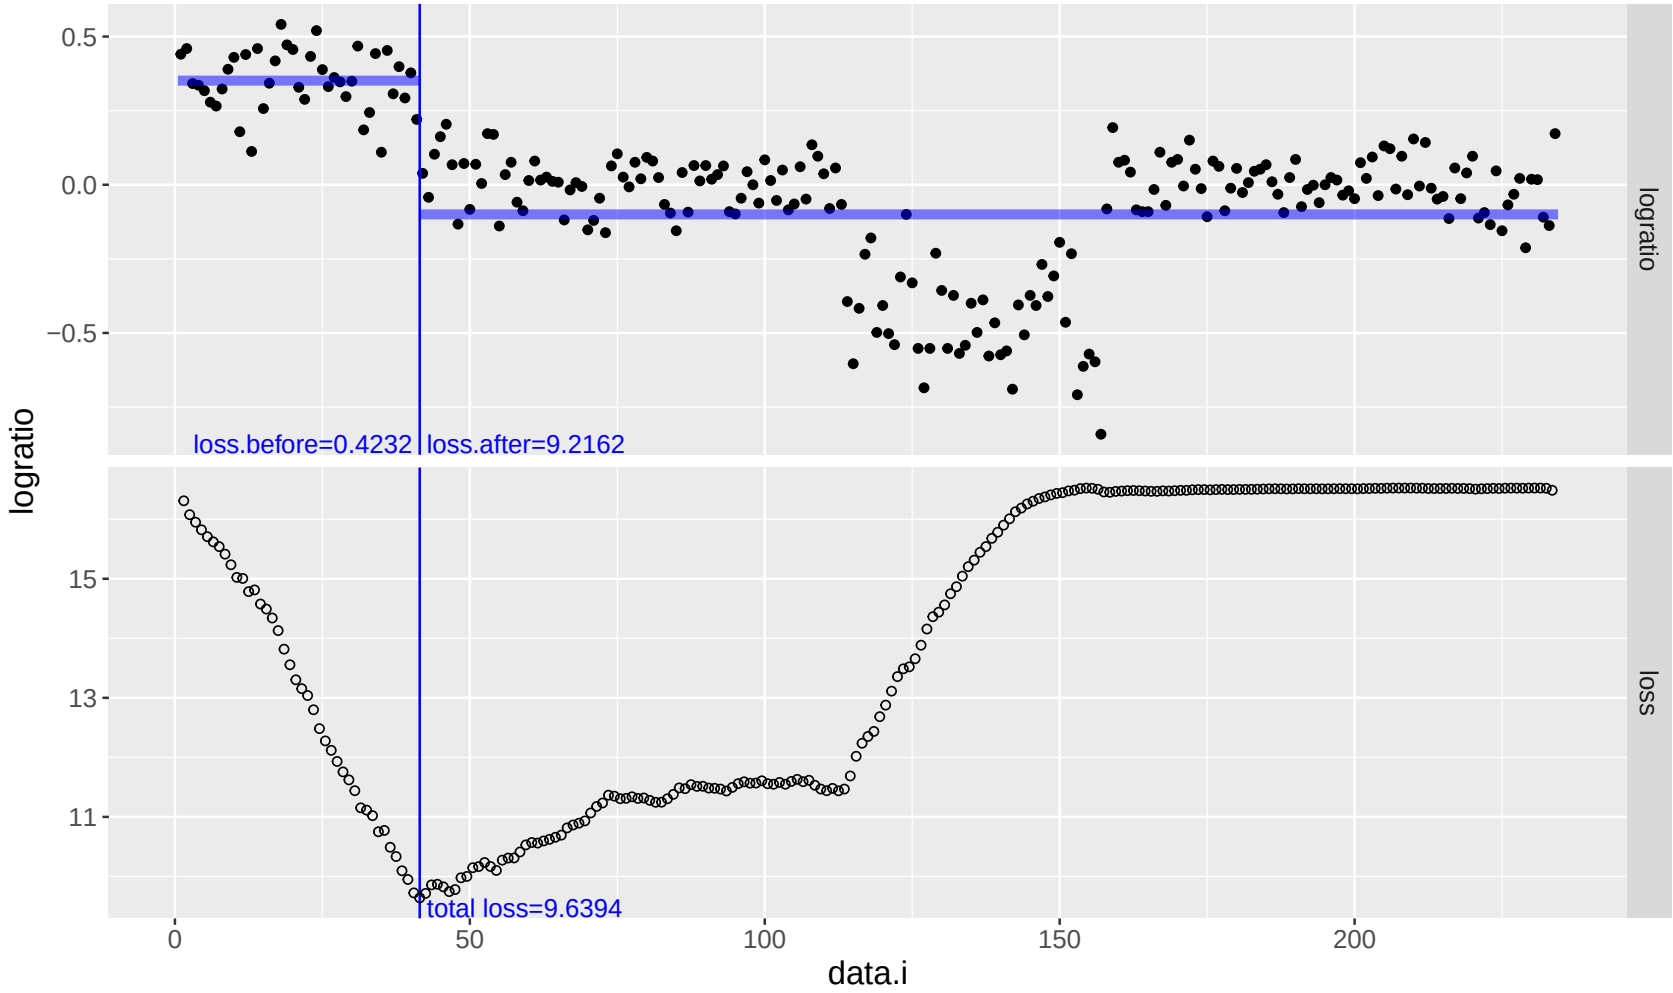

First step of binary segmentation

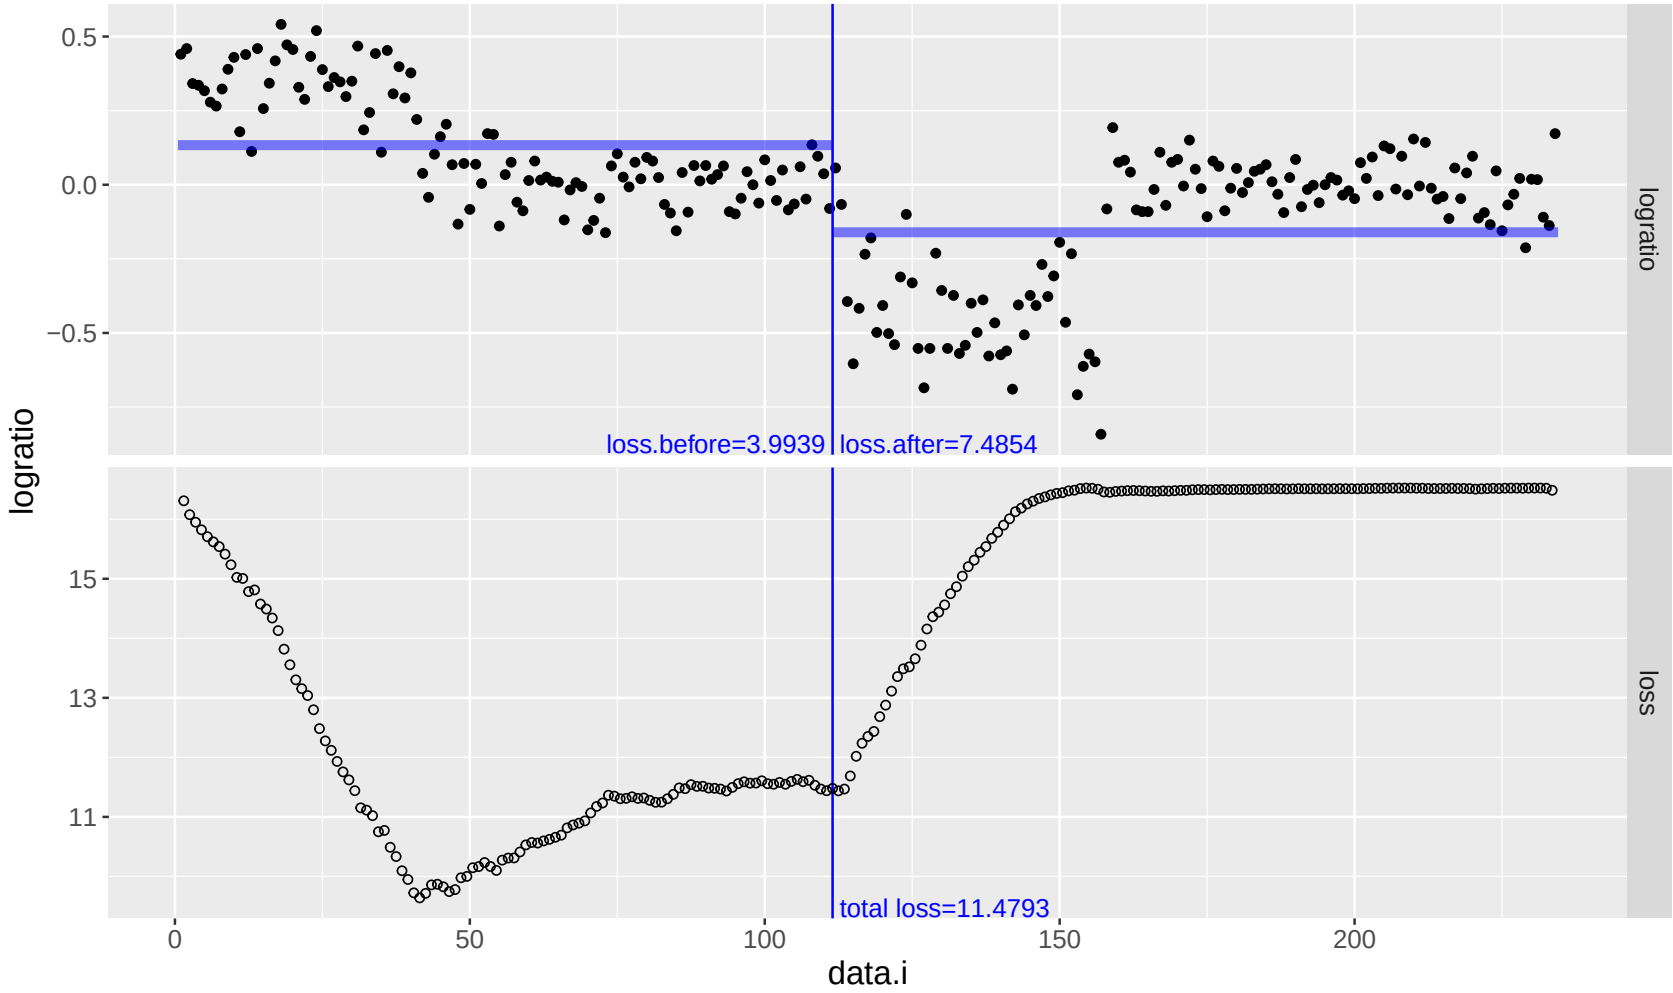

First step of binary segmentation

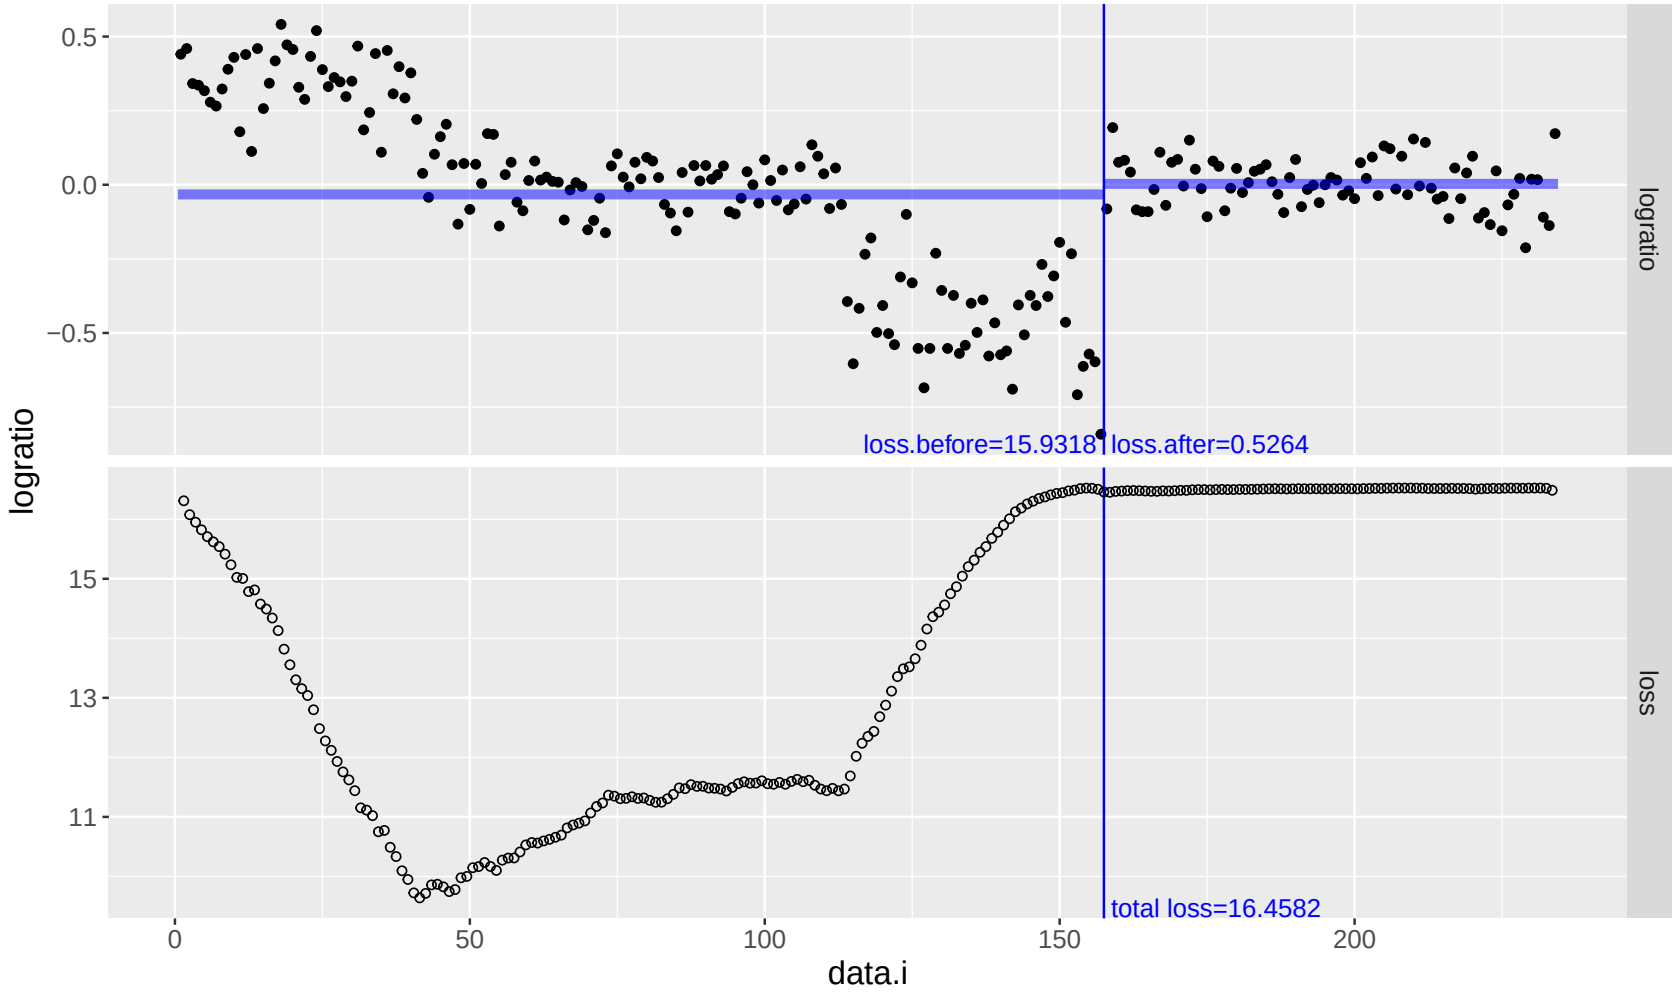

Second step of binary segmentation

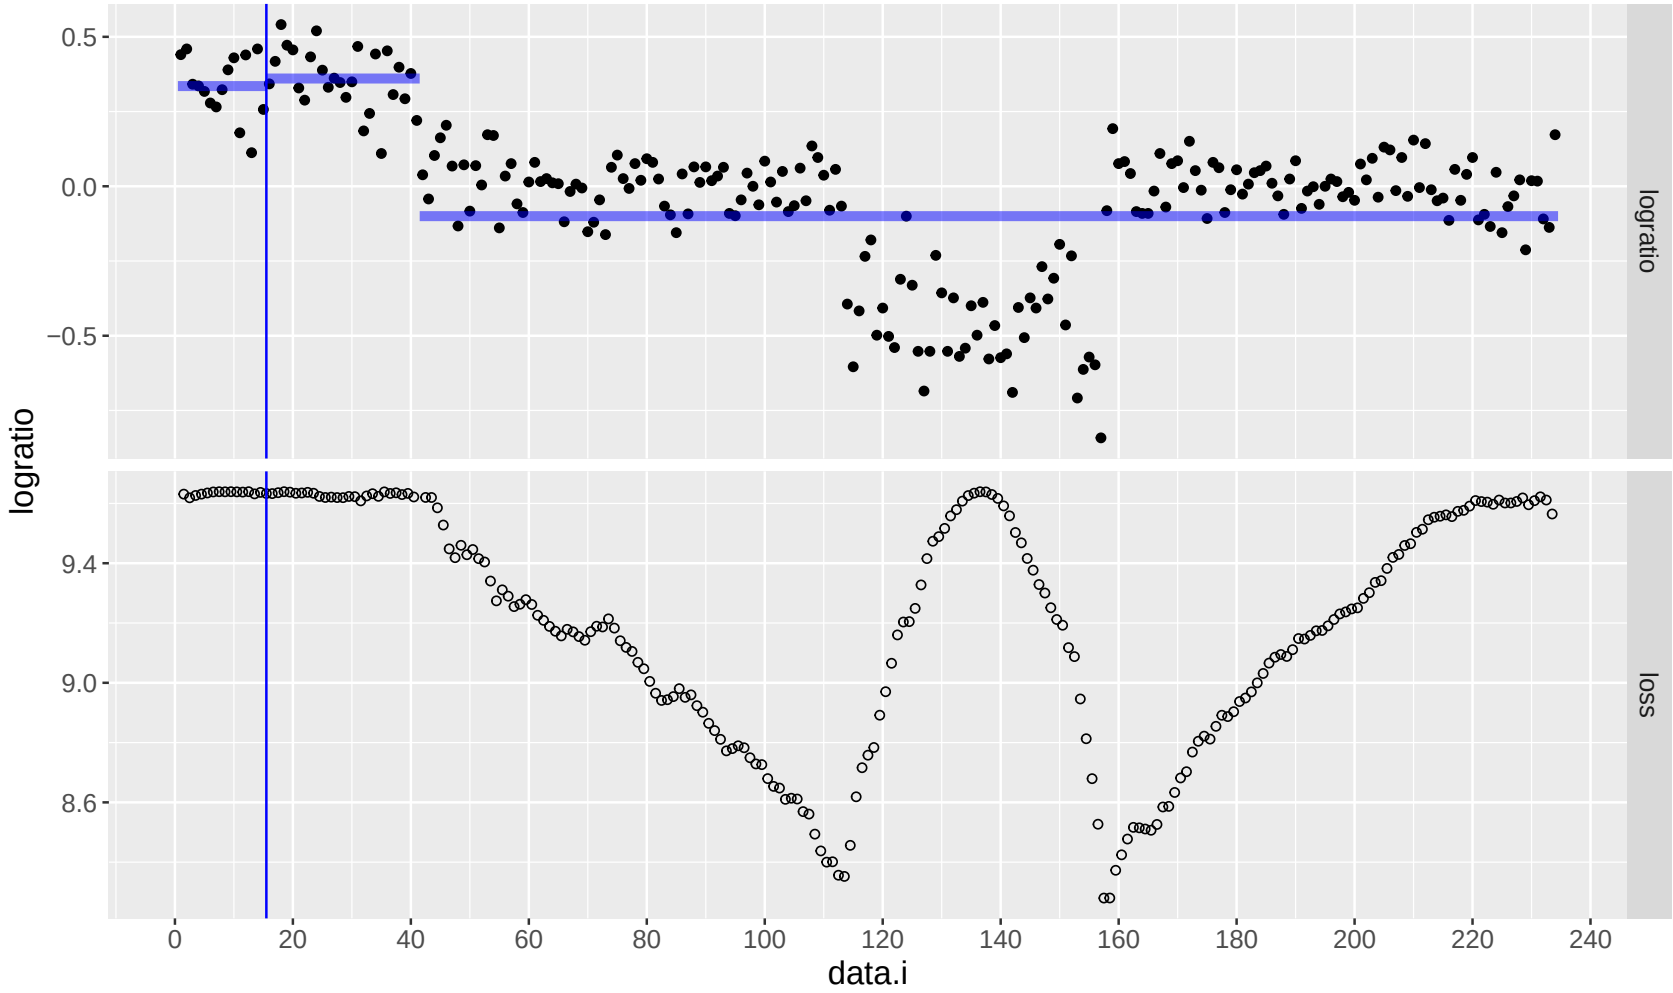

Second step of binary segmentation

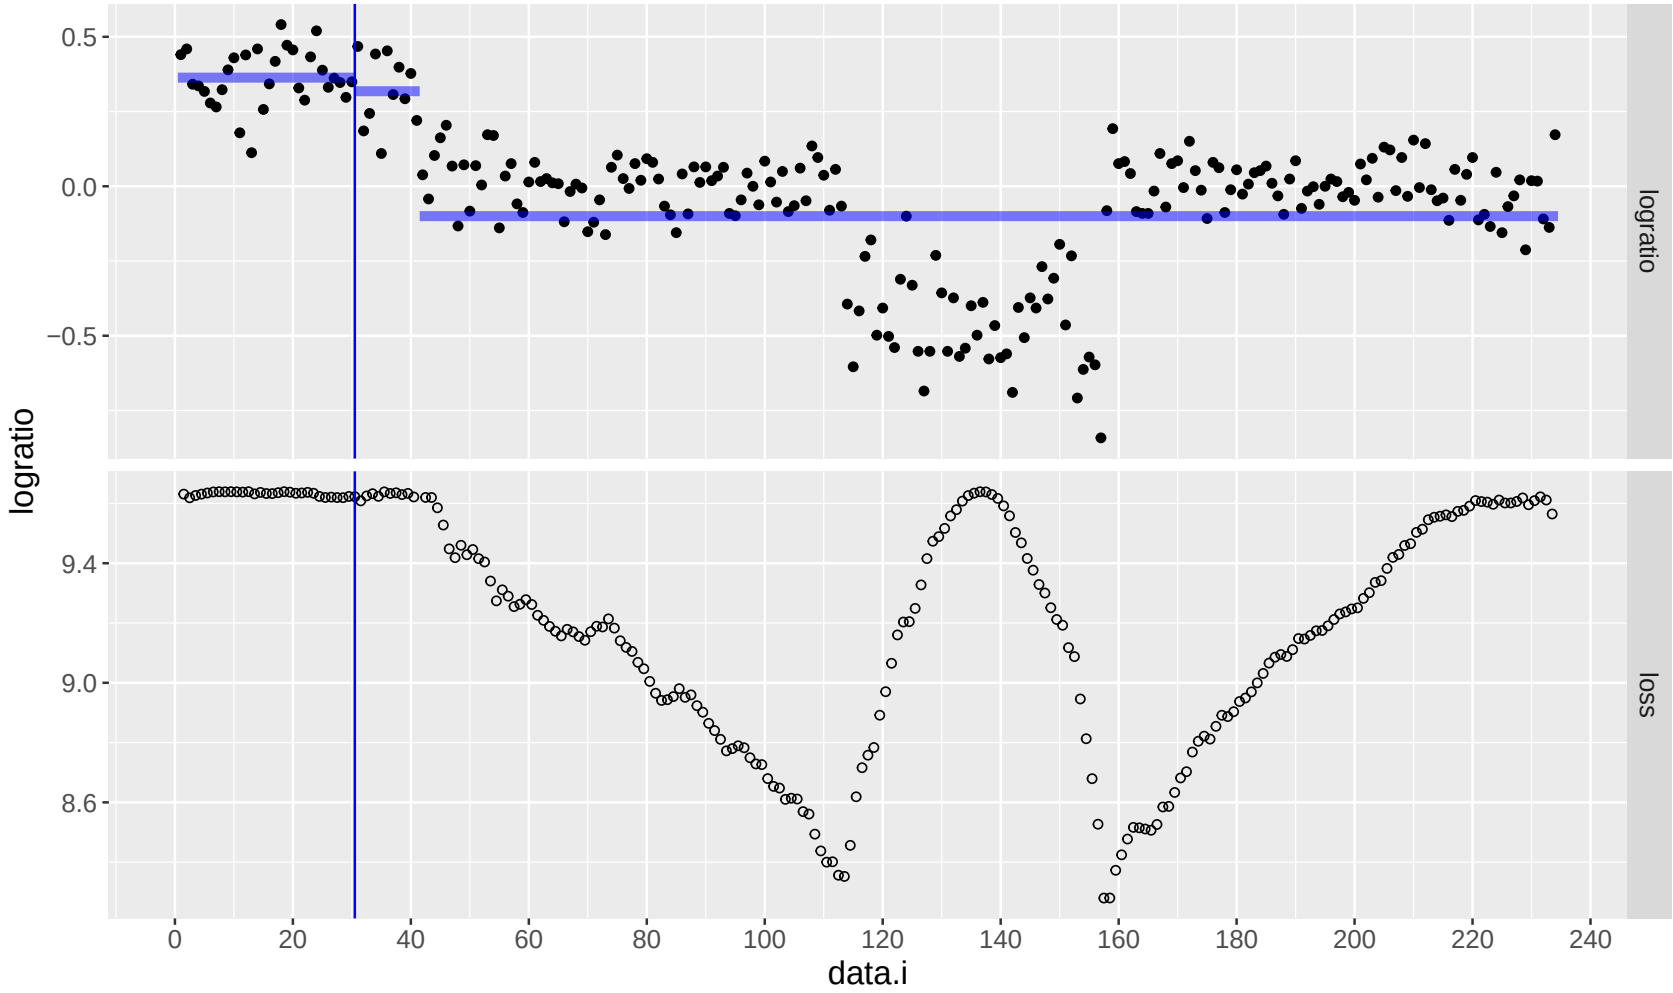

Second step of binary segmentation

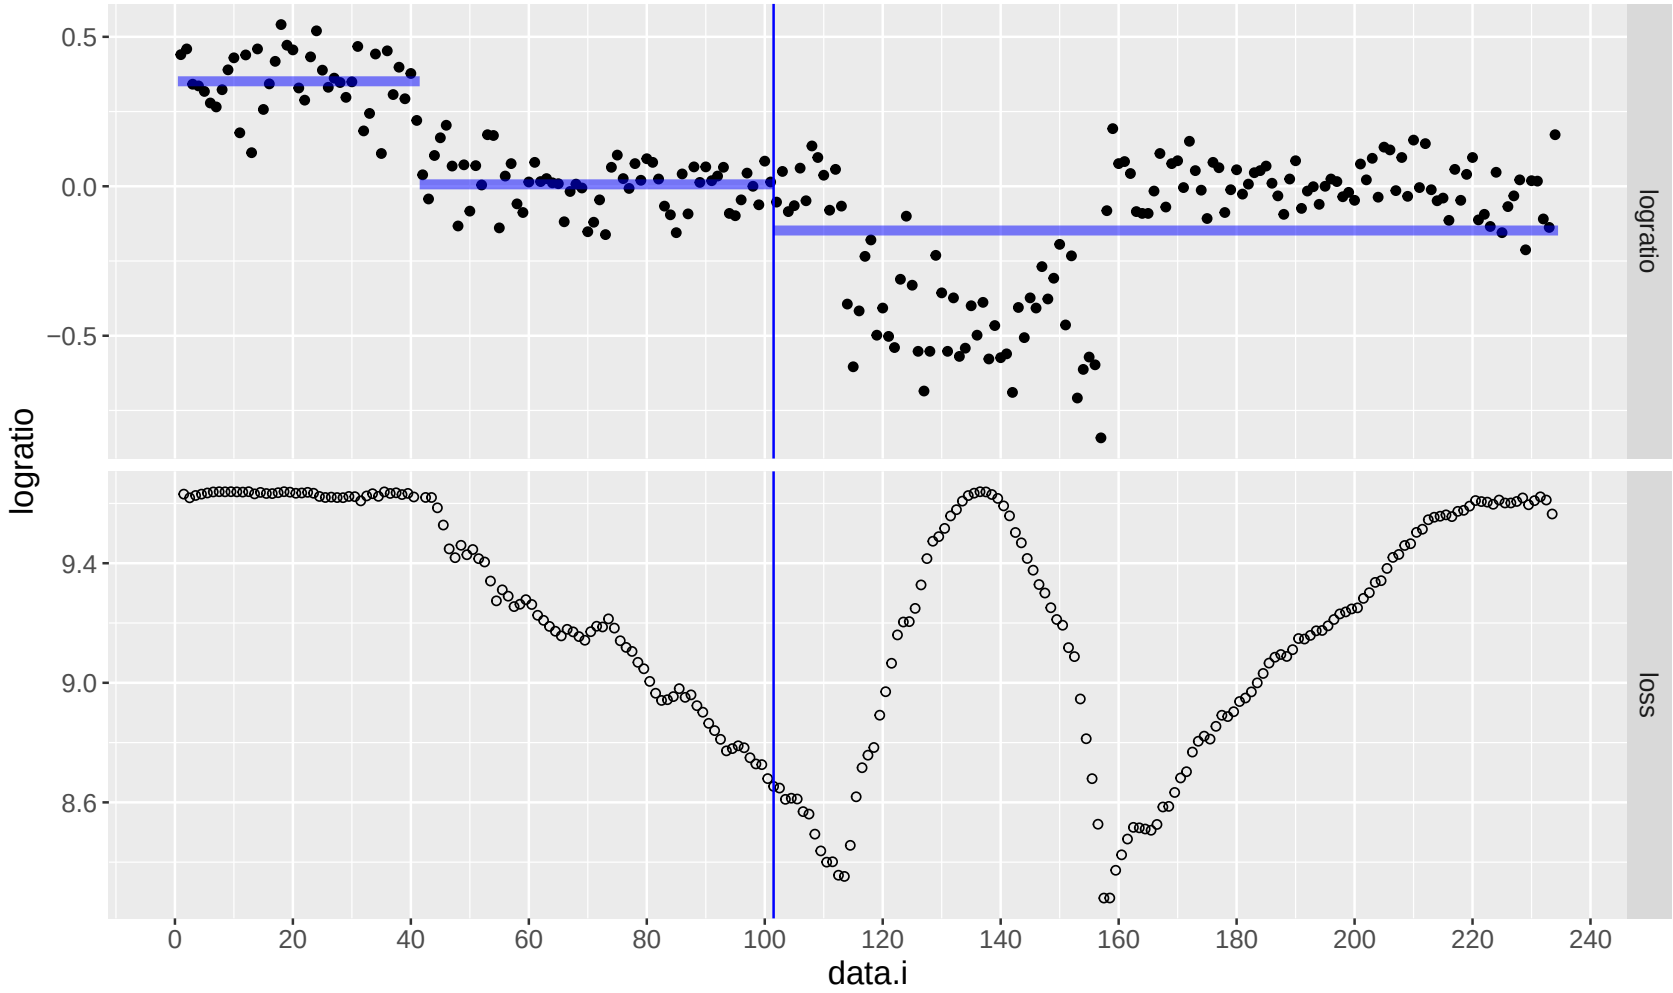

Second step of binary segmentation

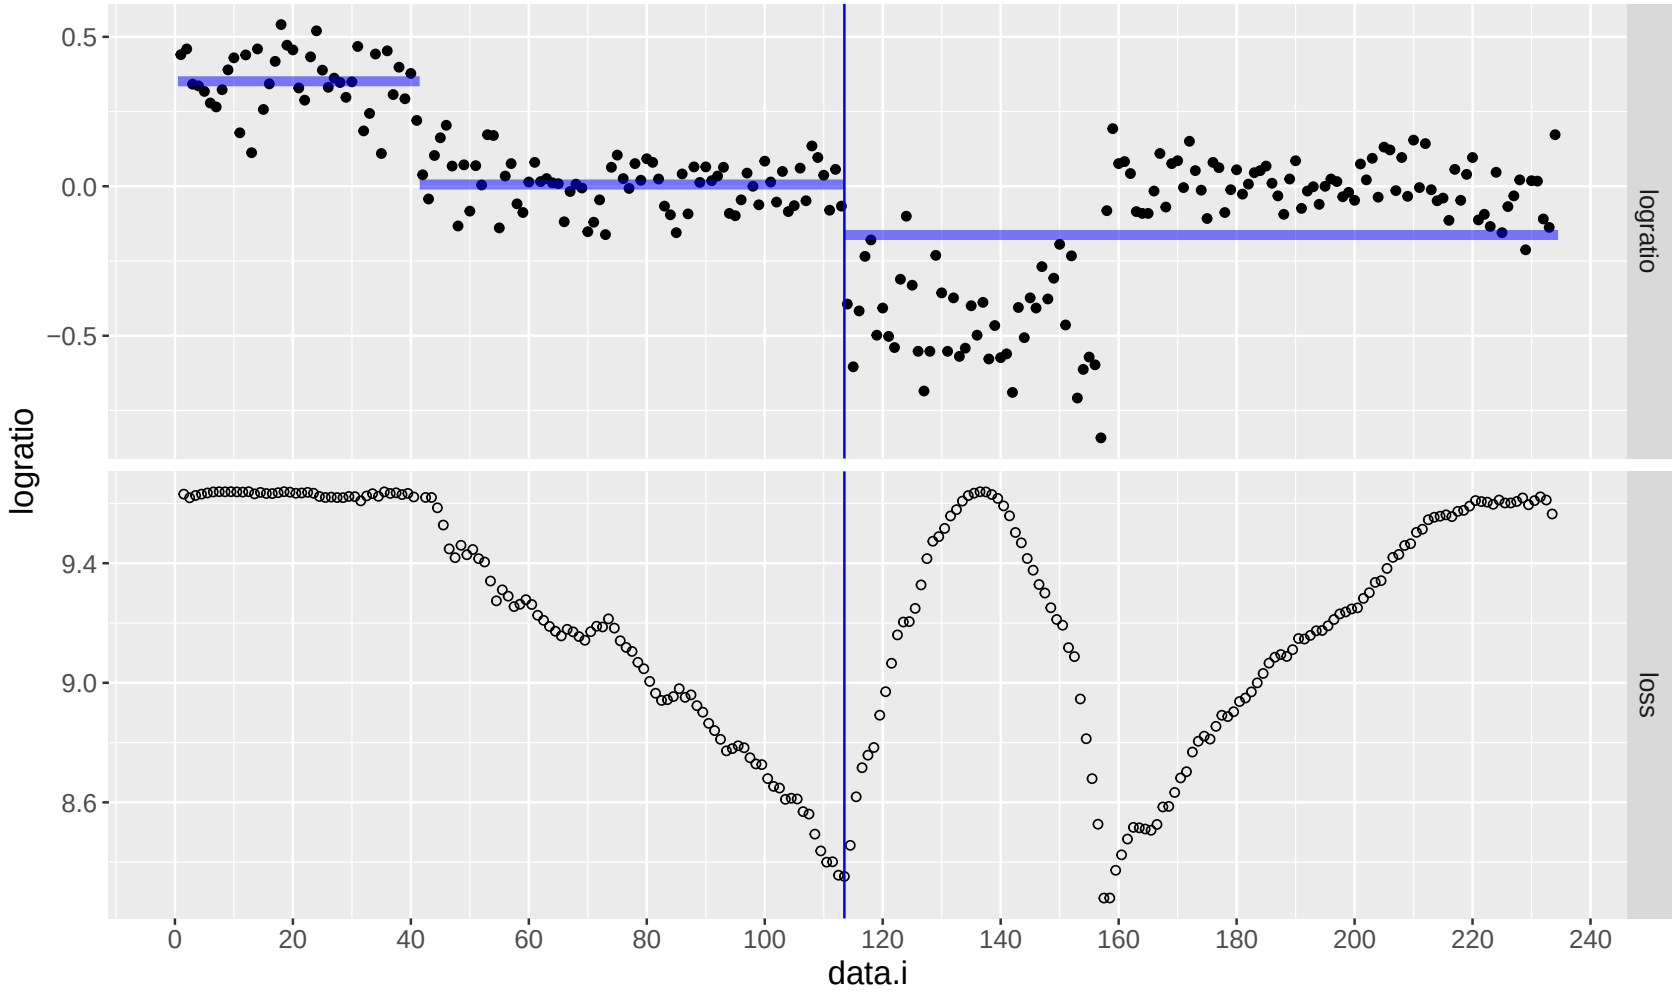

Second step of binary segmentation

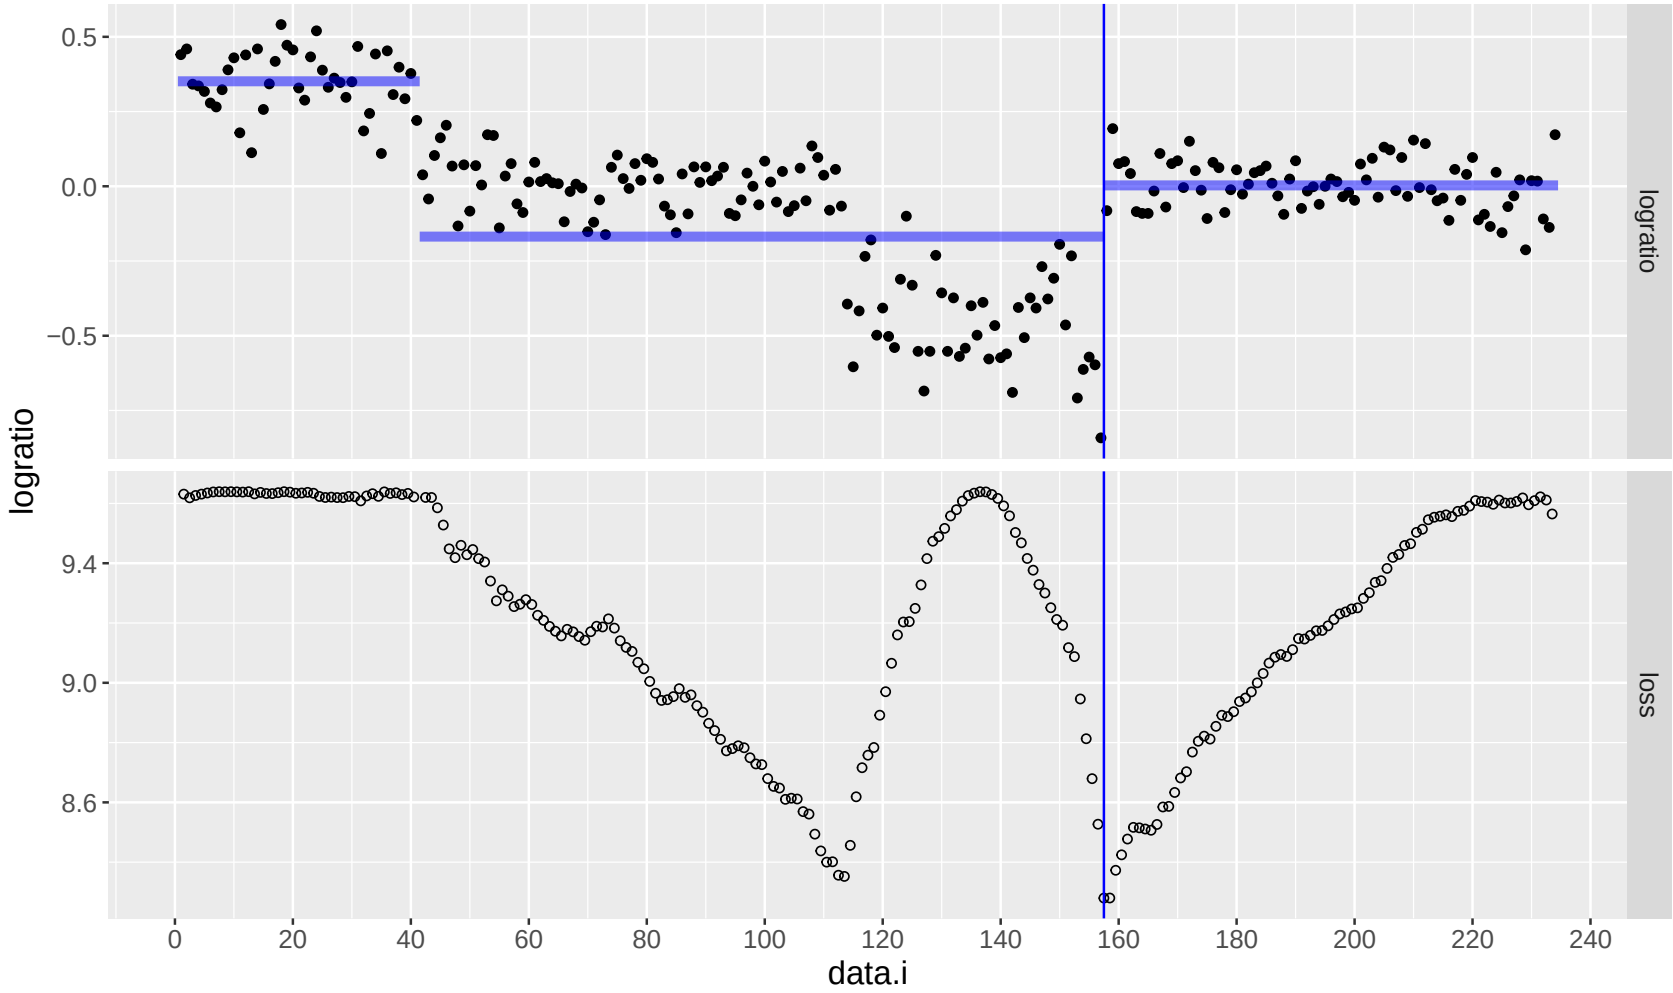

Second step of binary segmentation

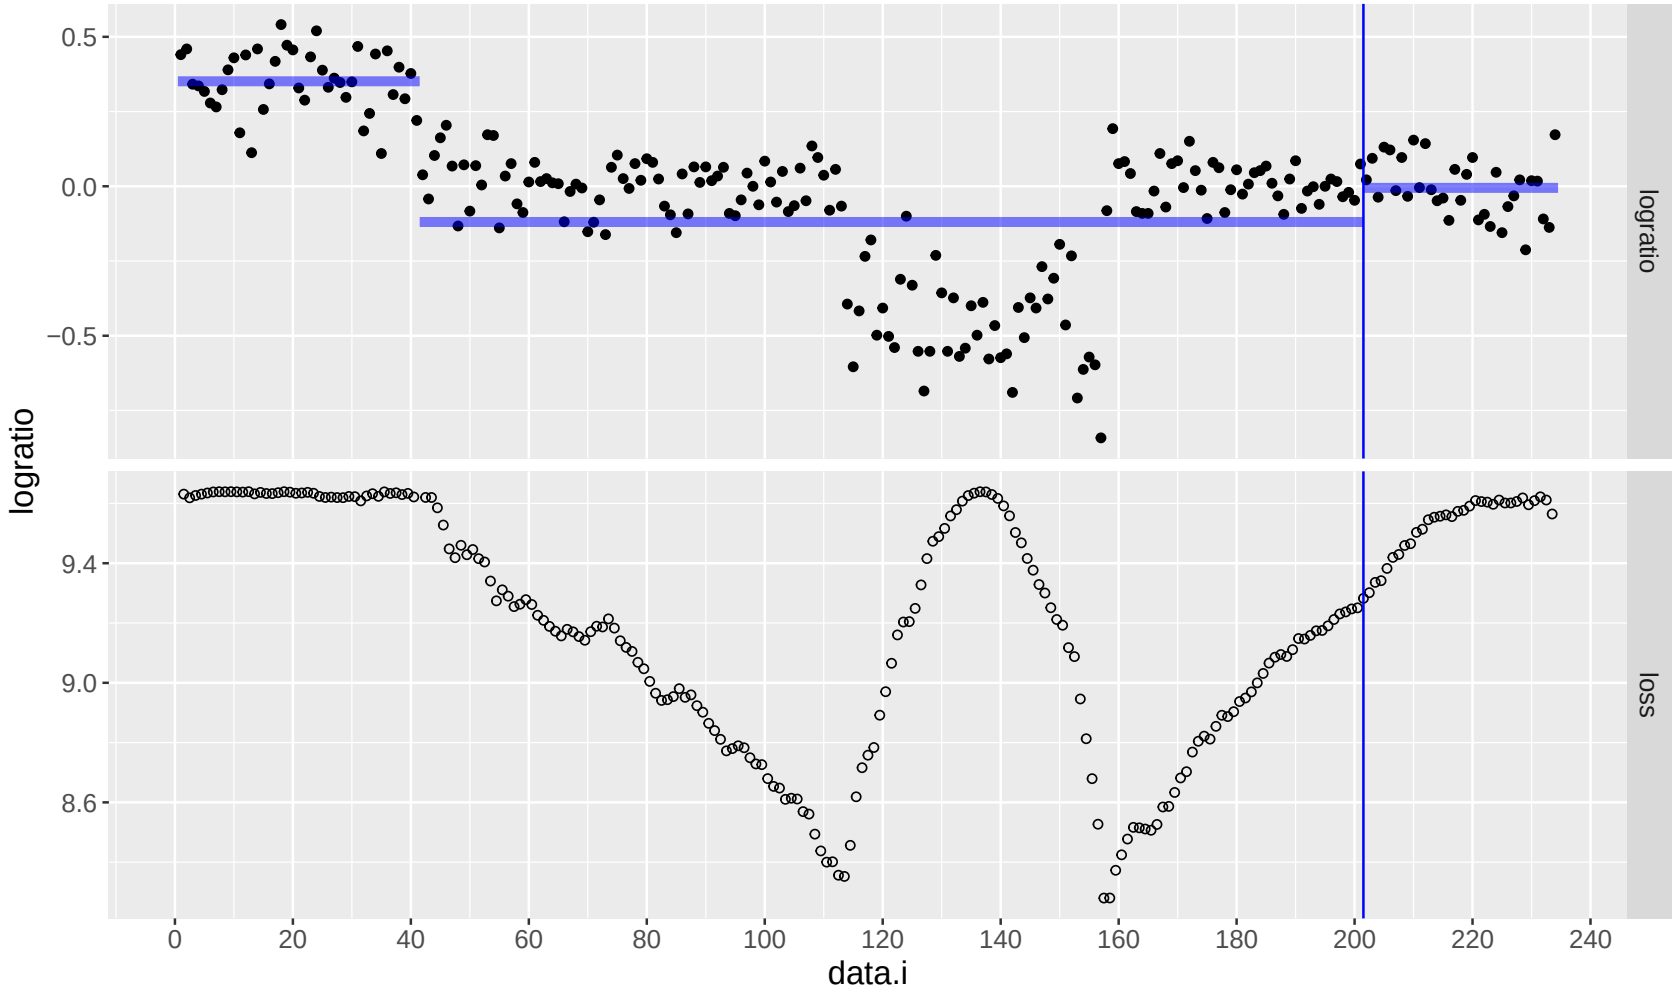

Efficient loss computation

- Minimization can be performed by choosing the split with loss (black point) which maximizes the decrease in loss with respect to previous model (red, with no split).

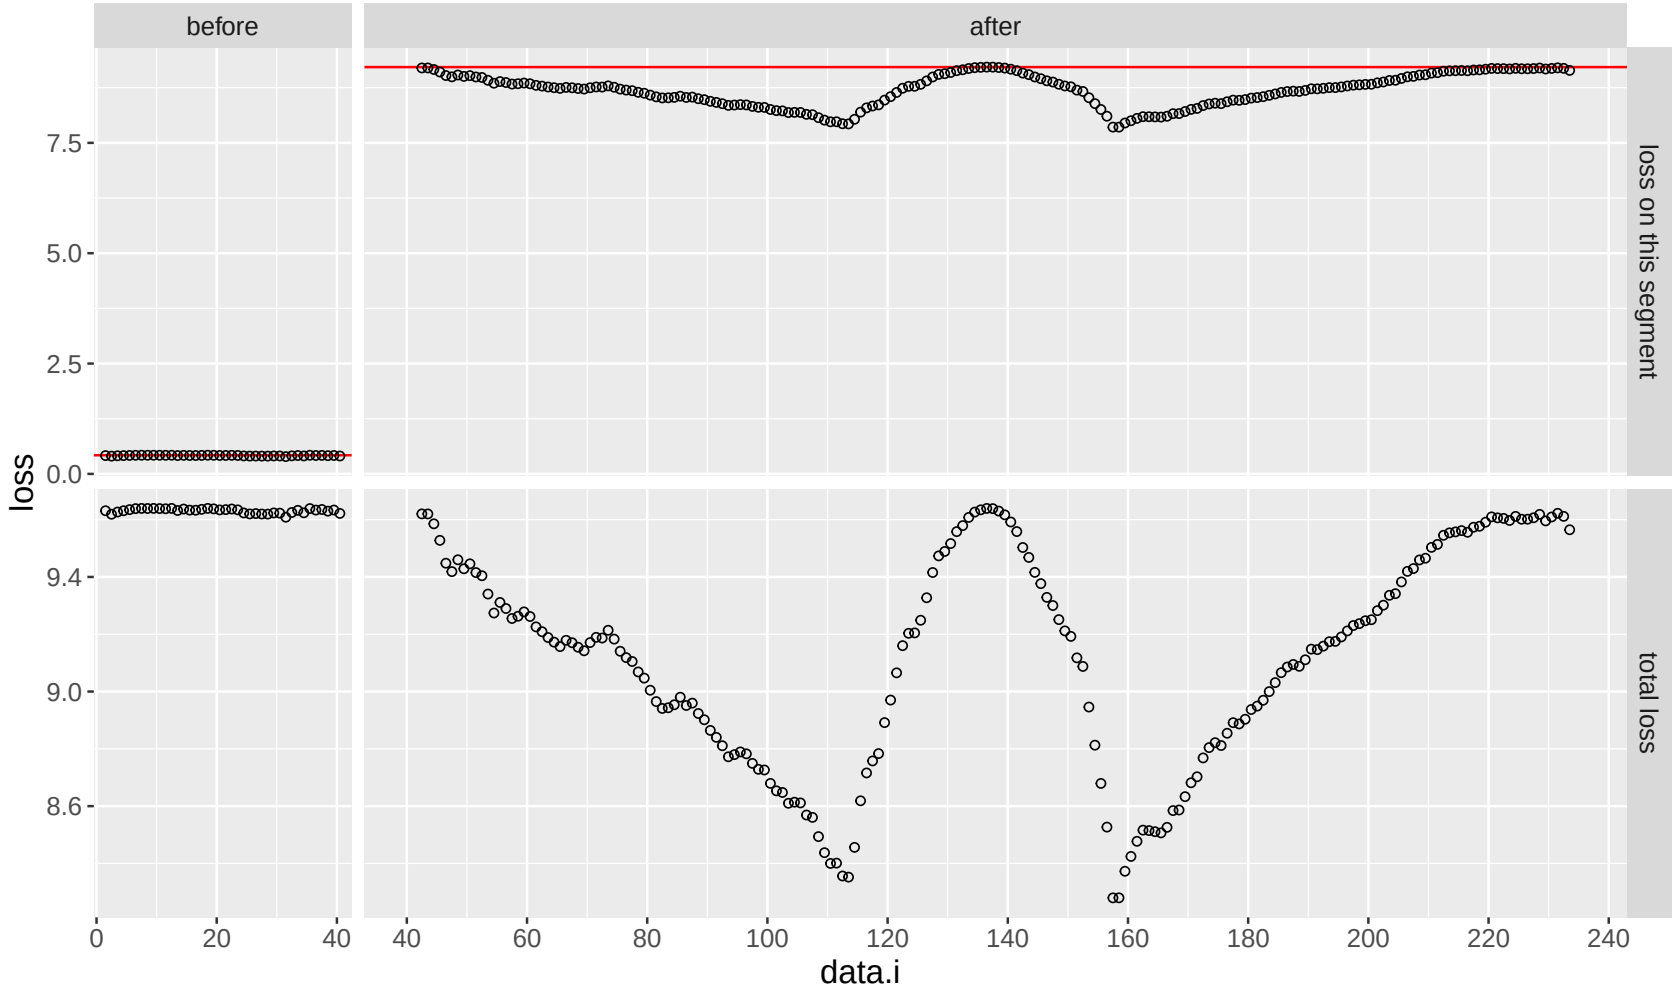

Learning algorithm, implementation details

First compute the vectors of cumulative sums of data and squares, $y_1 = z_1 = 0, y_t = \sum_{i=1}^{t-1} x_i, z_t = \sum_{i=1}^{t-1} x_i^2$, for all $t \in \{2, \ldots, n+1\}$.

Assume there is some set $\mathcal{S}$ of segments that could be split, each $(j, e) \in \mathcal{S}$ is a segment start j and end e (both in $1, \ldots, n$).

Then the next segment to split (j, e), and best split point t, are defined by the best loss decrease,

$$\max_{(j,e)\in\mathcal{S}} \underbrace{L_{j,e}}_{\text{loss before split}} - \min_{t\in\{j,\ldots,e-1\}} \underbrace{L_{j,t} + L_{t+1,e}}_{\text{loss after split}}$$

Use the cumsum trick to compute loss in constant time,

$$\sum_{i=j}^{t} x_i = y_{t+1} - y_j.$$

$$L_{j,t} = z_{t+1} - z_j - (y_{t+1} - y_j)^2/(t - j + 1).$$

Learning algorithm, recursion/pseudo-code

Notation. Let $D_{j,e}(t) = L_{j,t} + L_{t+1,e} - L_{j,e}$ be the loss difference after splitting segment (j, e) at t, and let $f(j, e) = \min, \arg\min_{t \in \{j,\ldots,e-1\}} D_{j,e}(t)$ be the best loss difference/split on segment (j, e).

Initialization. Let $\mathcal{L}_1 = L_{1,n}$ be the loss with one segment, and let $\mathcal{S}_1 = \{(1, n, f(1, n))\}$ be the initial segment to split.

Recursion. For all $k \in \{2, \ldots, K\}$: (max segments $K \leq n$)

- $j_k^*, e_k^*, d_k^*, t_k^* = \arg\min_{(j,e,d,t) \in \mathcal{S}_{k-1}} d$ (best segment to split)
- $\mathcal{L}_k = \mathcal{L}_{k-1} + d_k^*$, (loss)
- $\mathcal{N}_k = \{(j_k^*, t_k^*), (t_k^* + 1, e_k^*)\}$ (new segments)
- $\mathcal{V}_k = \{(j, e, f(j, e)) \mid (j, e) \in \mathcal{N}_k, j < e\}$ (splittable segments)
- $\mathcal{S}_k = [\mathcal{S}_{k-1} \setminus (j_k^*, e_k^*, d_k^*, t_k^*)] \cup \mathcal{V}_k$ (segments to search)

Modification for min segment length

Sometimes there is prior knowledge that there should be no segments with fewer data points than $\ell \in \{1, 2, \dots\}$, and in that case there is a simple modification:

- $f(j, e) = \min, \arg\min_{t \in \{j+\ell-1, \dots, e-\ell\}} D_{j,e}(t)$ (best split)
- $\mathcal{V}_k = \{(j, e, f(j, e)) \mid (j, e) \in \mathcal{N}_k,\ e - j + 1 \geq 2\ell\}\}$ (splittable segments)

Introduction

Analysis of time complexity

Comparing empirical timings between implementations

Visualization of computations at each iteration

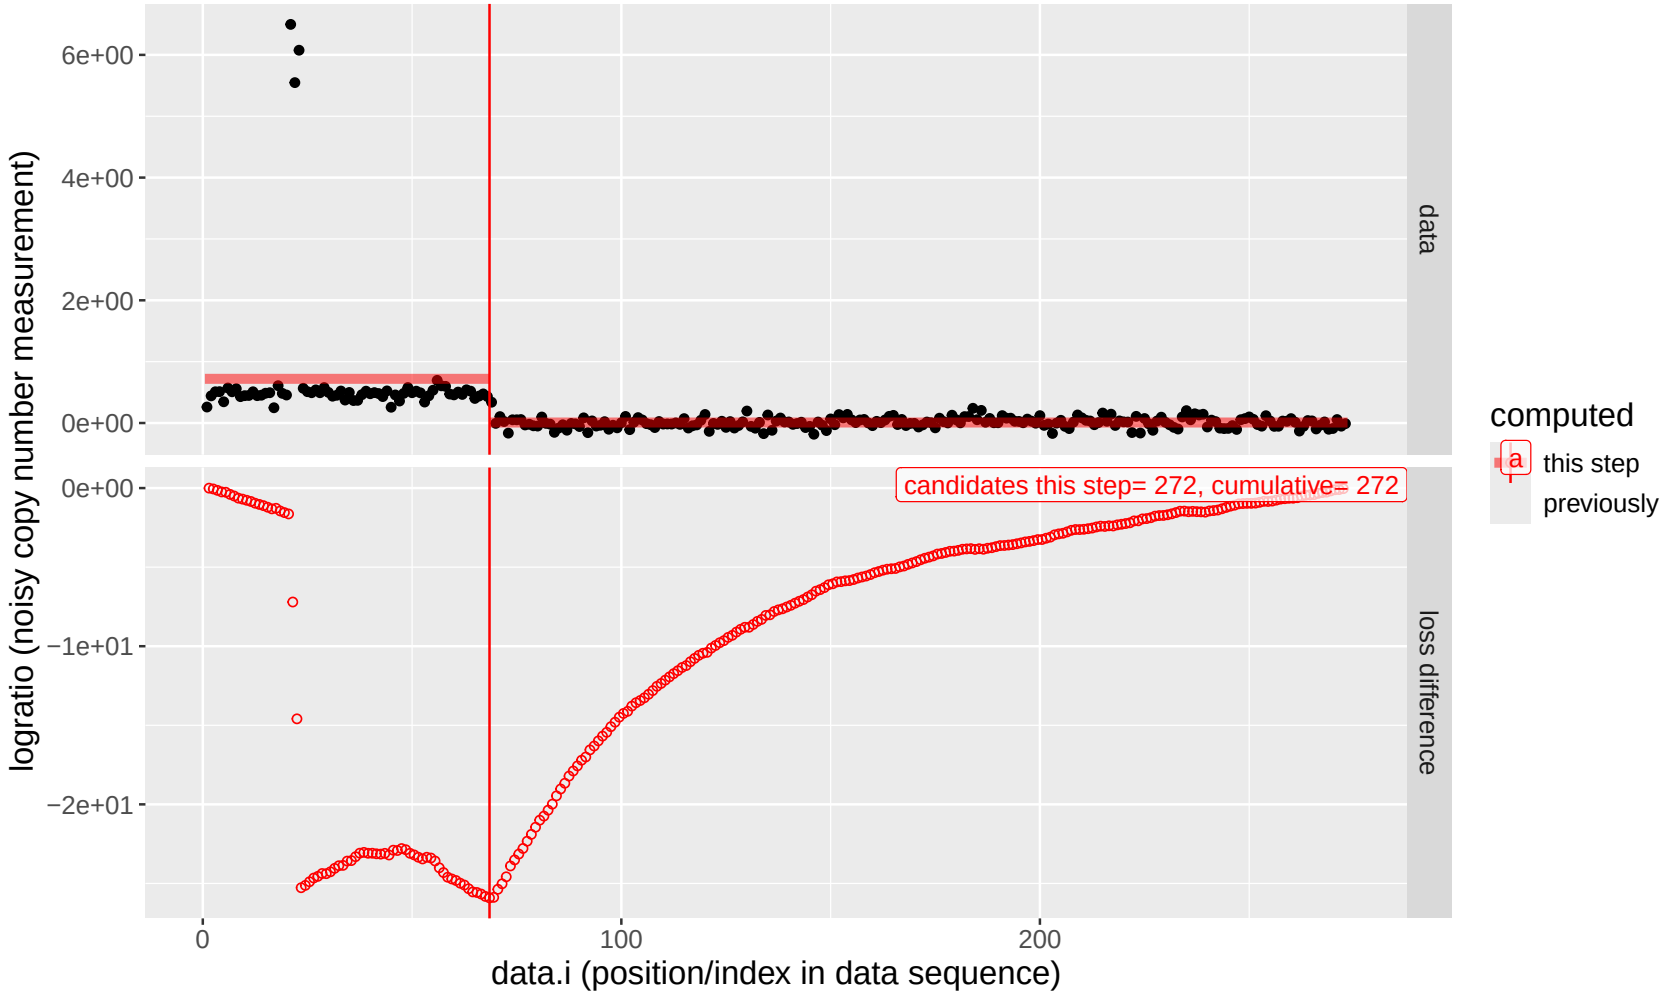

Visualization of computations at each iteration

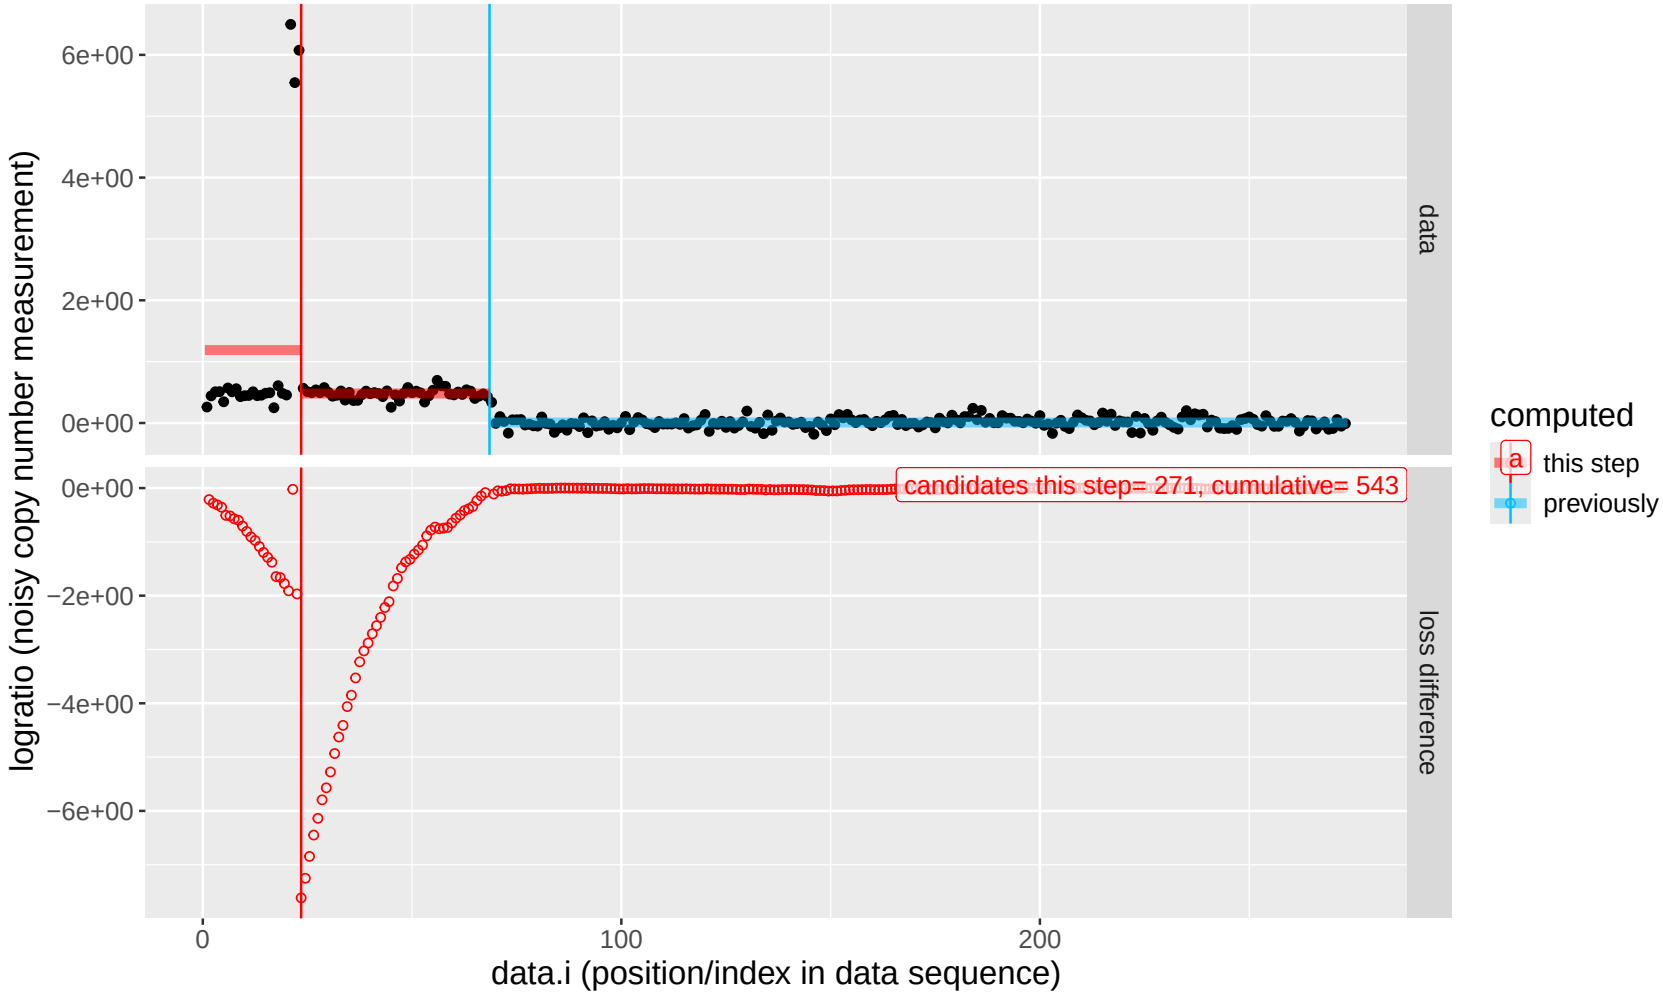

Visualization of computations at each iteration

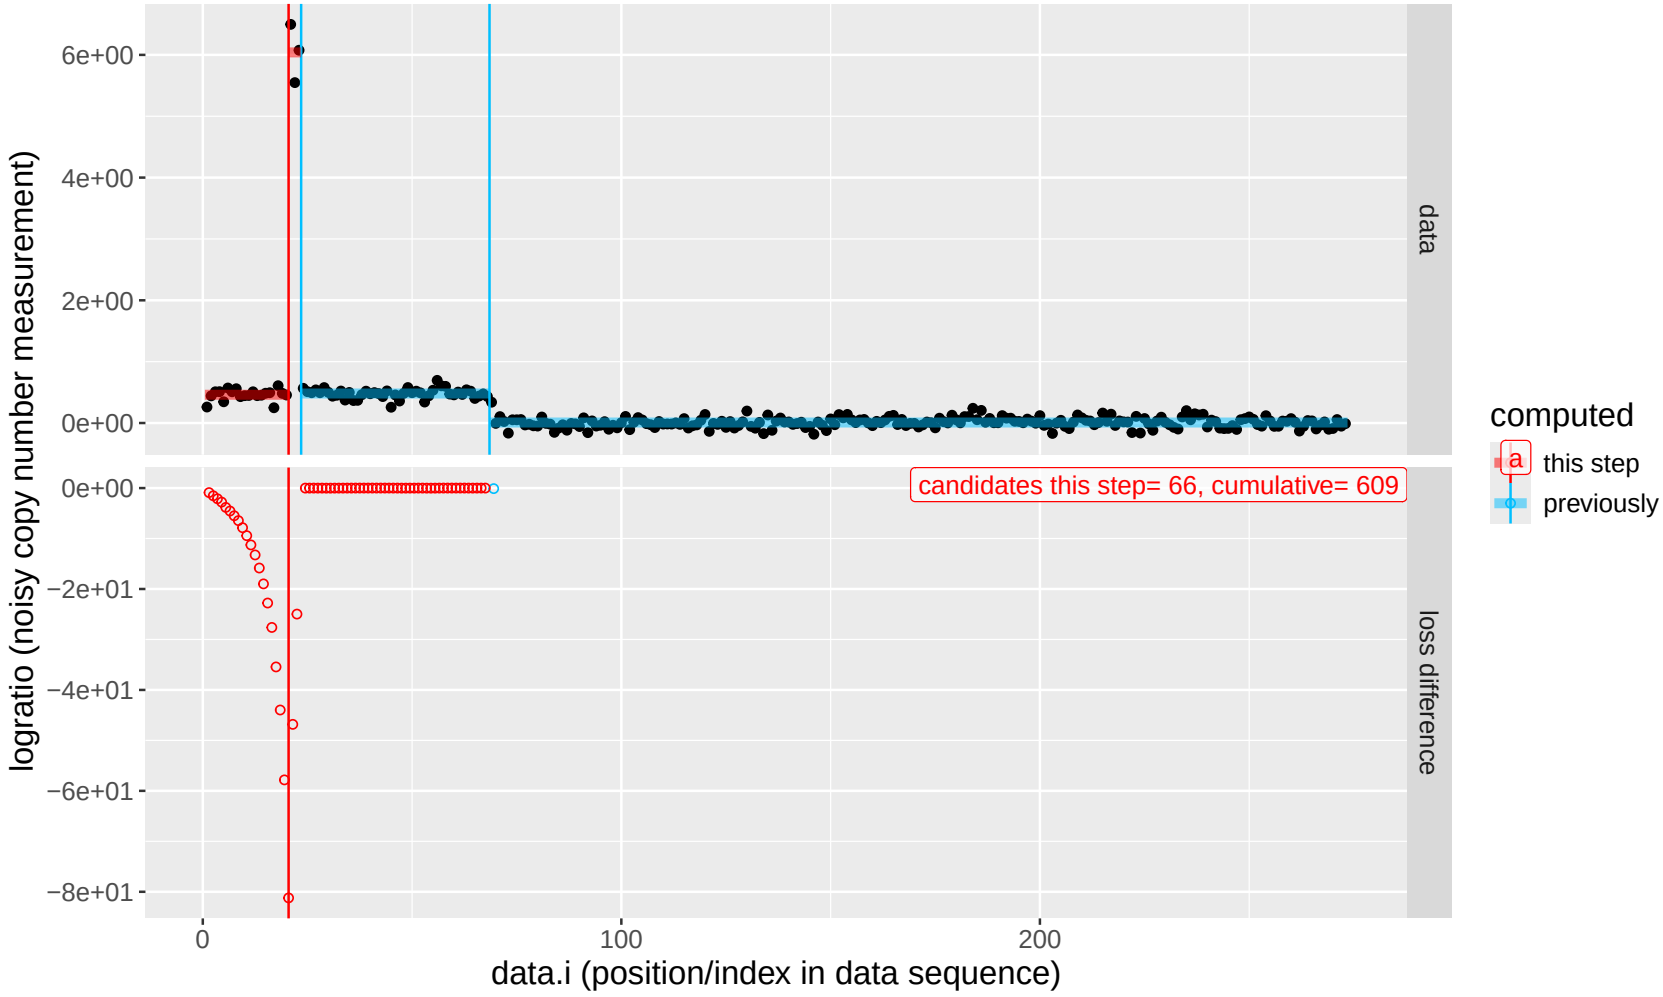

Visualization of computations at each iteration

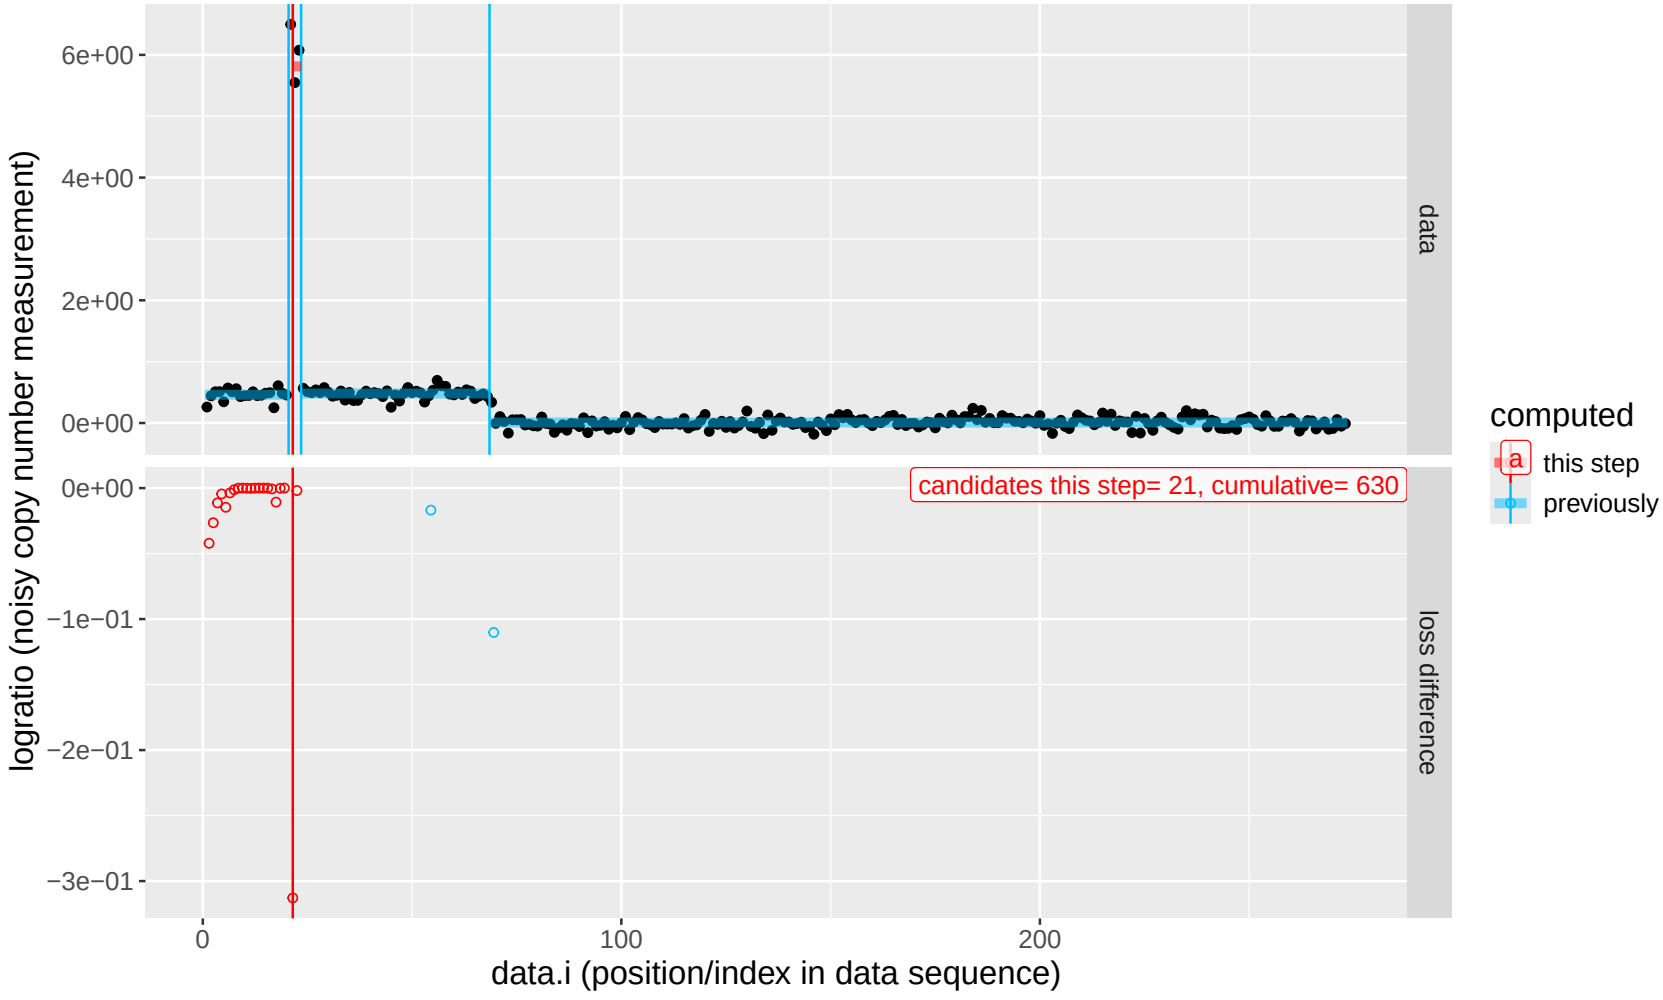

Visualization of computations at each iteration

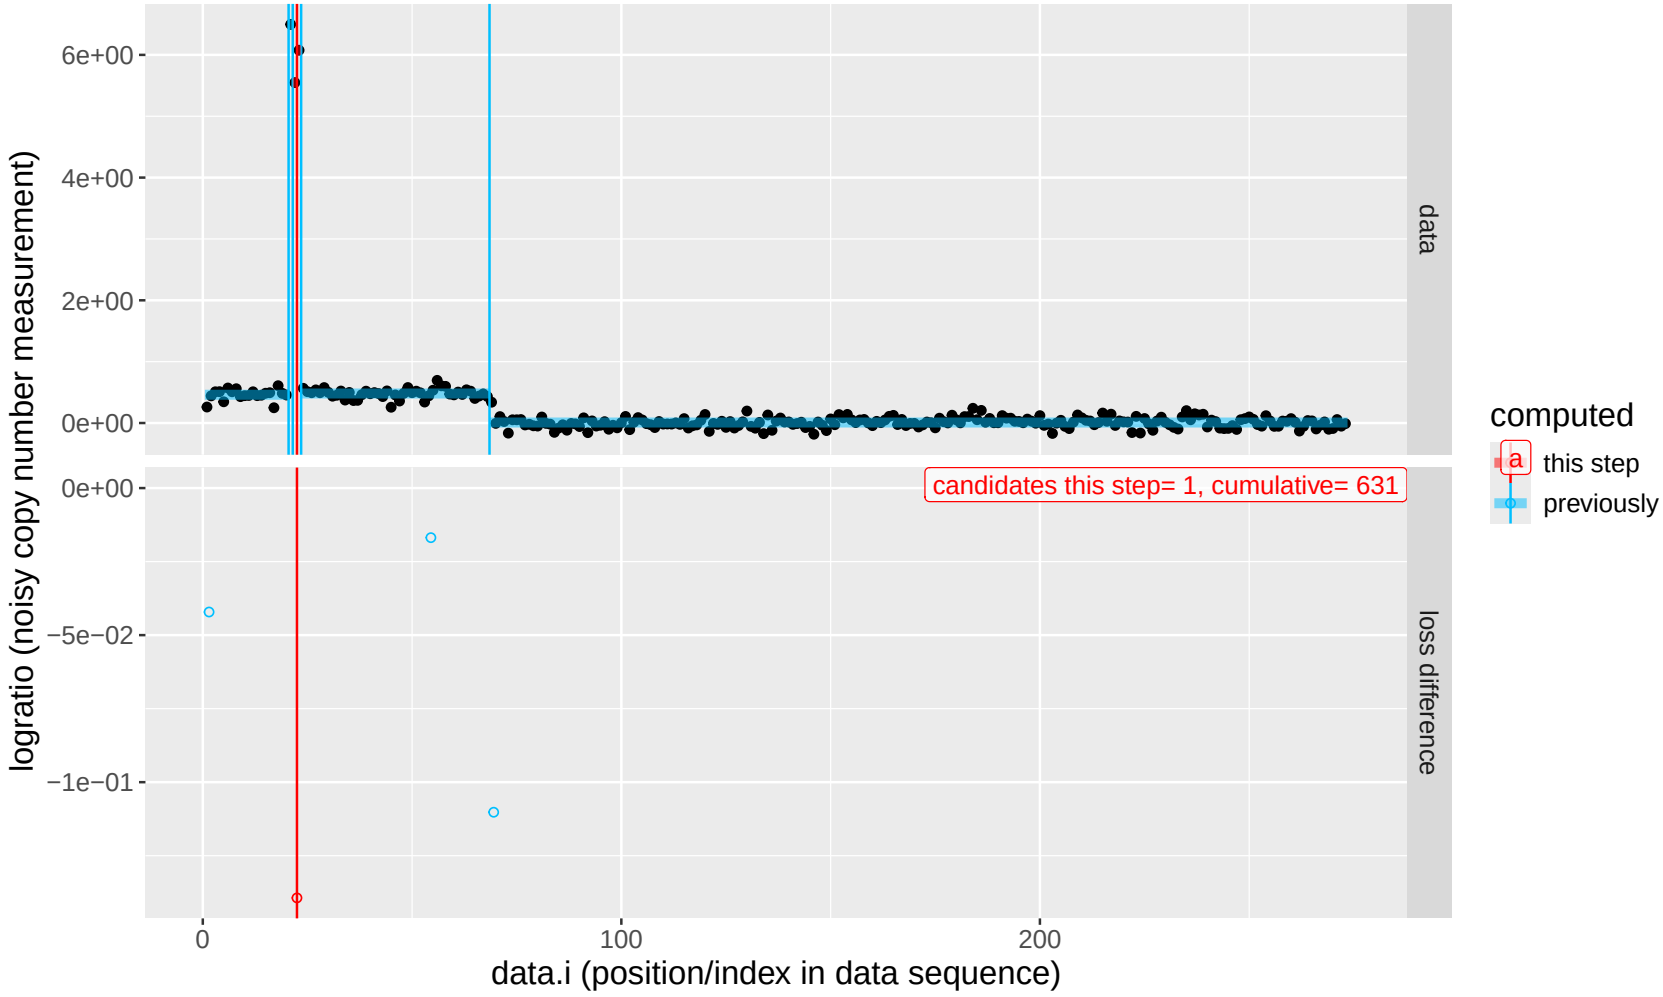

Visualization of computations at each iteration

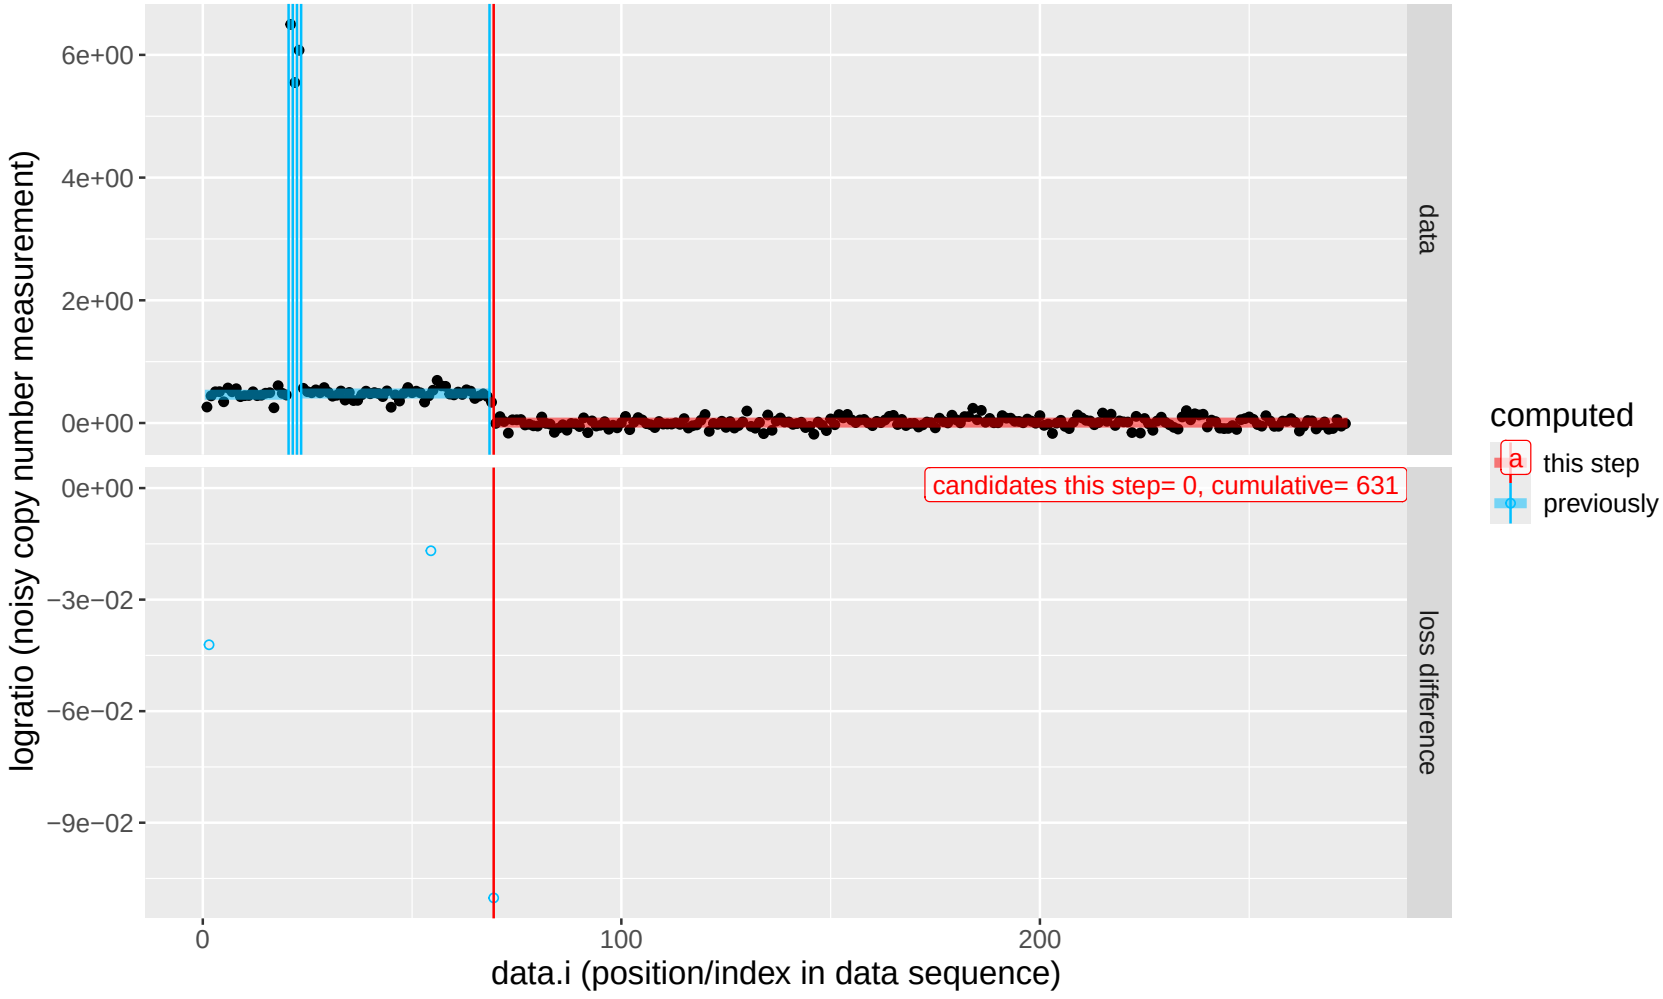

Visualization of computations at each iteration

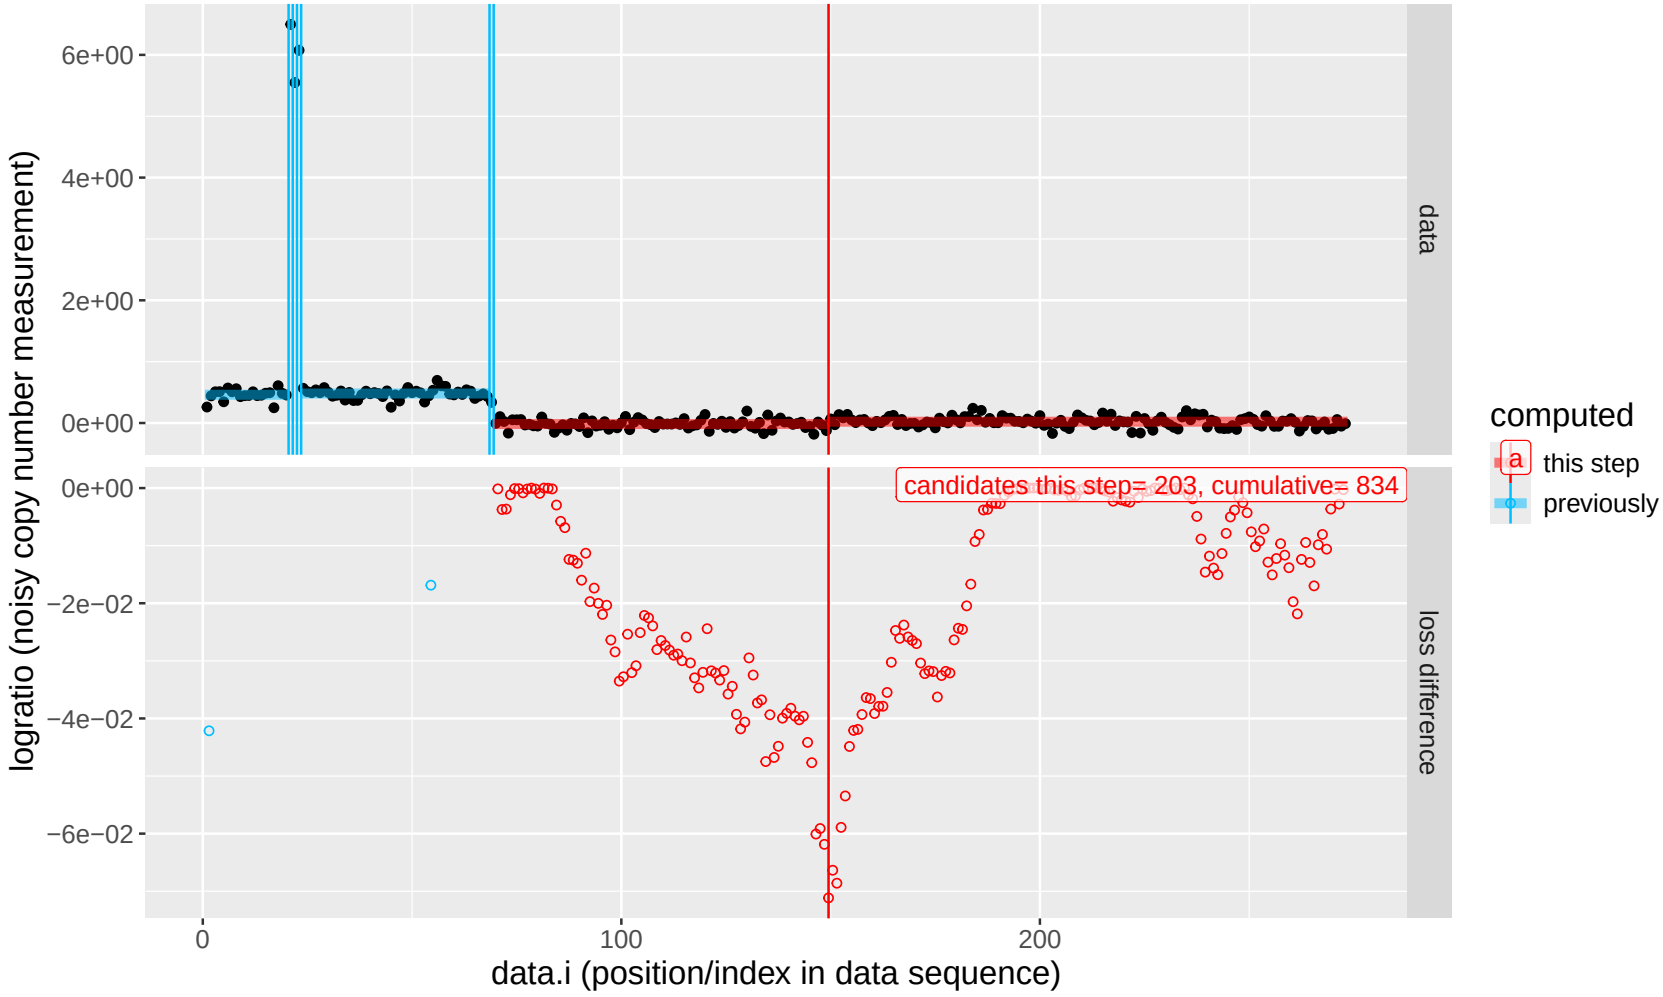

Visualization of computations at each iteration

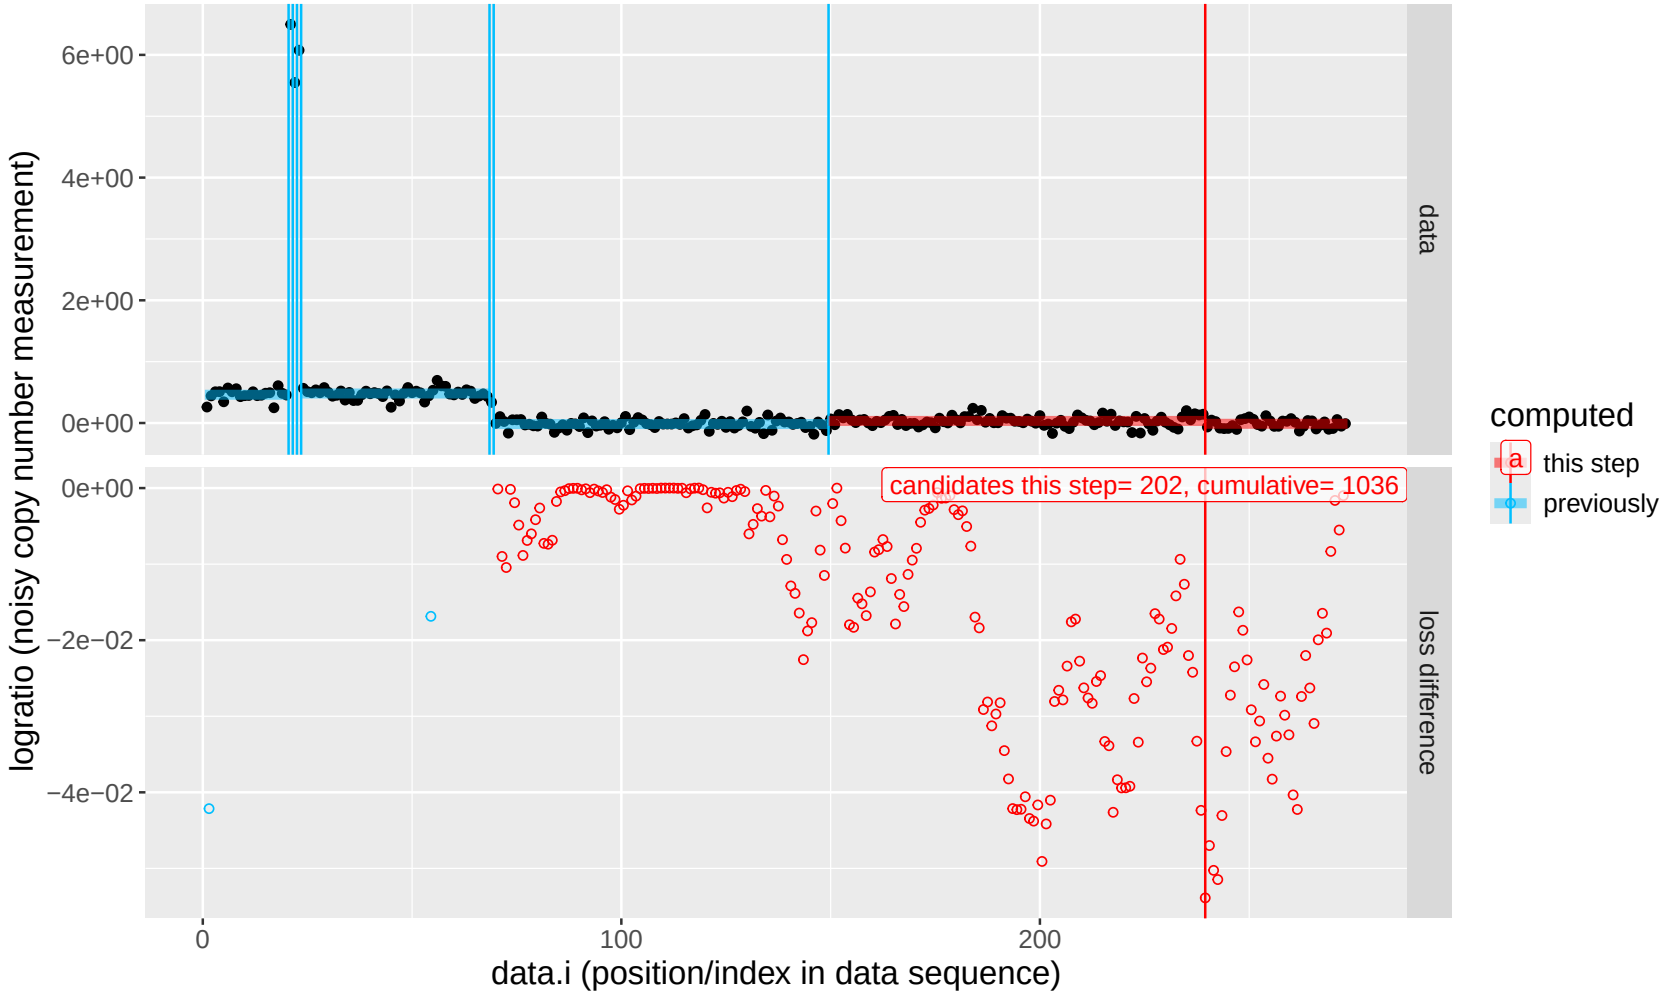

Visualization of computations at each iteration

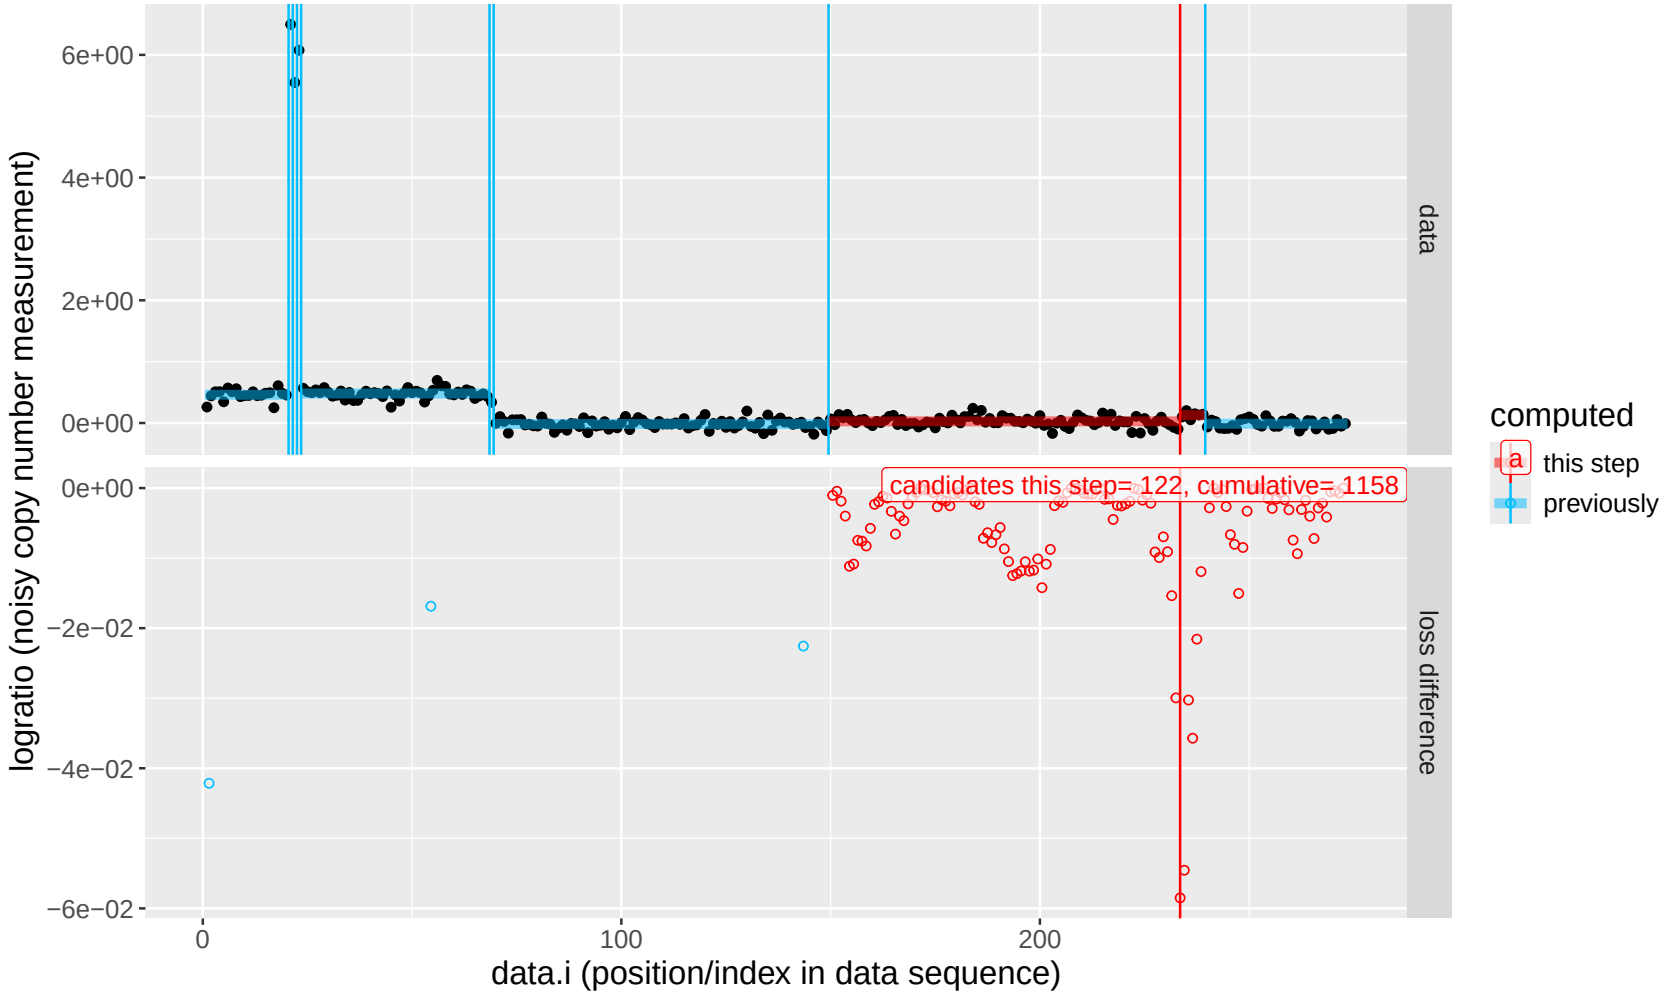

Visualization of computations at each iteration

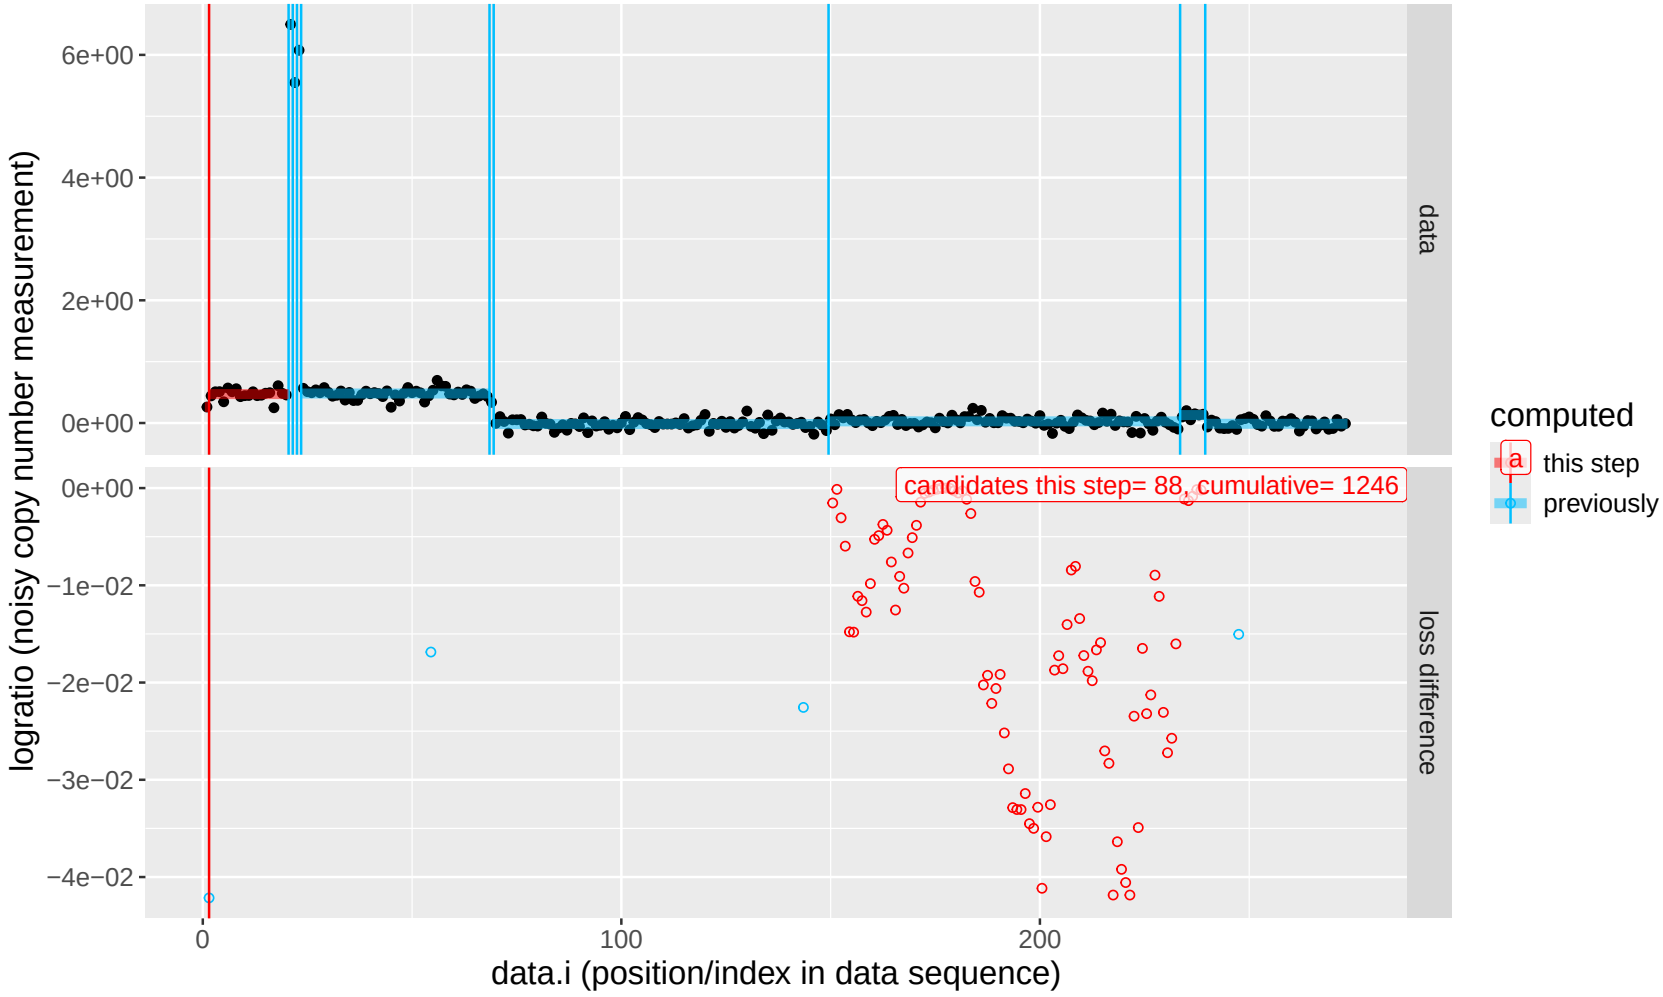

Visualization of computations at each iteration

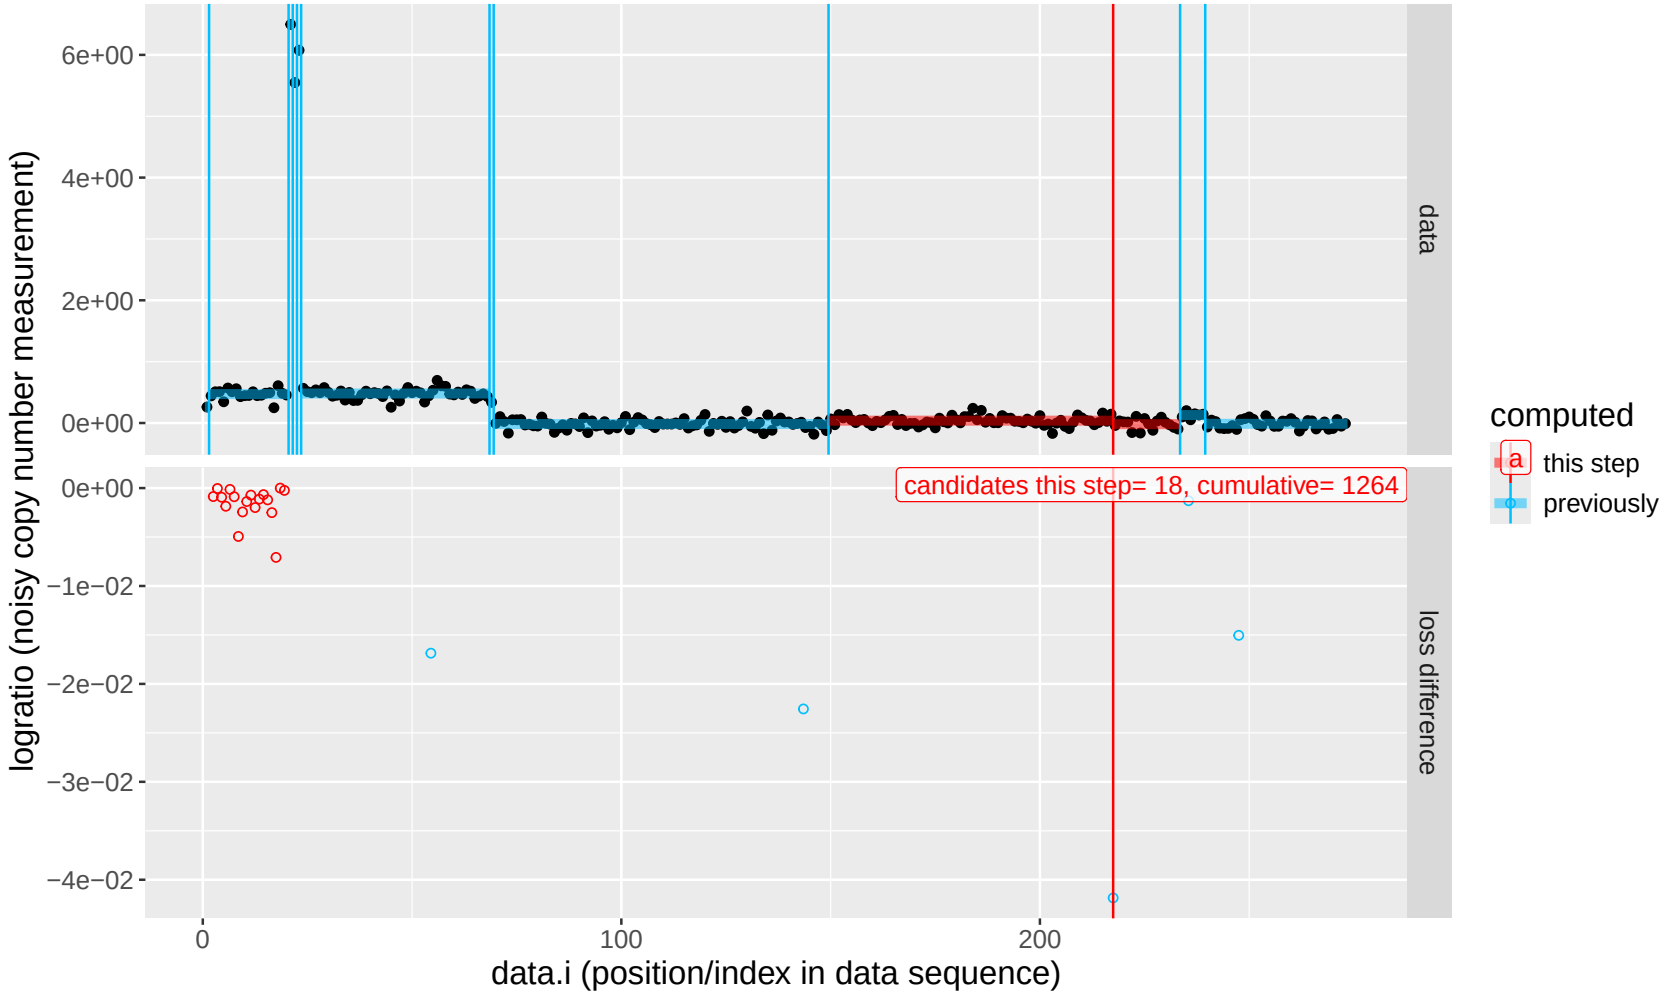

Complexity analysis

- Assume n data and K segments.
- Computing best loss decrease and split point for a segment with t data takes $O(t)$ time.
- Keep a list of segments which could be split, sorted by loss decrease values.
- Best case is when segments get cut in half each time, $O(n \log K)$ time. (minimize number of possible splits for which we have to recompute loss)
- Worst case is when splits are very unequal $(1, t-1)$, $O(nK)$ time. (maximize number of possible splits for which we have to recompute loss)

Detailed complexity analysis

- Let $n = 2^J$ for some $J \in \{1, 2, \dots\}$, for example $J = 6 \Rightarrow n = 64$.
- For any $j \in \{1, \dots, J+1\}$ if we do $K = 2^{j-1}$ iterations then how many split cost values to compute?
- Best case: $nj - 2^j + 1 = n(1 + \log_2 K) - K/2 + 1 \Rightarrow O(n \log K)$.
- Worst case: $nK - K(1+K)/2 \Rightarrow O(nK)$.

j	K	best	total	worst	total
1	1	$n-1=63$	$n-1=63$	$n-1=63$	$n-1=63$
2	2	$n-2=62$	$2n-3=125$	$n-2=62$	$2n-3=125$
	3	$n/2-2=30$		$n-3=61$	$3n-6=186$
3	4	$n/2-2=30$	$3n-7=185$	$n-4=60$	$4n-10=246$
⋮	⋮	⋮	⋮	⋮	⋮
4	8	$n/4-2=14$	$4n-15=241$	$n-8=56$	$8n-36=476$
⋮	⋮	⋮	⋮	⋮	⋮
7	64	$n/32-2=0$	$7n-128=321$	$n-64=0$	$64n-2080=2016$

Real data time complexity analysis

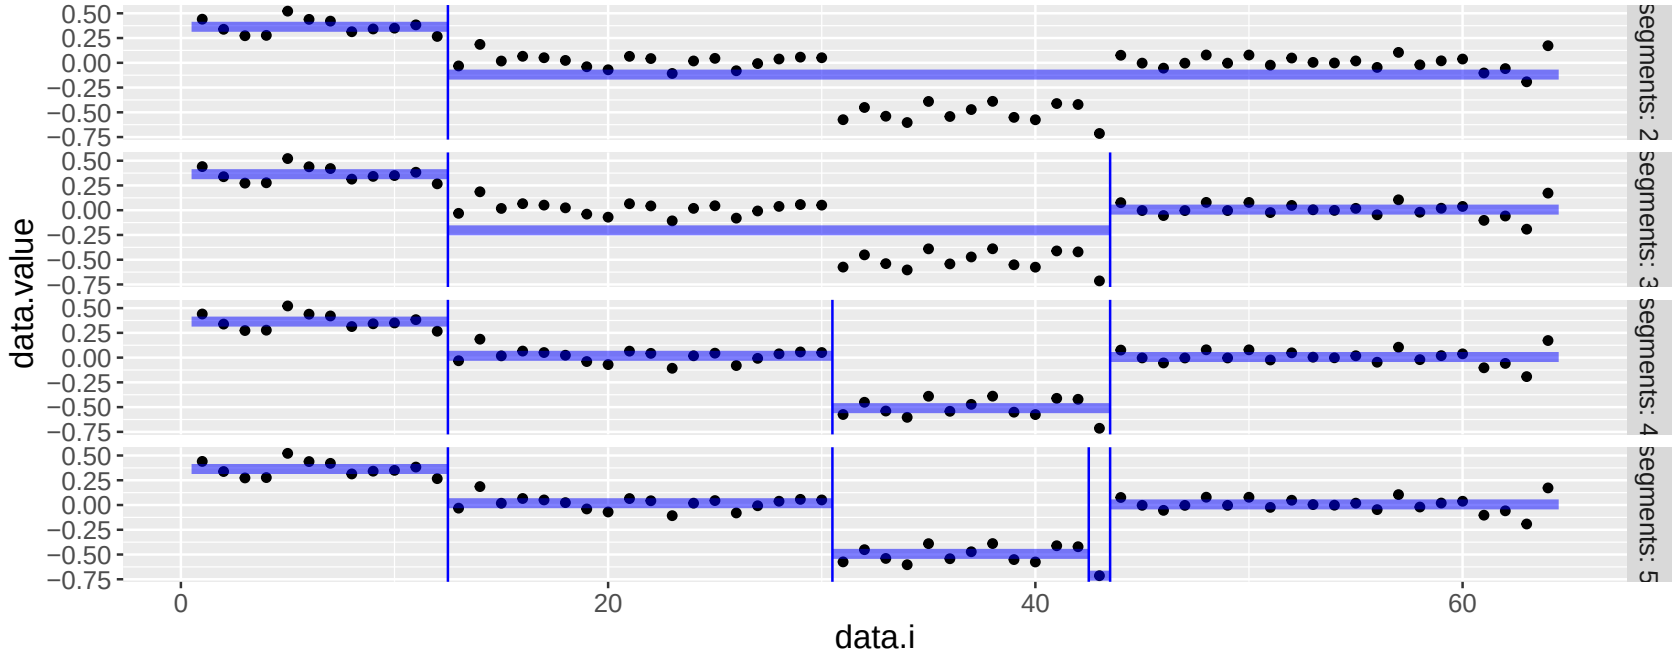

Number of splits for which loss is computed, empirical data type=real

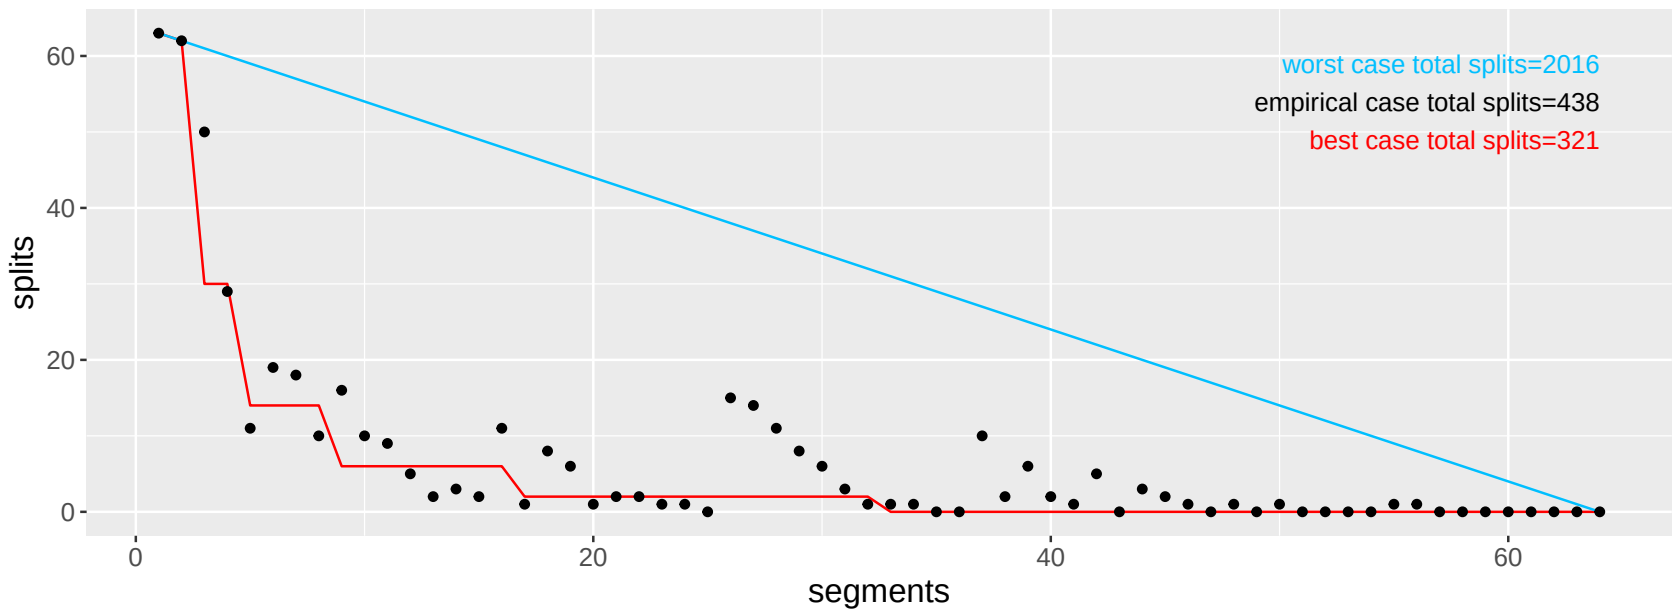

Synthetic data time complexity analysis

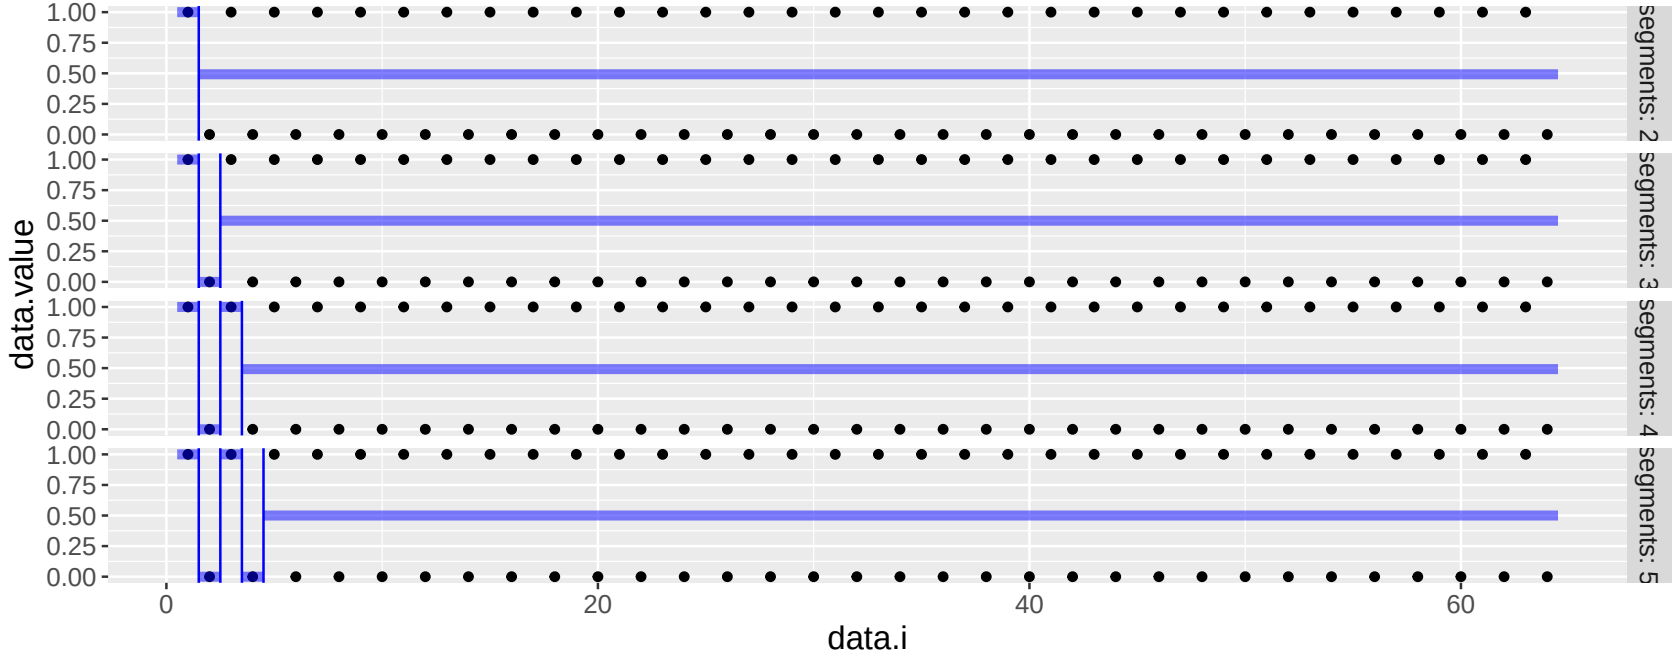

Number of splits for which loss is computed, empirical data type=up.and.down

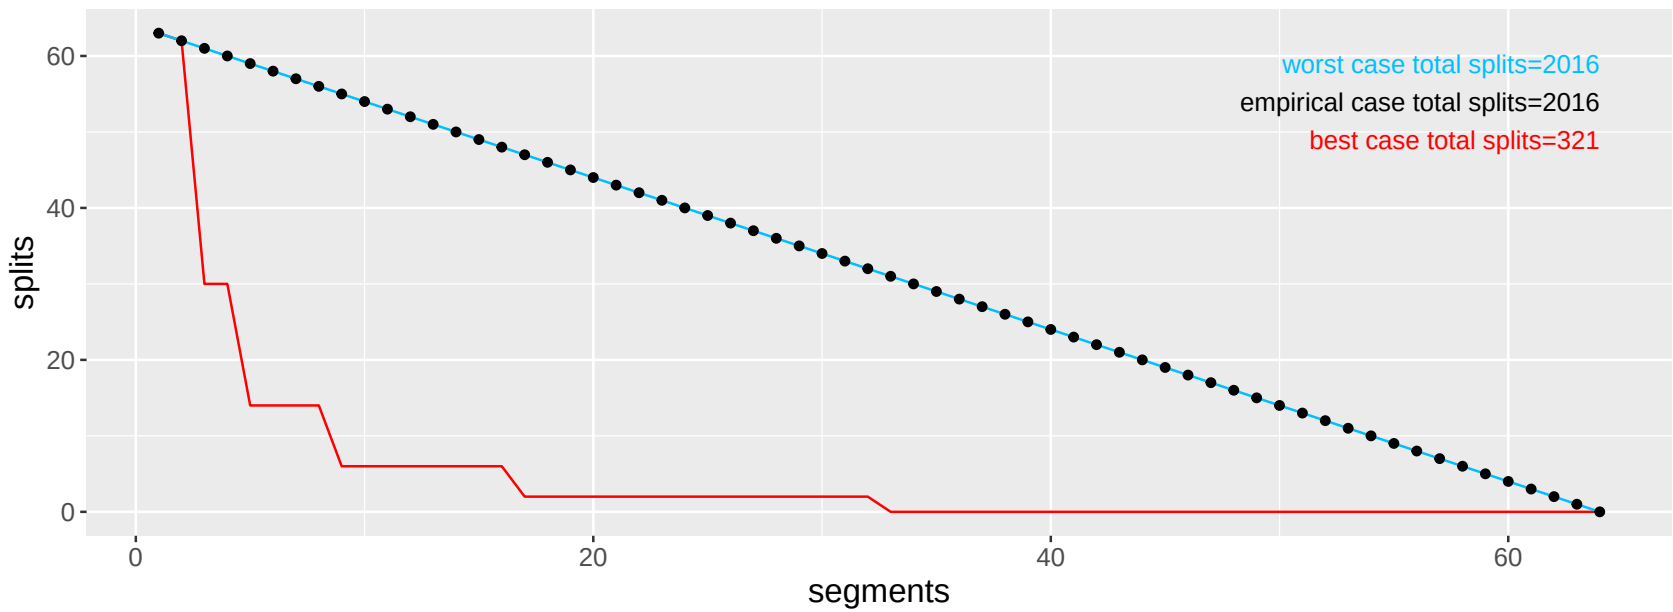

Synthetic data time complexity analysis

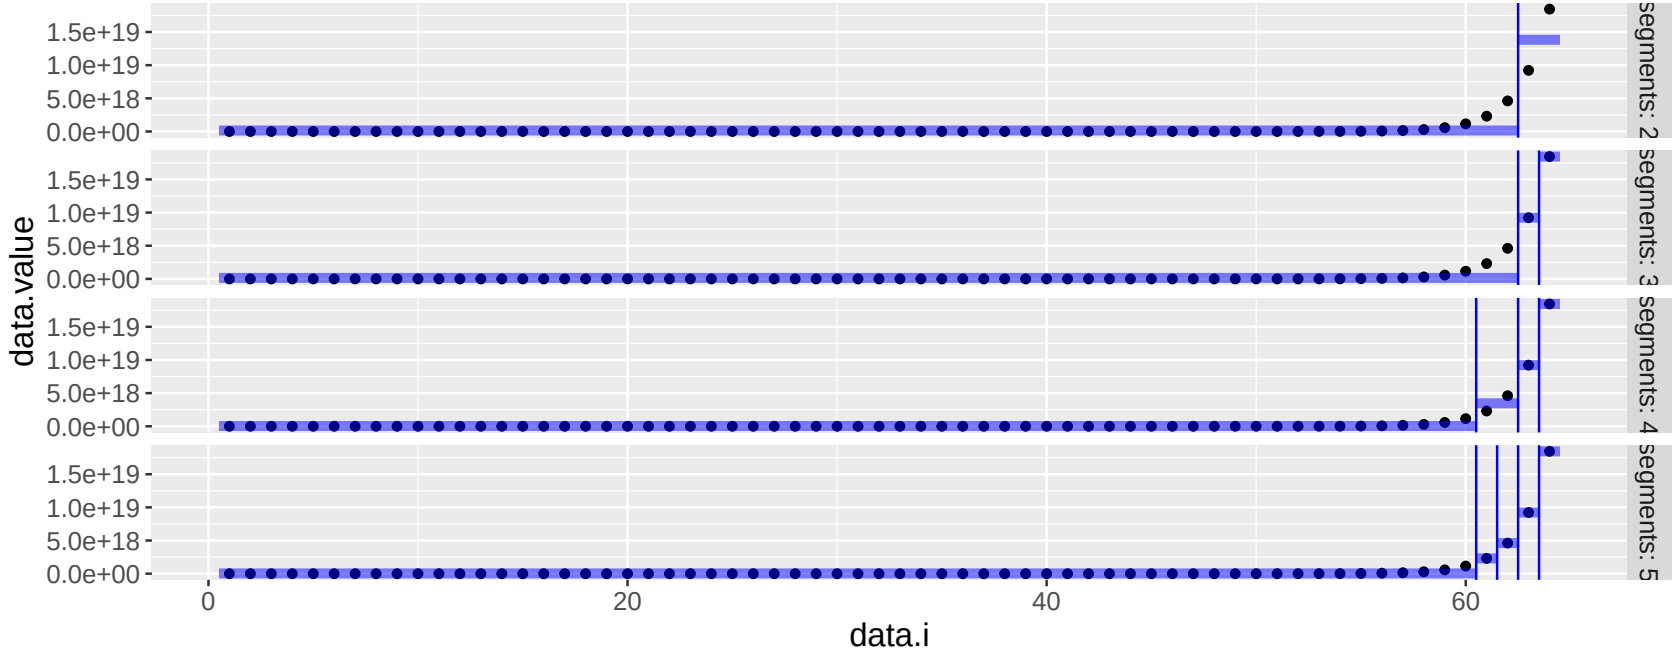

Number of splits for which loss is computed, empirical data type=exponential

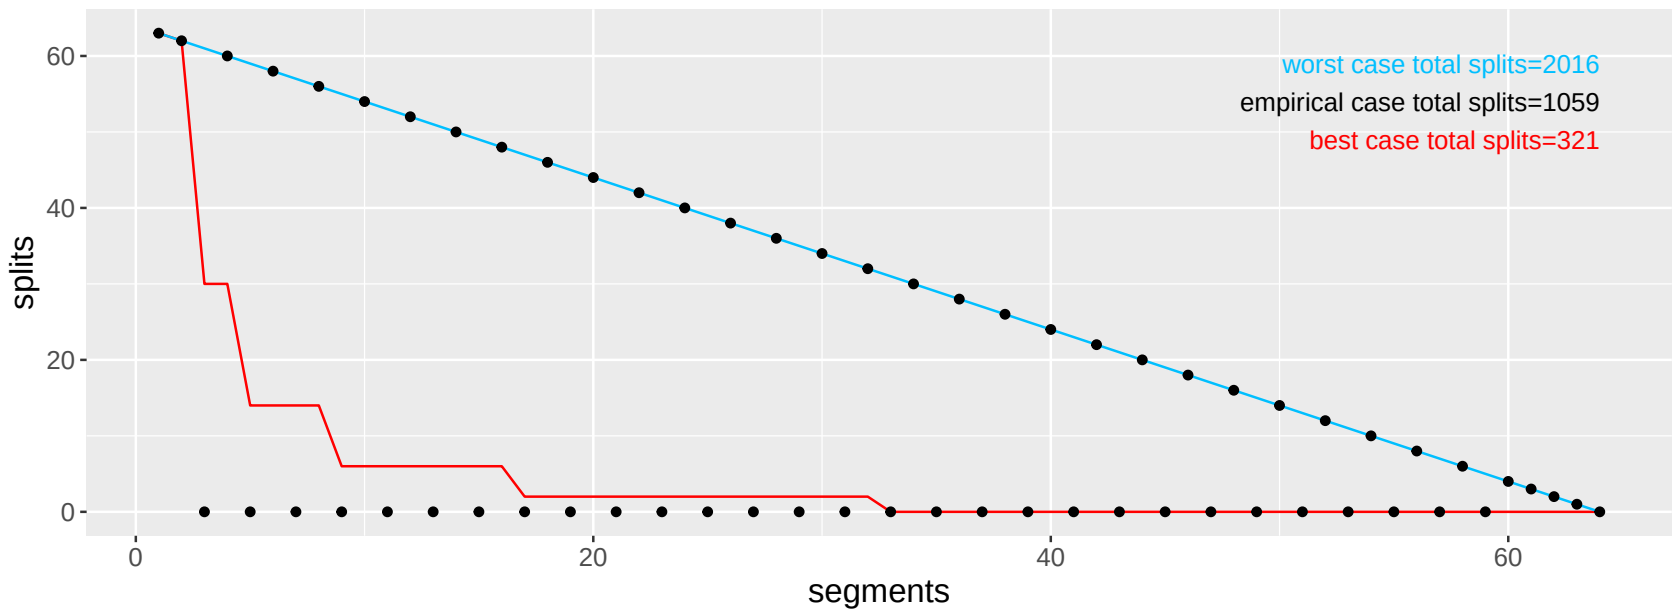

Synthetic data time complexity analysis

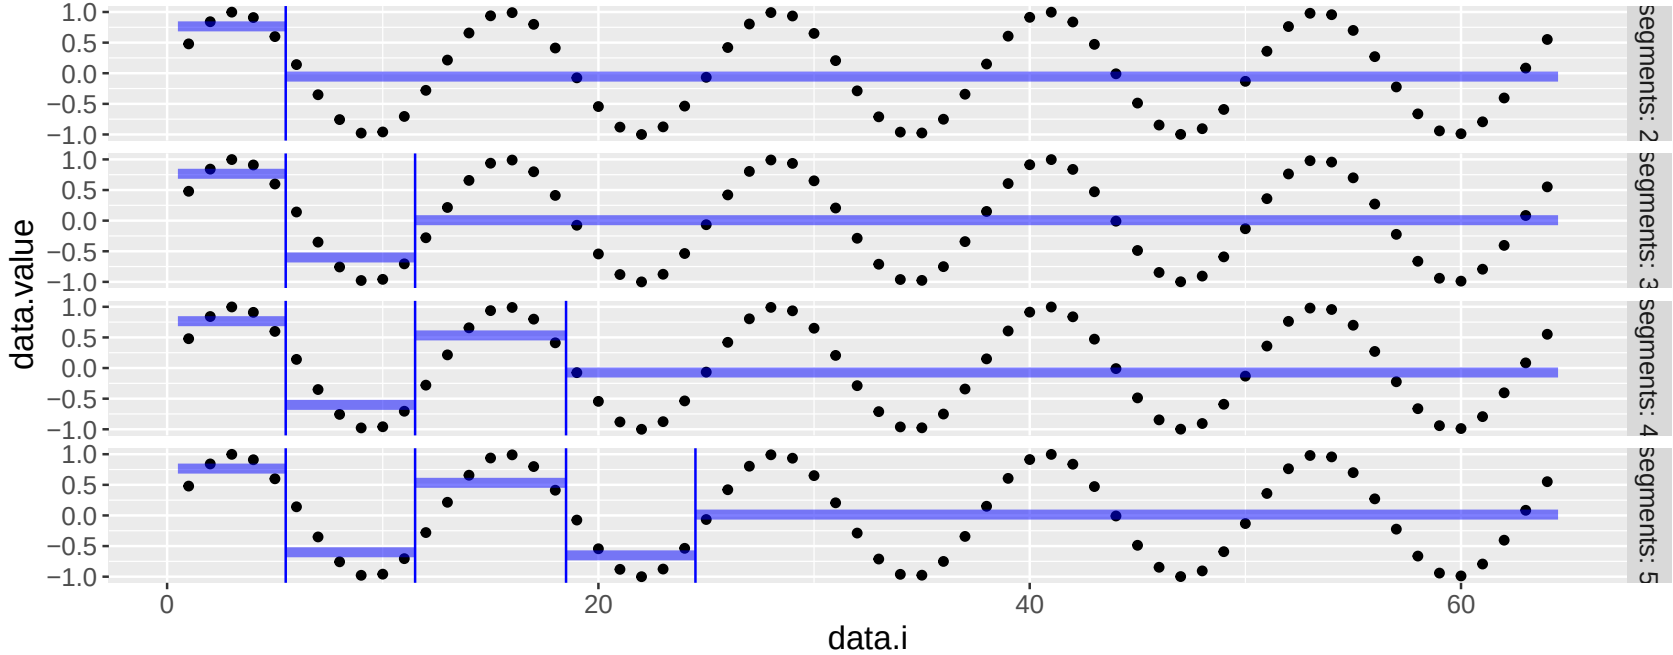

Number of splits for which loss is computed, empirical data type=sin

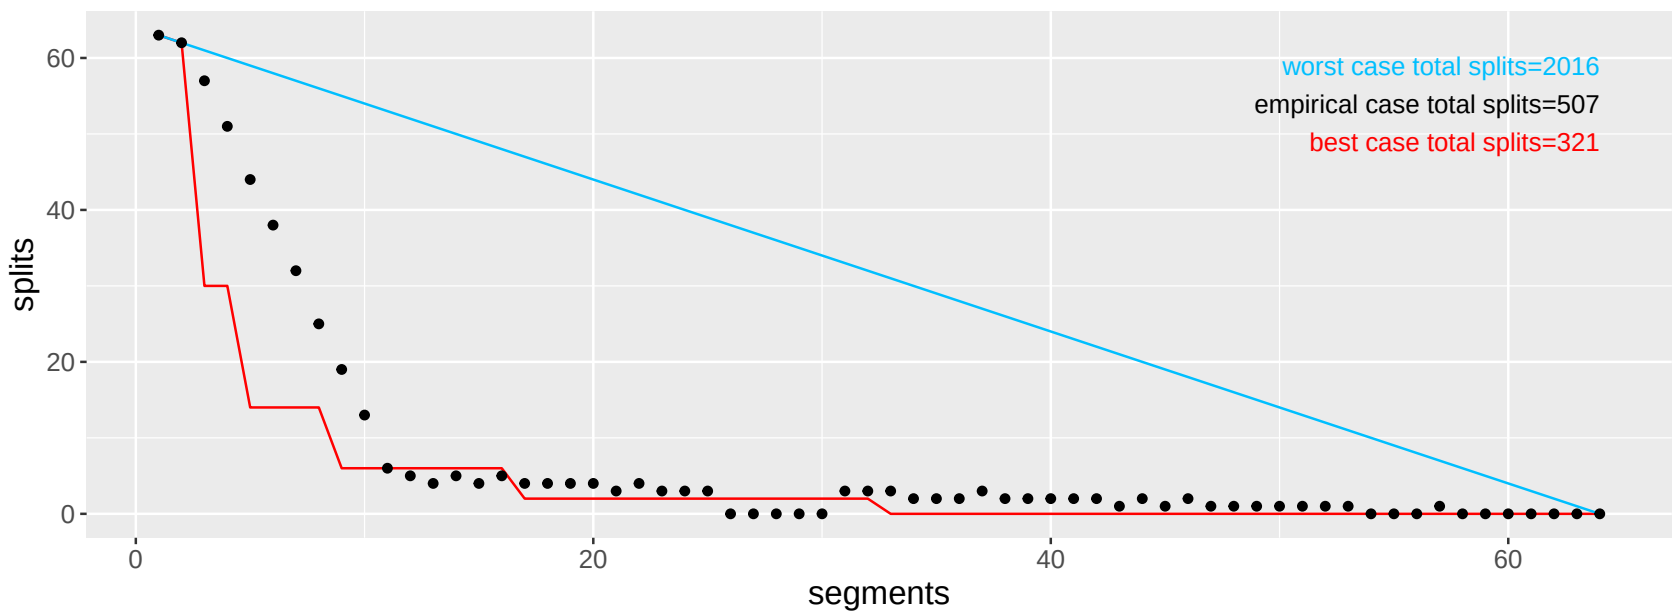

Synthetic data time complexity analysis

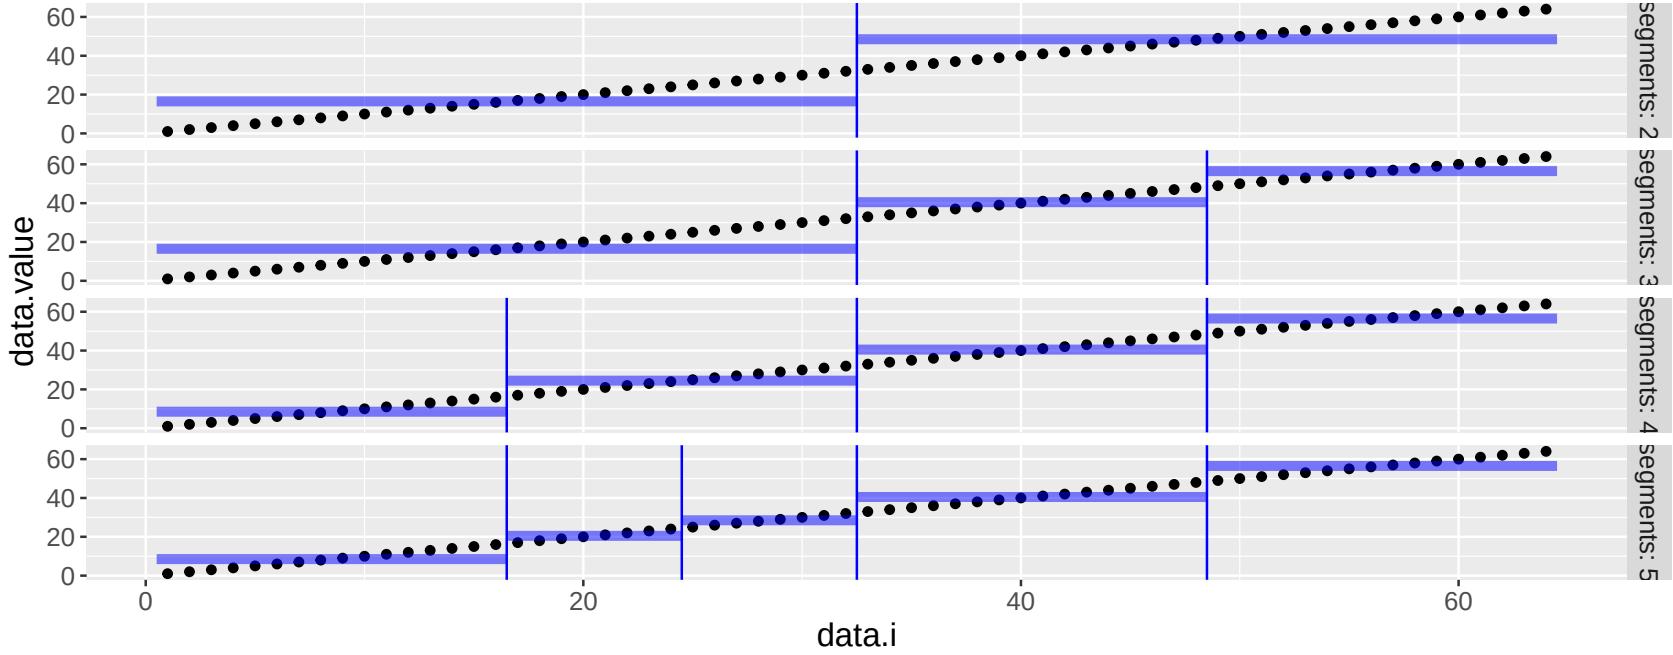

Number of splits for which loss is computed, empirical data type=linear

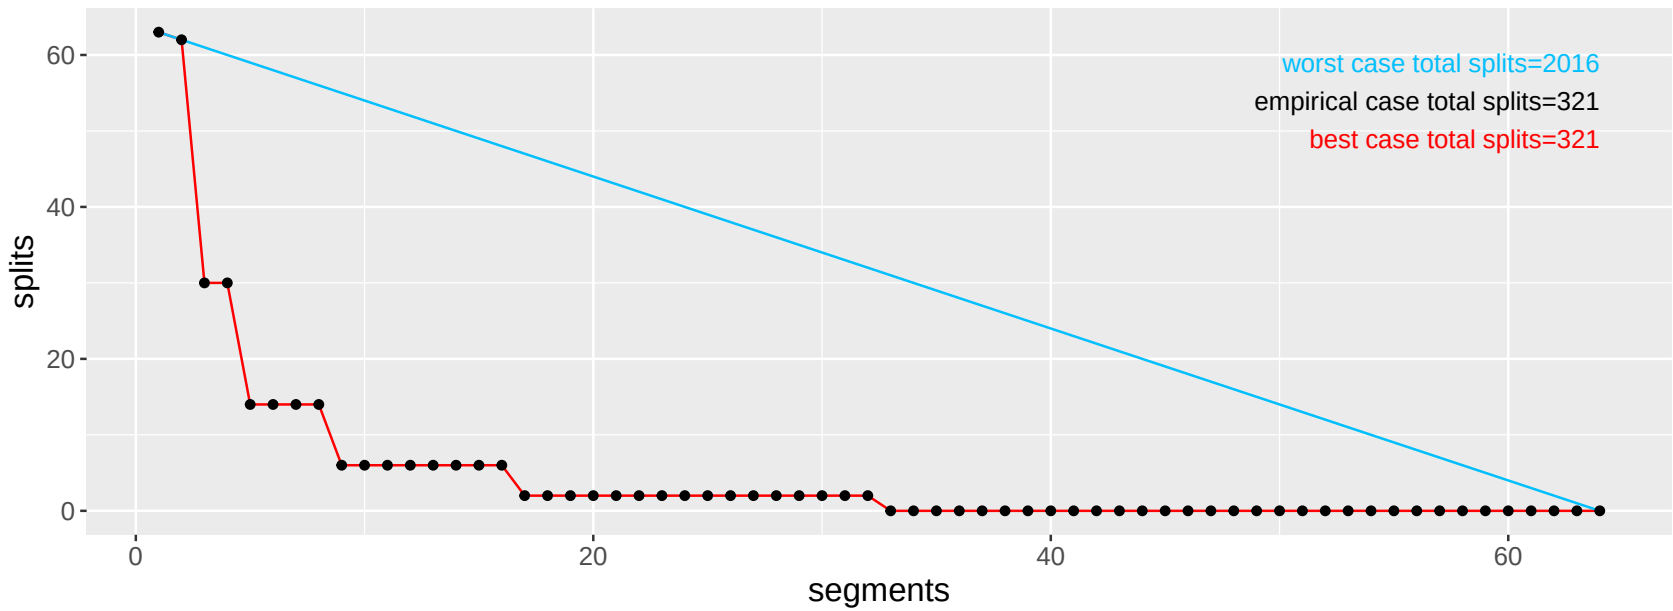

Synthetic data time complexity, vary max segments

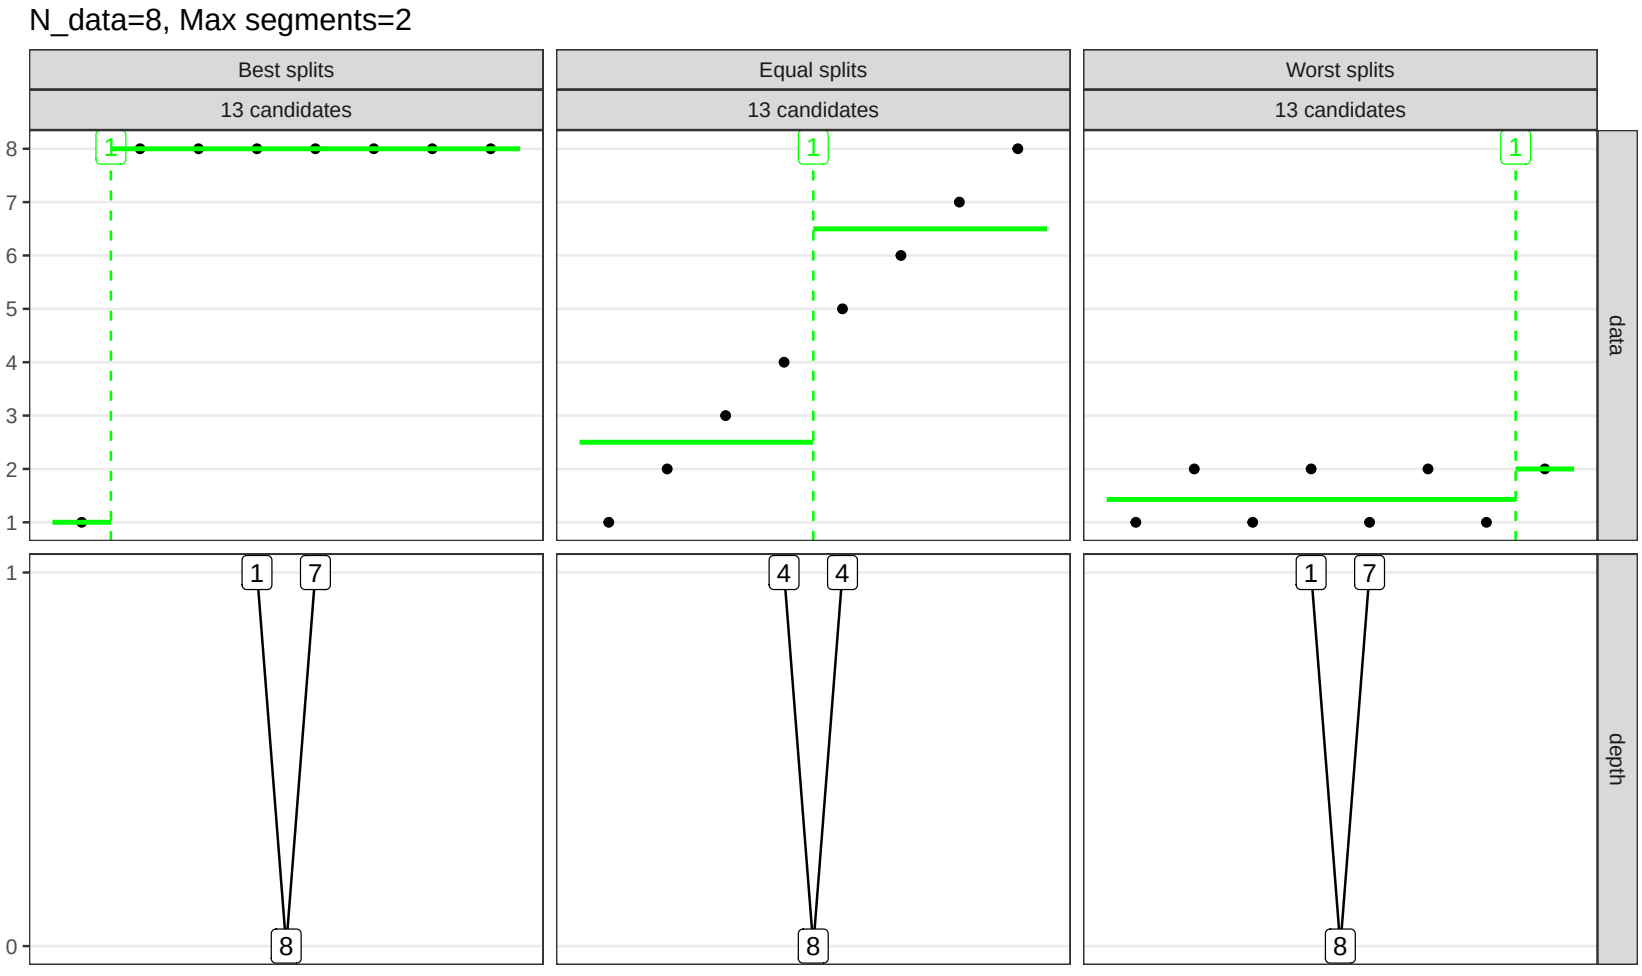

Synthetic data time complexity, vary max segments

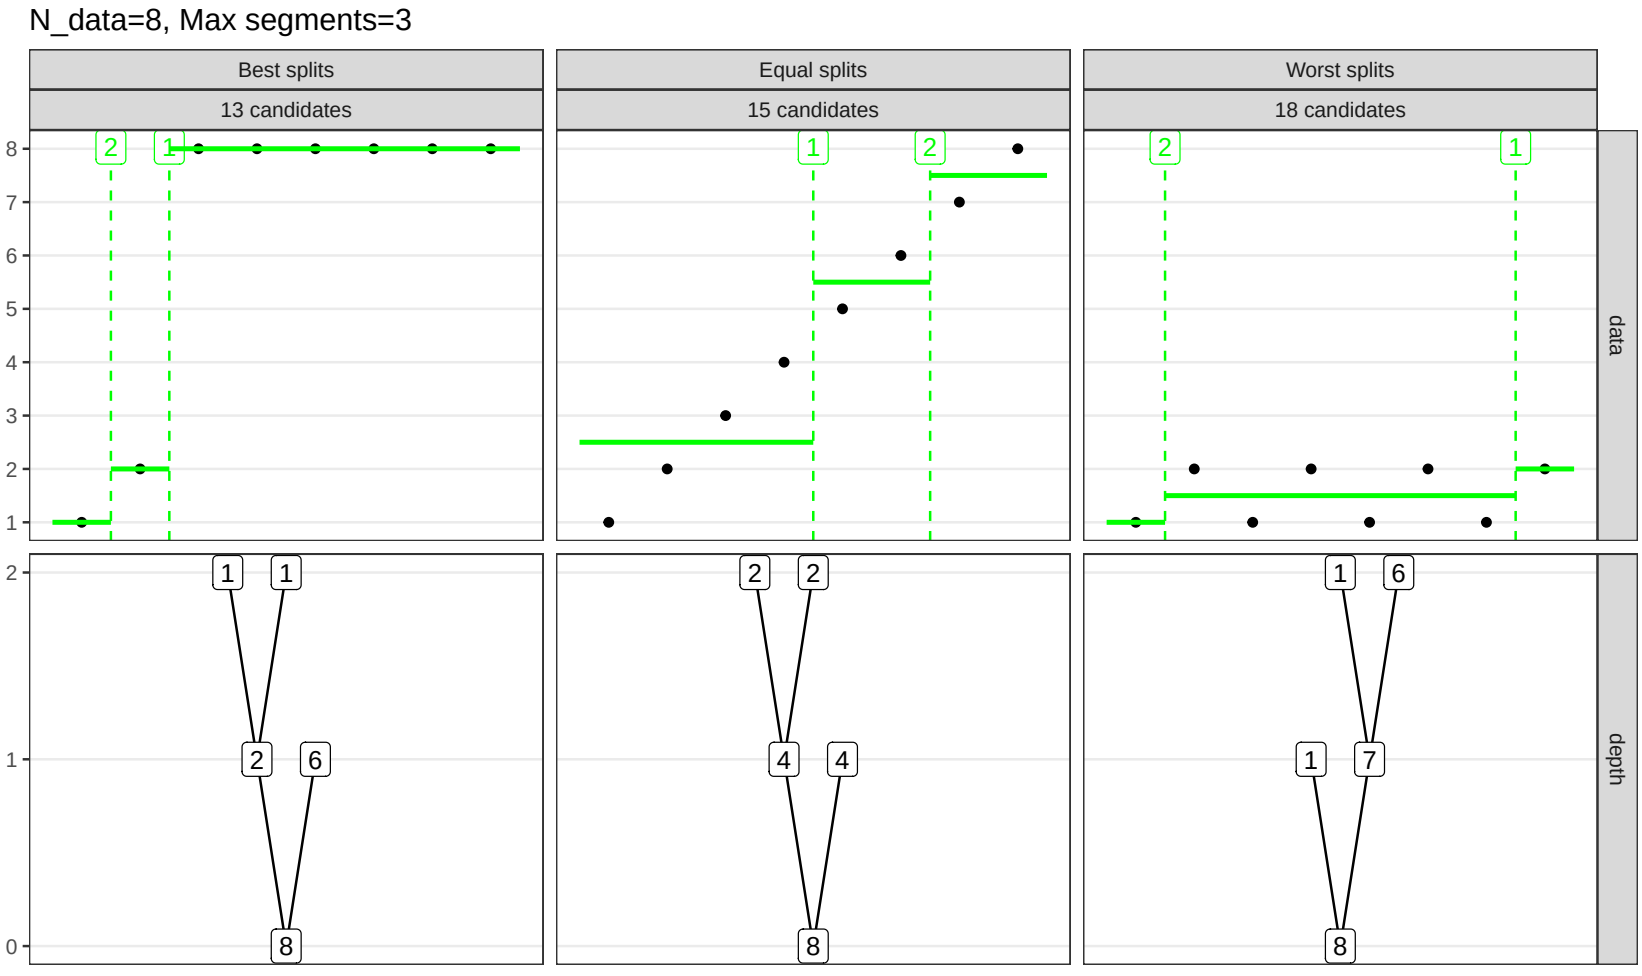

Synthetic data time complexity, vary max segments

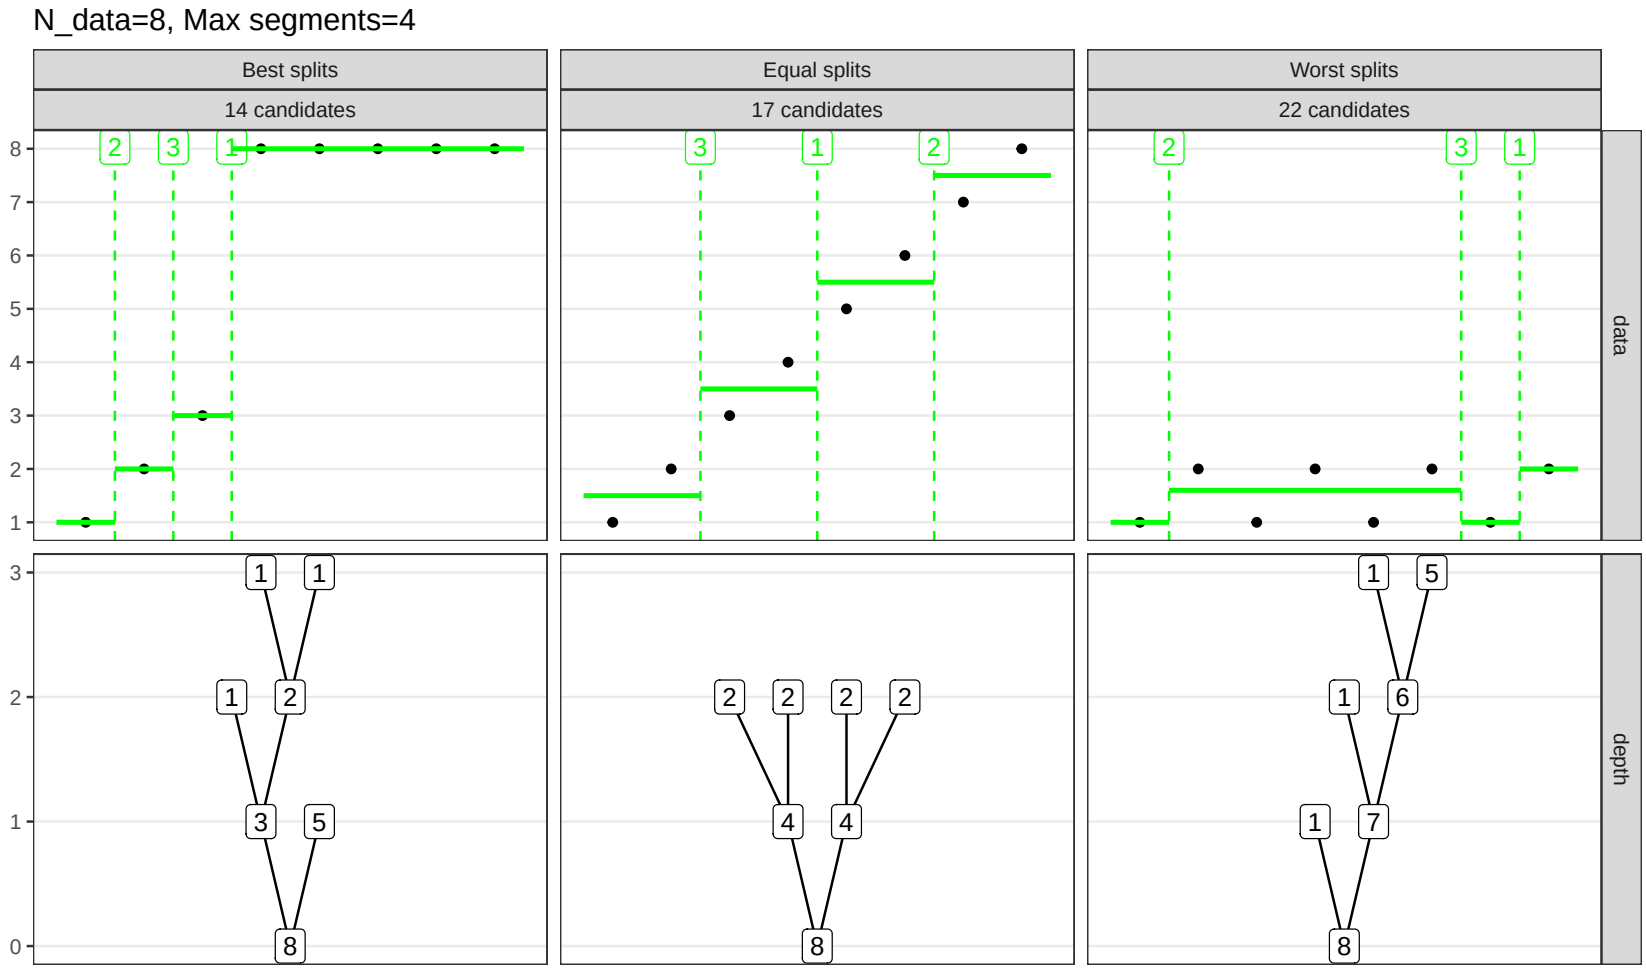

Synthetic data time complexity, vary max segments

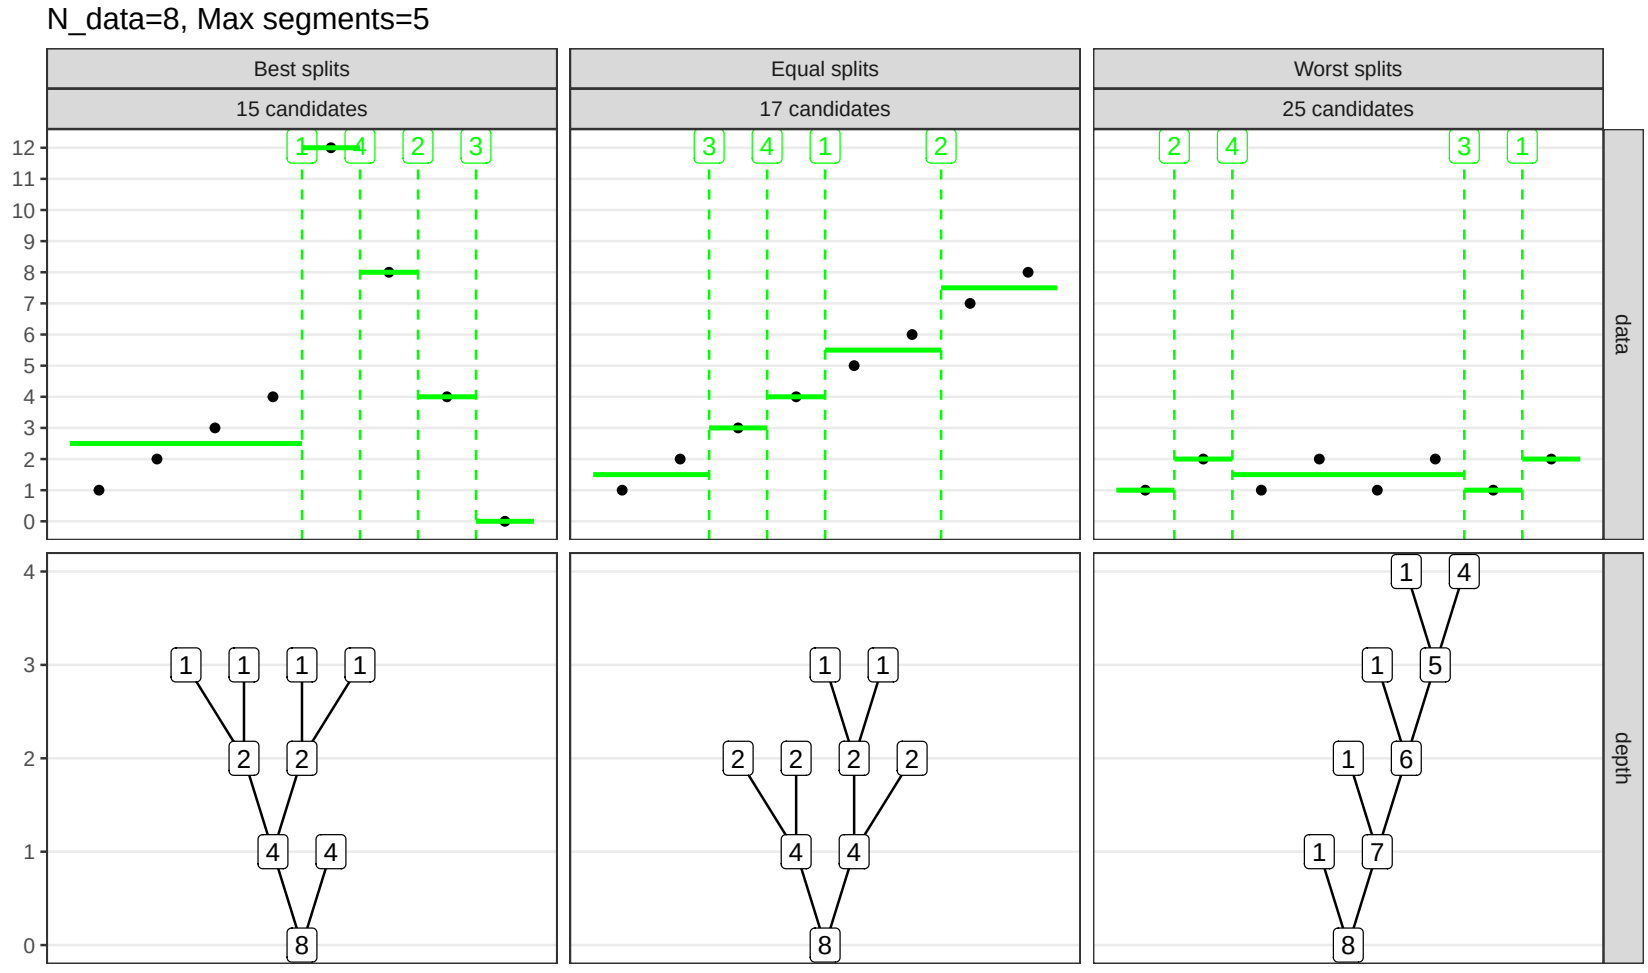

Synthetic data time complexity, vary max segments

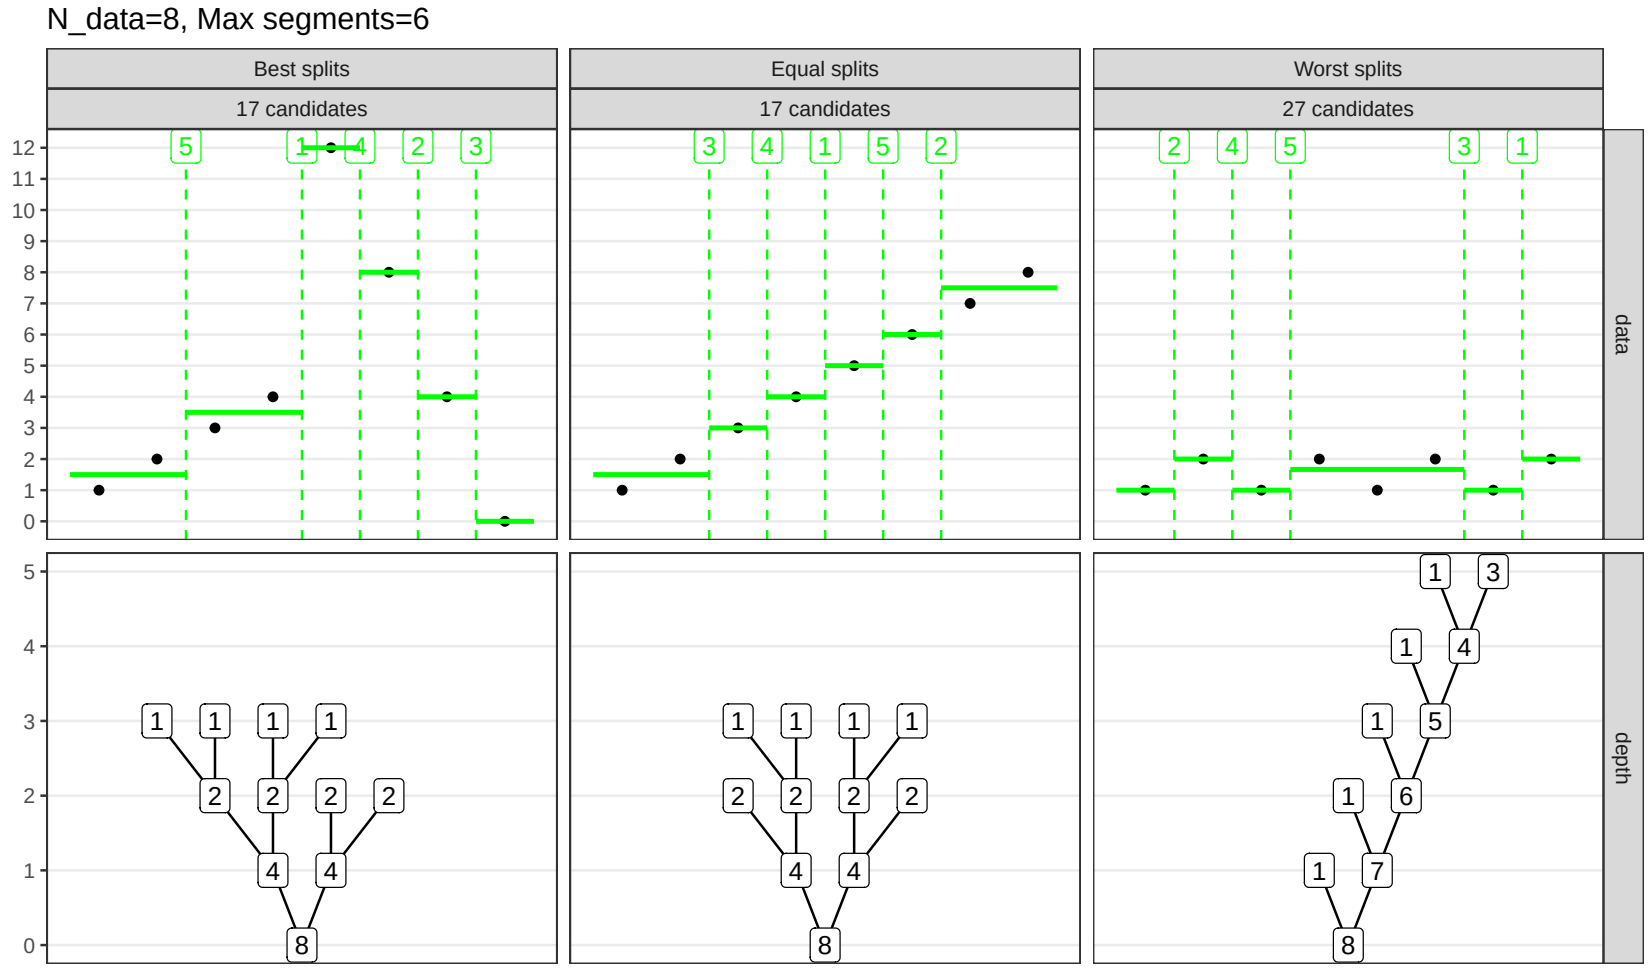

Synthetic data time complexity, vary max segments

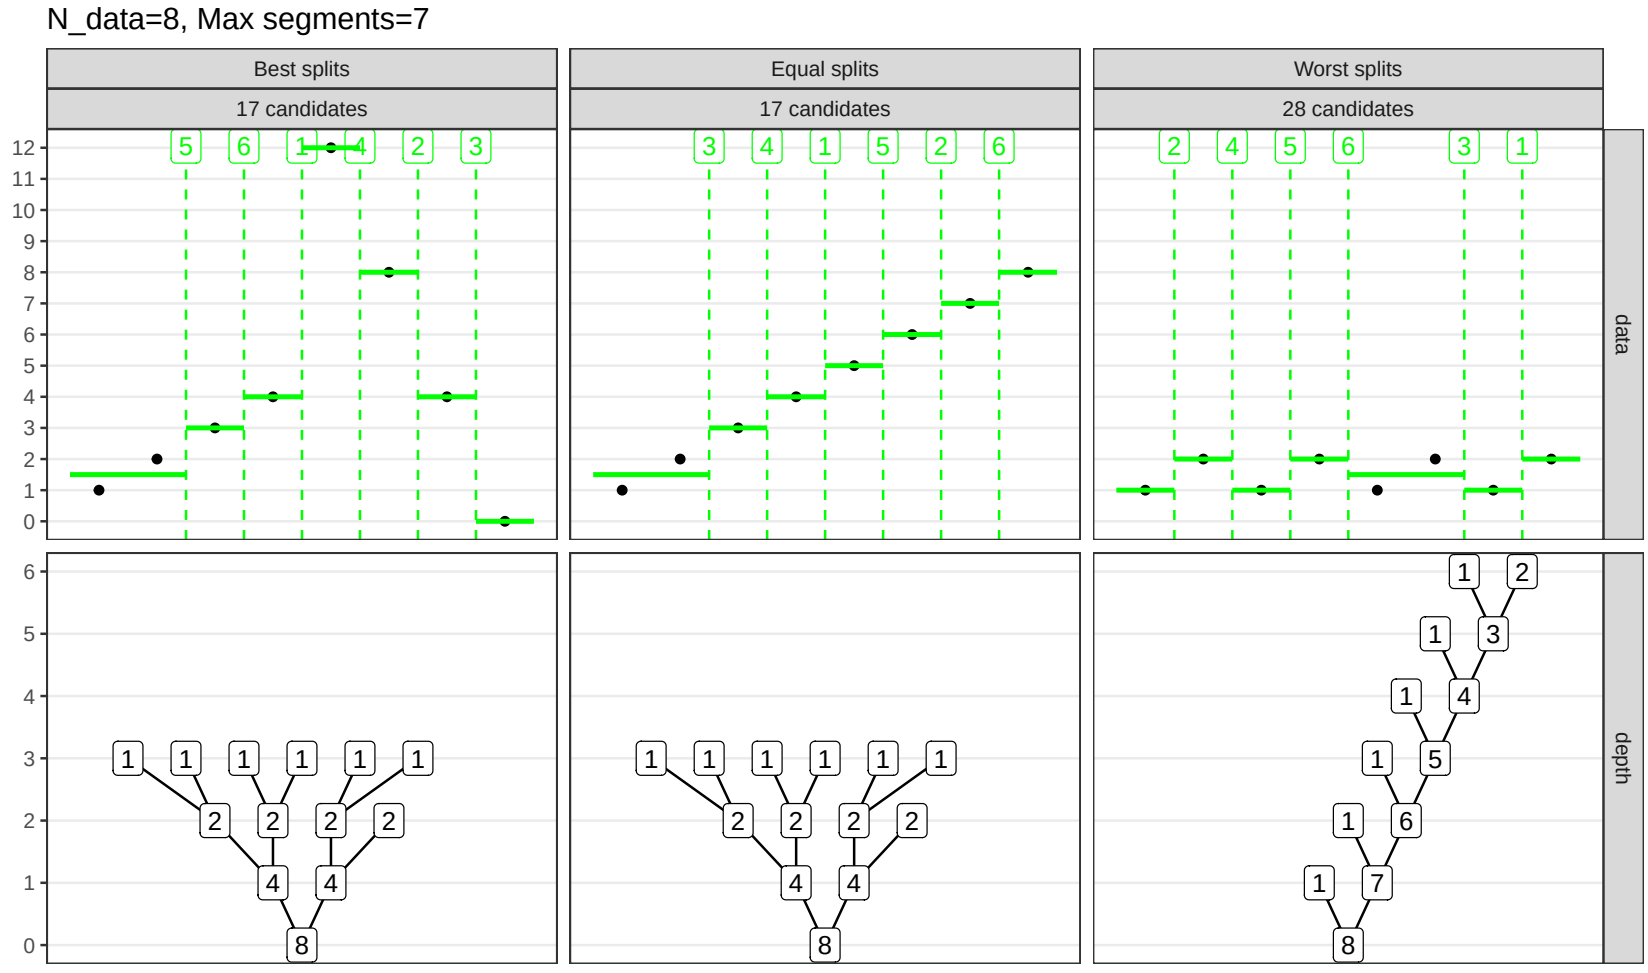

Synthetic data time complexity, vary max segments

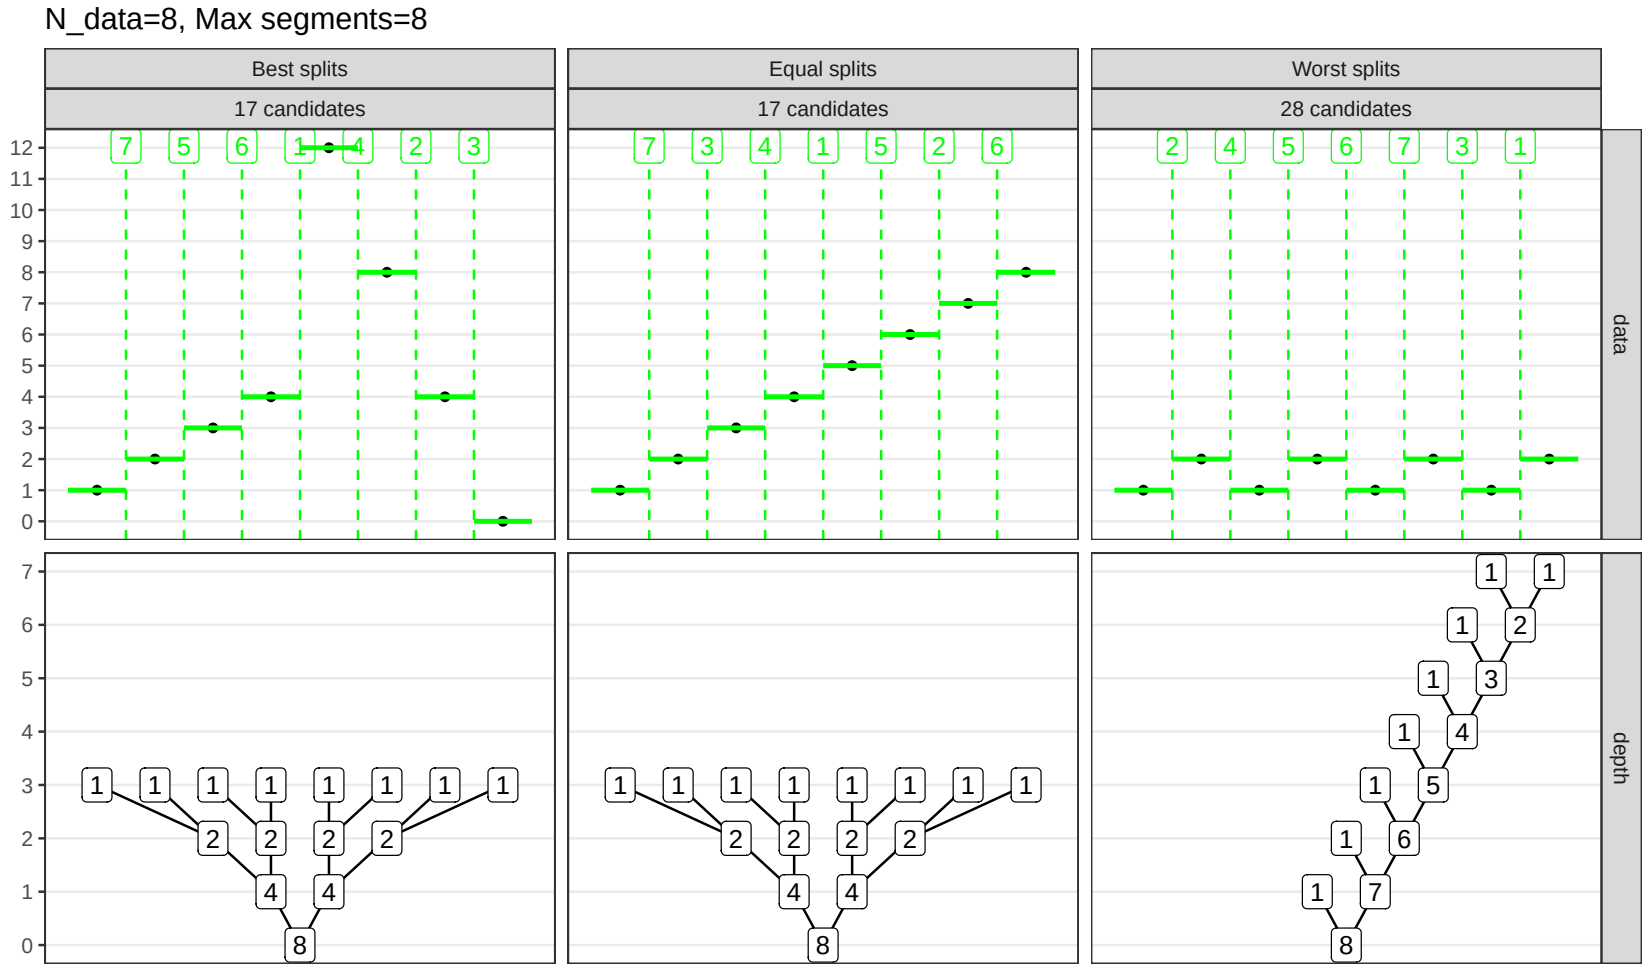

Optimal binary trees

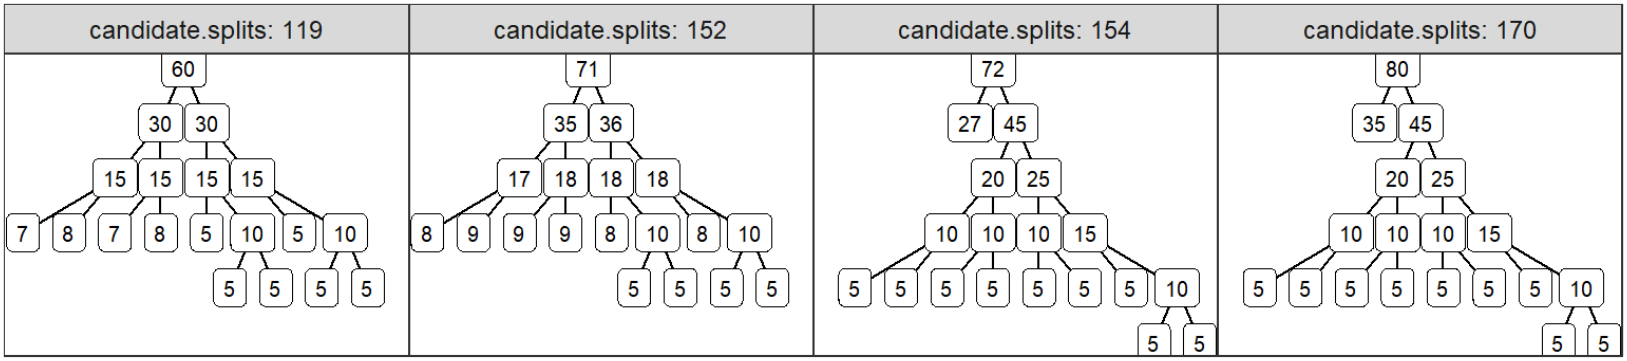

- $n \in \{60, 71, 72, 80\}$ data, min segment length $m = 5$, and number of splits/iterations $K = 9$.
- Smaller $n \in \{60, 71\}$ values result in a balanced first split,
- Larger $n \in \{72, 80\}$ values result in an unbalanced first split (one small child with no splits, one large child with all remaining splits).
- Hocking arXiv:2410.08654 Finite Sample Complexity Analysis of Binary Segmentation propose algorithm to compute optimal binary tree.
- Is binary segmentation asympototically optimal speed in real data?

Empirical number of candidates (Poisson loss)

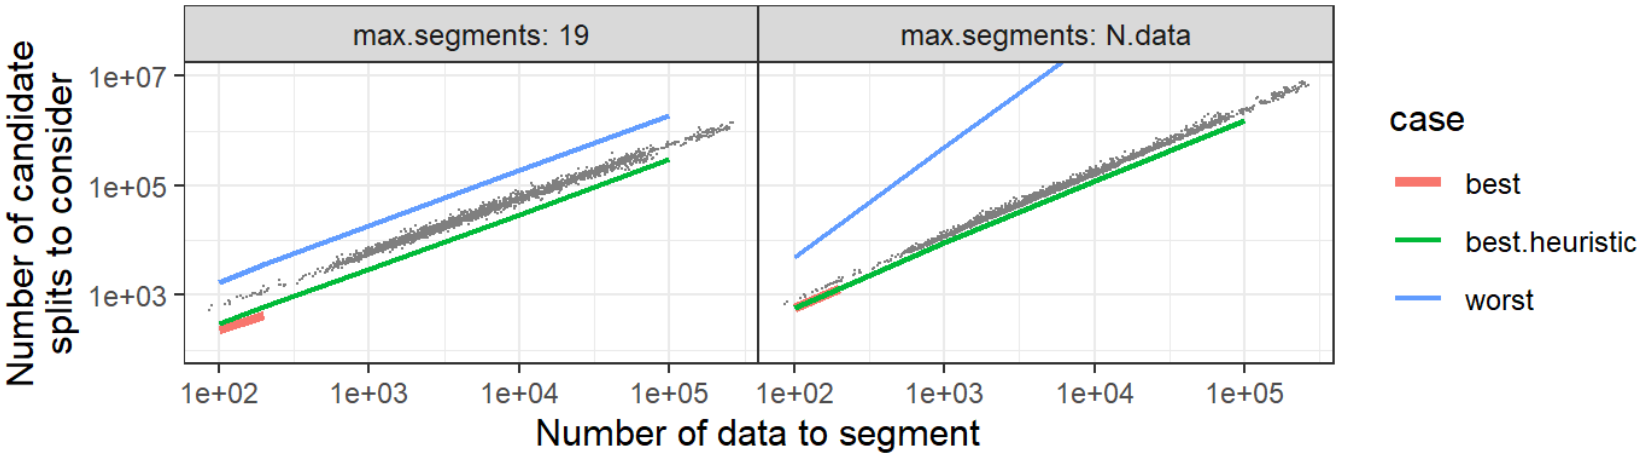

- 2752 real genomic count data sets from McGill benchmark of data size n from 87 to 263169, using binary segmentation with the Poisson loss.
- Number of candidate splits to consider in real data (grey dots) achieves the asymptotic best case, $O(n \log n)$.

Empirical number of candidates (Square loss)

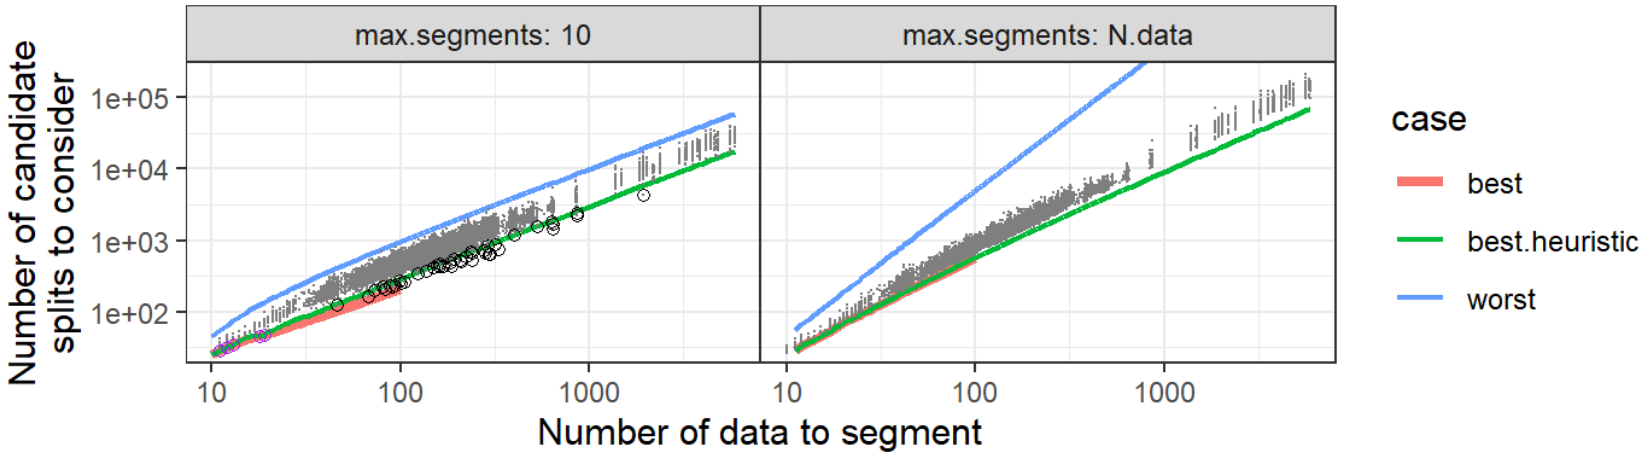

- 13721 real genomic data sets from neuroblastoma benchmark of data size n from 11 to 5937, using binary segmentation with the square loss.
- Number of candidate splits achieves the asymptotic best case, $O(n \log n)$, and in 45 instances (black circles) requires fewer candidate splits than predicted by the best case heuristic.
- For max segments $= 10$, we used dynamic programming to compute the best case number of splits for all data sizes between 11 and 100 (orange line), and the heuristic (green line) was exact for only 5 data sizes $n \in \{11, 12, 13, 18, 19\}$ (purple circles).

Introduction

Analysis of time complexity

Comparing empirical timings between implementations

Analysis of insert time and storage (STL list)

To store previously computed best loss/split for each segment, use C++ Standard template library list as baseline.

- If list has p items then insert takes $O(1)$ time.
- Find best segment to split: $O(p)$ time.
- Total $O(K^2)$ time for K iterations (equal splits), larger than $O(n \log K)$ time for computing candidate loss values.

splits K	items in list (equal splits)	items in list (unequal splits)
1	1	1
2	2	1
3	3	1
⋮	⋮	⋮
32	32	1
33	31	1
34	30	1
⋮	⋮	⋮
62	2	1
63	1	1

Analysis of insert time and storage (STL multiset)

To achieve asymptotic best case timings, use a C++ Standard template library multiset, keyed by loss decrease. If multiset has p items then insert takes $O(\log p)$ time. Below: inserts column shows p for each insert, and size column shows p after inserts. Total $-\log(K-1) + \sum_{p=1}^{K-1} 2 \log p \in O(K \log K)$ time over all inserts, smaller than $O(n \log K)$ time for split computation.

	equal splits		unequal splits	
iteration K	inserts	size	inserts	size
1	0	1	0	1
2	0,1	2	0	1
3	1,2	3	0	1
⋮	⋮	⋮	⋮	⋮
32	30,31	32	0	1
33		31	0	1
34		30	0	1
⋮	⋮	⋮	⋮	⋮
62		2	0	1
63		1	0	1

Timings (black) with asymptotic references (violet)

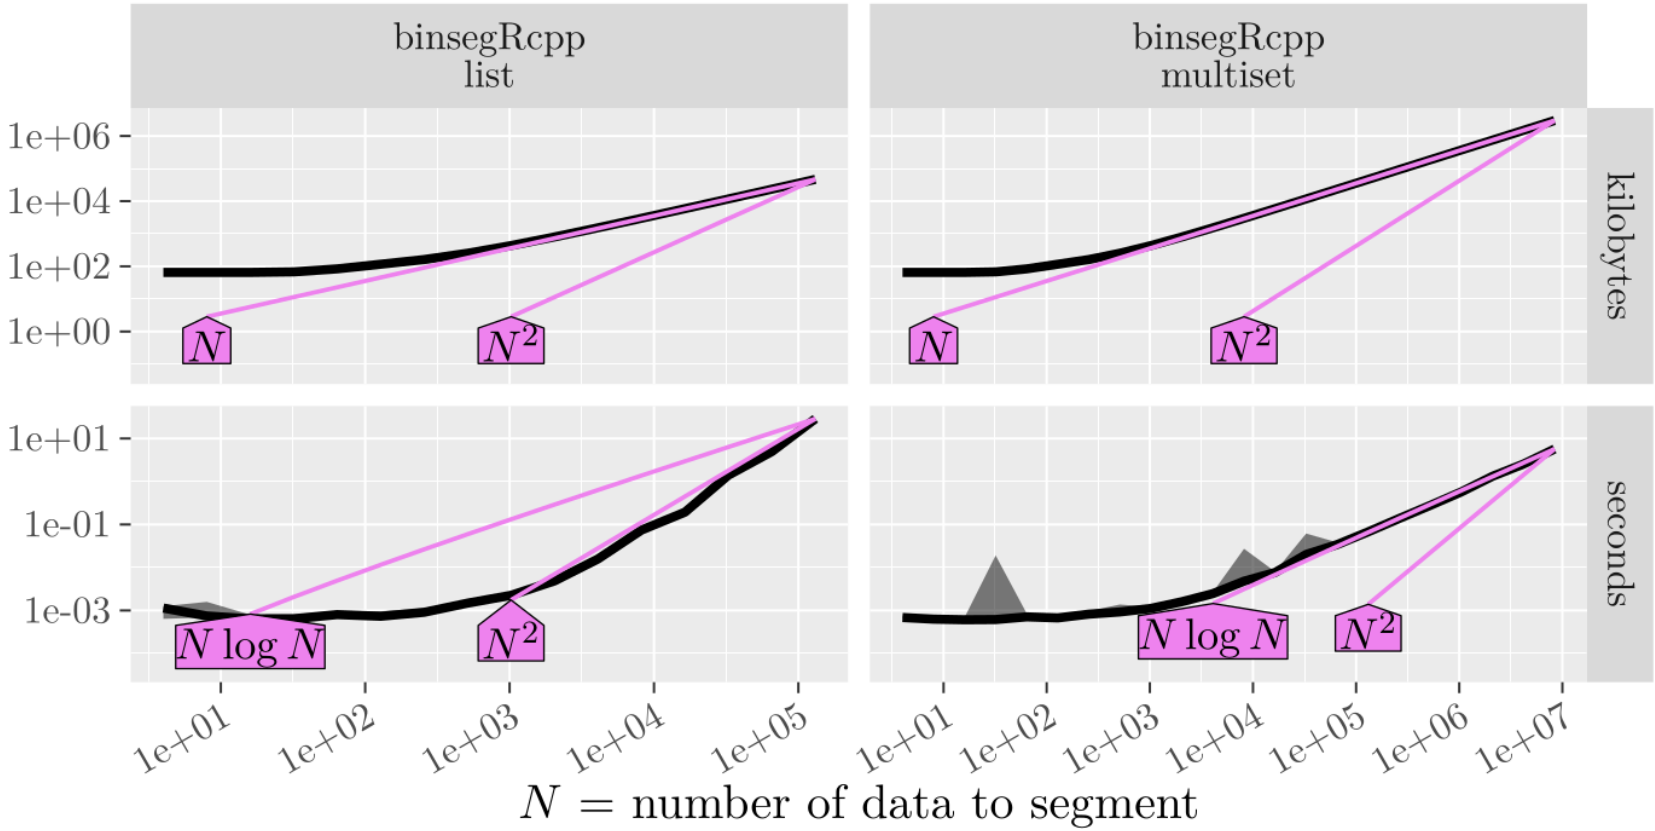

- Number of segments grows with data size, $K = O(N)$.
- Best case data, square loss.
- Linear space, $O(N)$.
- binsegRcpp multiset is $O(N \log N)$ time, asymptotically faster than binsegRcpp list, $O(N^2)$ time.

Timings (black) with asymptotic references (violet)

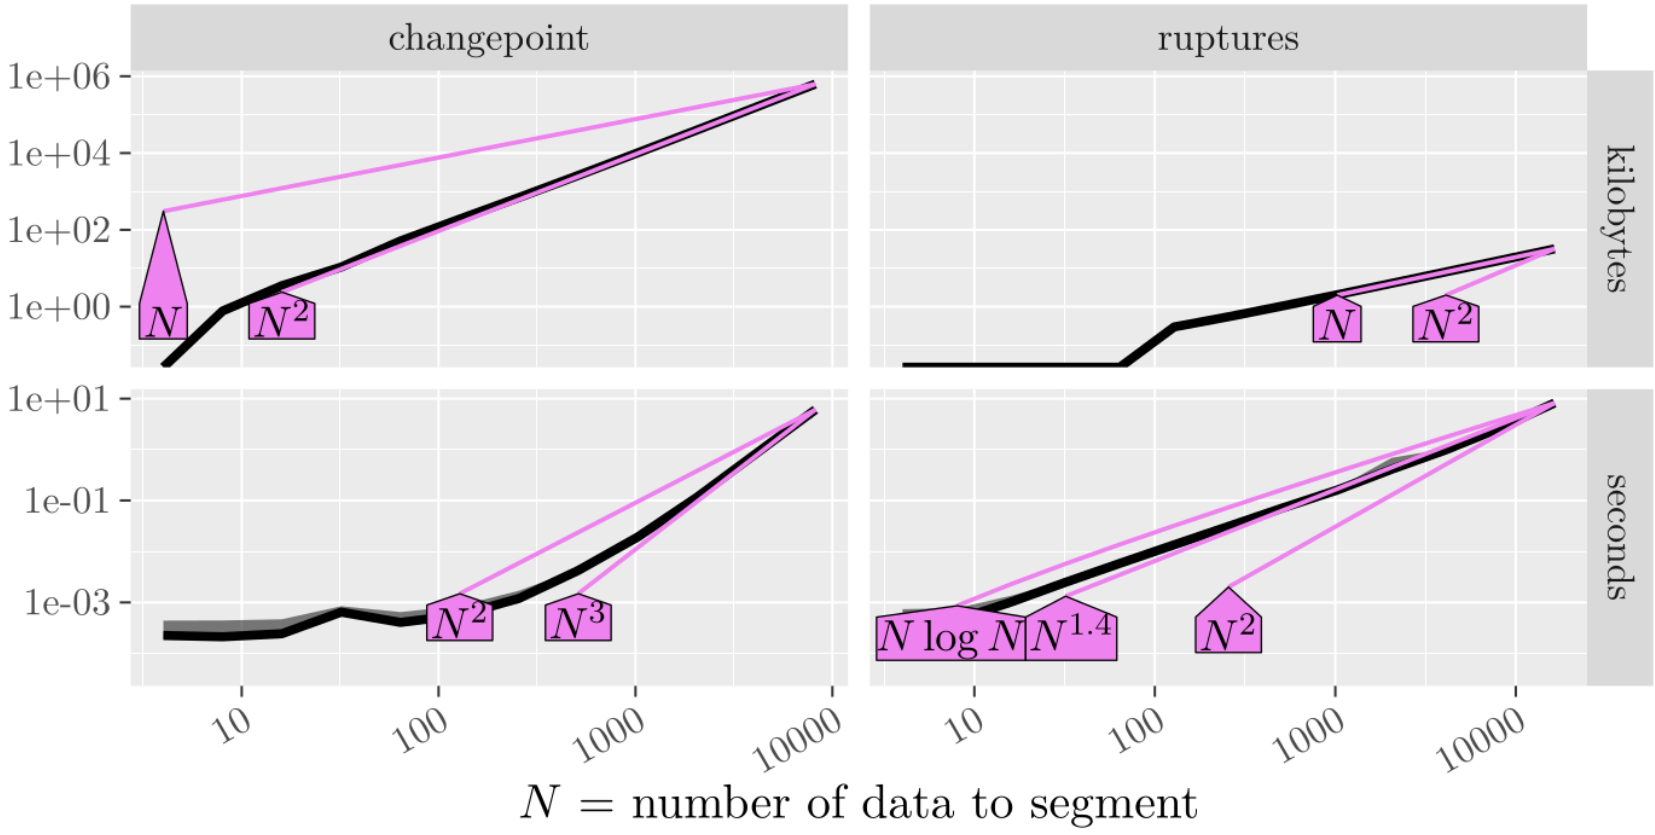

- Number of segments grows with data size, $K = O(N)$.
- Best case data, square loss.
- changepoint is $O(N^2)$ space, $O(N^3)$ time.
- ruptures is $O(N)$ space, $O(N^{1.4})$ time (unclear why)
- Not asymptotic optimal complexity.
- Sub-optimal implementation.

Timing and throughput comparisons

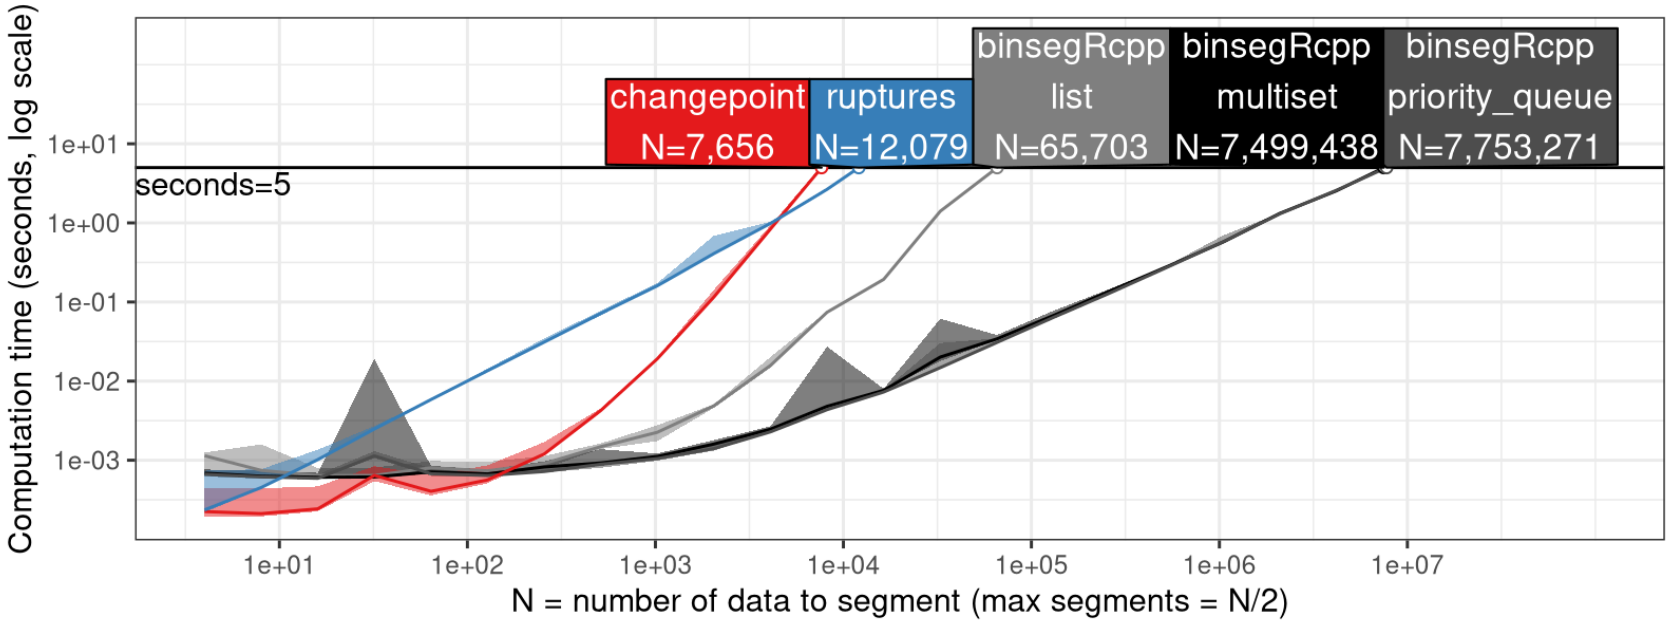

- Number of segments grows with data size, $K = O(N)$.
- Best case data, square loss.
- in 5 seconds, binsegRcpp using multiset or priority queue has 100x larger throughput than list.
- about 500x larger throughput than ruptures.
- about 1000x larger throughput than changepoint.

Comparing implementations of binary segmentation

package	binsegRcpp	changepoint	wbs	fpop	ruptures	blockcpd
function	binseg	cpt.mean	sbs	multiBinSeg	Binseg	fit_blockcpd
version	2022.3.29	2.2.3	1.4	2019.8.26	1.1.9	1.0.0
weights	yes	no	no	no	no	no
min len	yes	yes	no	no	yes	yes
max segs	yes	yes	no	yes	yes	yes
dim	one	one	one	multi	multi	multi
correct	yes	no	yes	yes	yes	yes
losses	5	6	1	1	10	5
language	C++	C	C	C++	Python	C++
storage	heap	arrays	recursion	heap	LRU cache	heap
space	$O(S)$	$O(S^2)$	$O(S)$	$O(S)$	$O(S)$	$O(S)$
cumsum	yes	yes	yes	yes	no	yes
best	$O(N \log N)$	$O(N^3)$	$O(N \log N)$	$O(N \log N)$	$O(N^{1.4})$	$O(N \log N)$
worst	$O(N^2)$	$O(N^3)$	$O(N^2)$	$O(N^2)$	$O(N^2)$	$O(N^2)$
CV	yes	no	no	no	no	no
params	all	one	no	no	no	one

Hocking 2025 Comparing binsegRcpp to other implementations of binary segmentation, in progress.

Timings using square loss

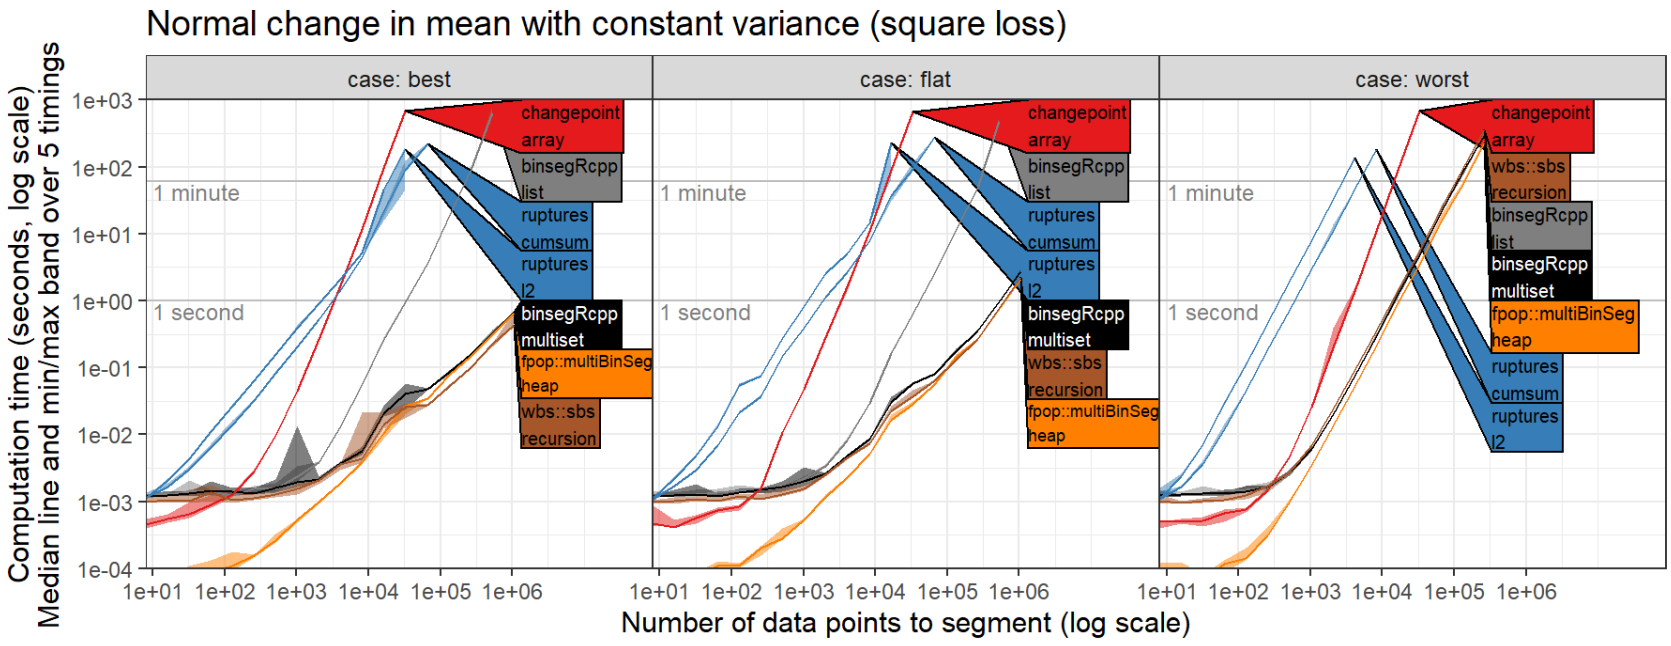

- binsegRcpp map, `wbs::sbs`, `fpop::multiBinSeg` optimal: $O(n \log n)$ time in best case.
- ruptures is between $O(n \log n)$ and $O(n^2)$ time, but with large constant factors.
- binsegRcpp list has larger slope: $O(n^2)$ time.
- changepoint has even larger slope: $O(n^3)$ time.

Timings using L1 loss

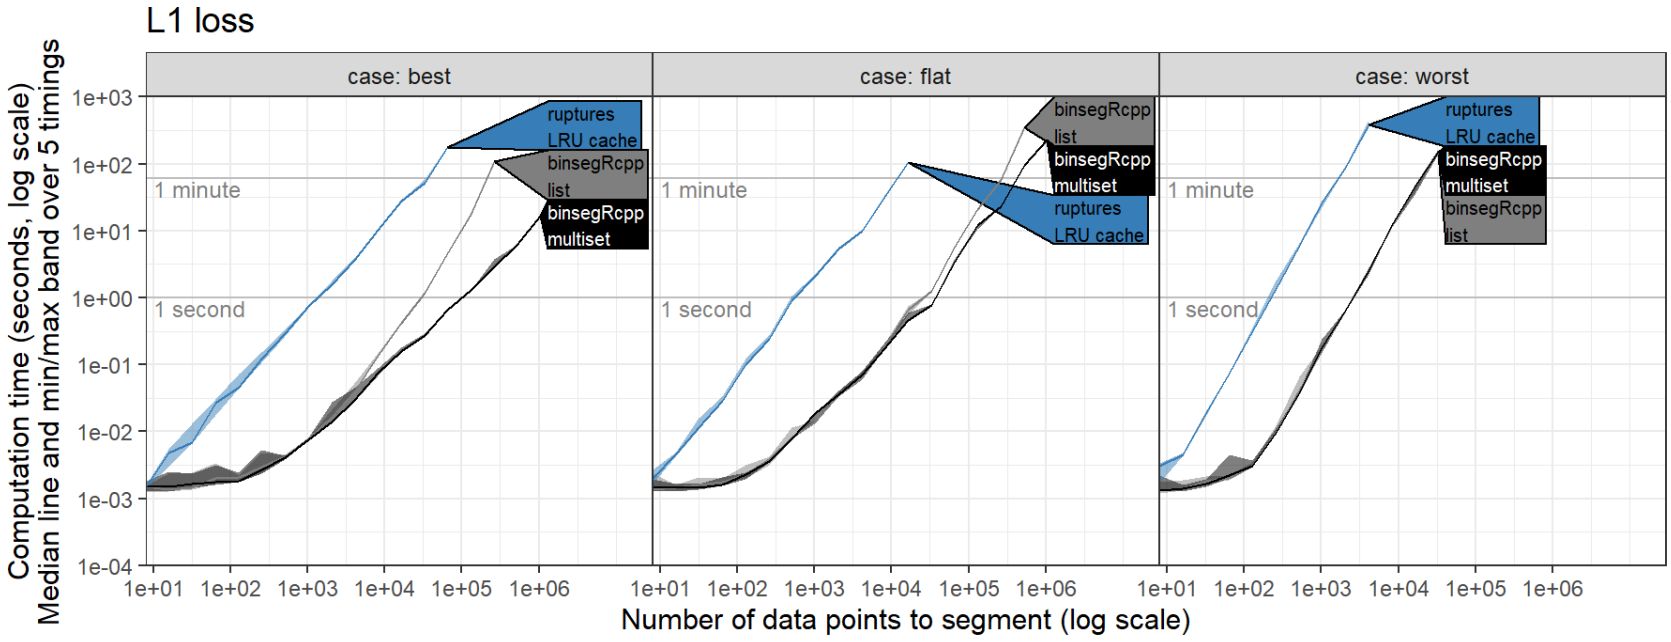

- Number of segments grows with data size, $K = O(n)$.
- Best case: binsegRcpp multiset is $O(n \log n)$, asymptotically faster than binsegRcpp list, $O(n^2)$.
- ruptures is significantly slower (100x).
- Worst case: all algorithms are quadratic, because that is the number of candidate change-points that must be considered.

Timings using Gaussian change in mean and variance model

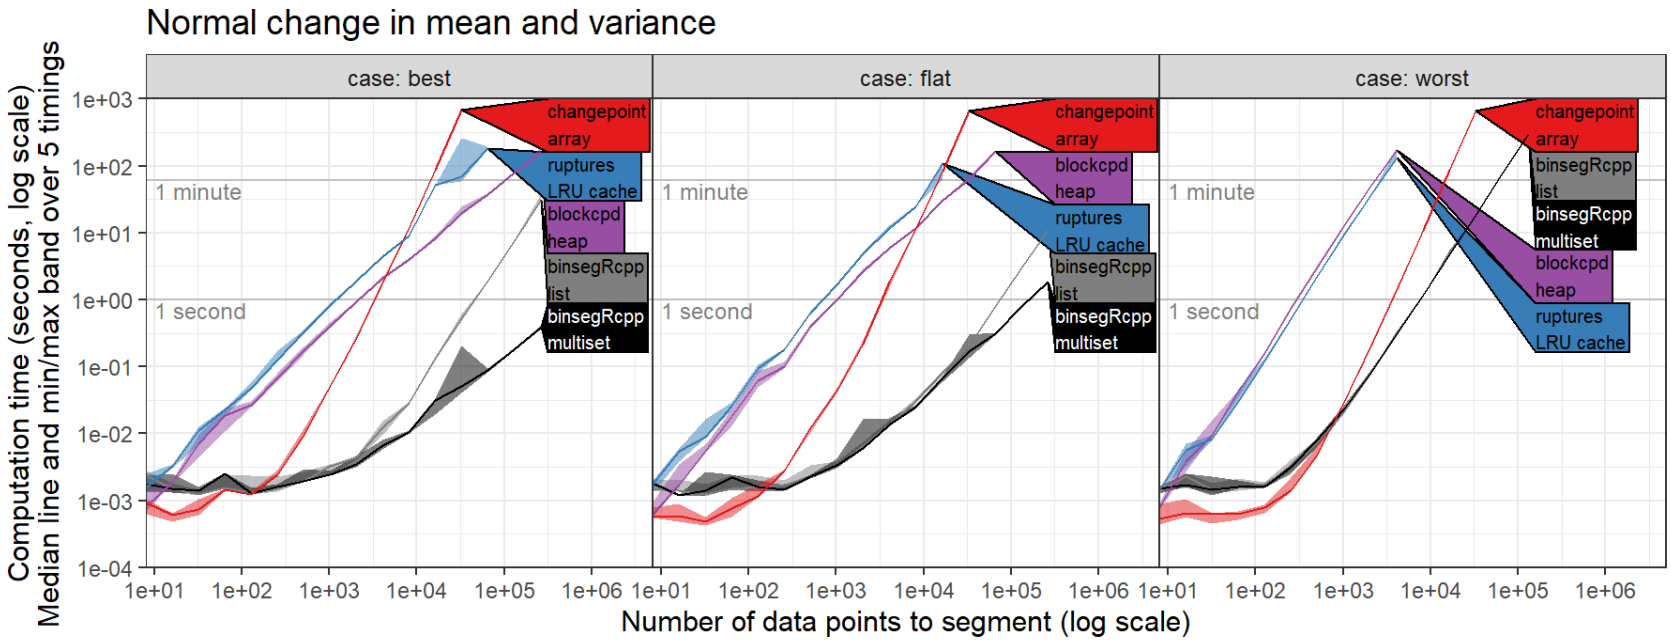

Timings using Poisson loss

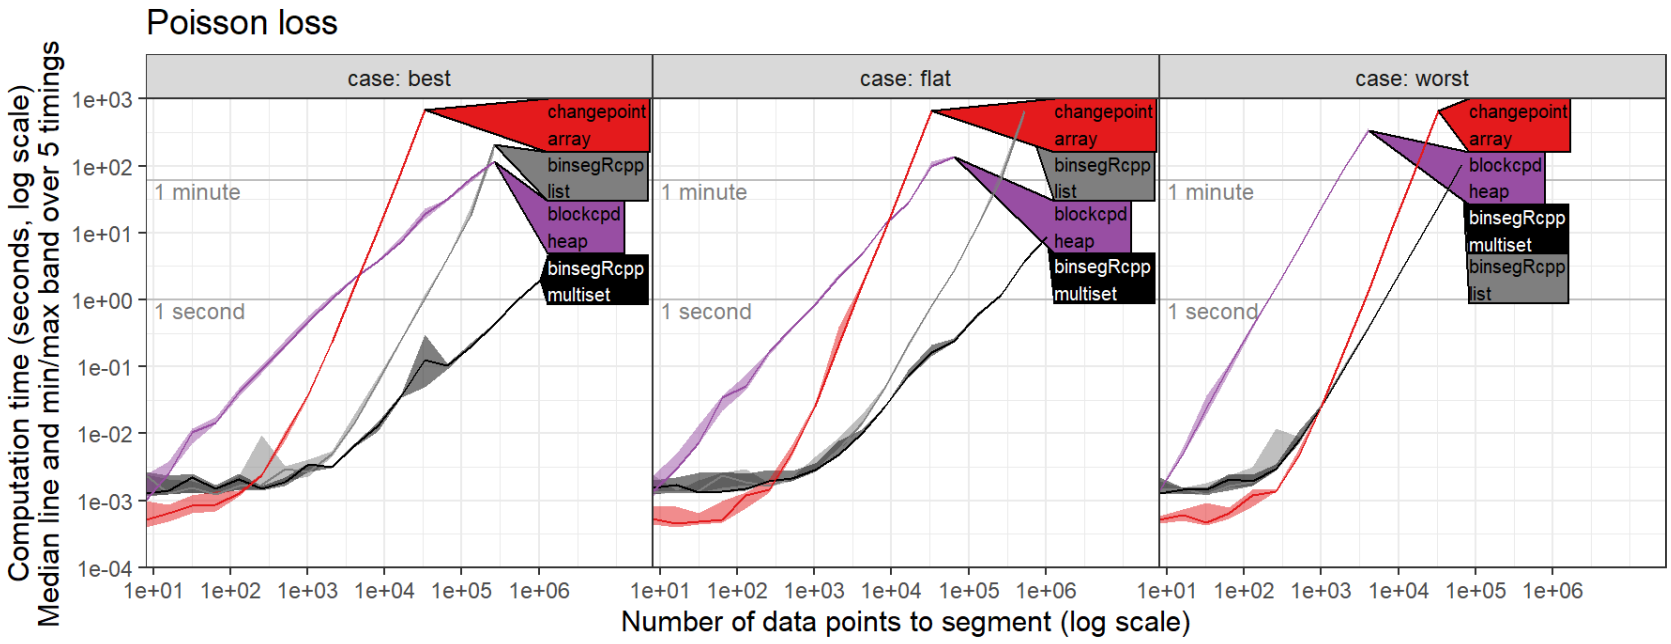

Comparing loss functions, binsegRcpp multiset

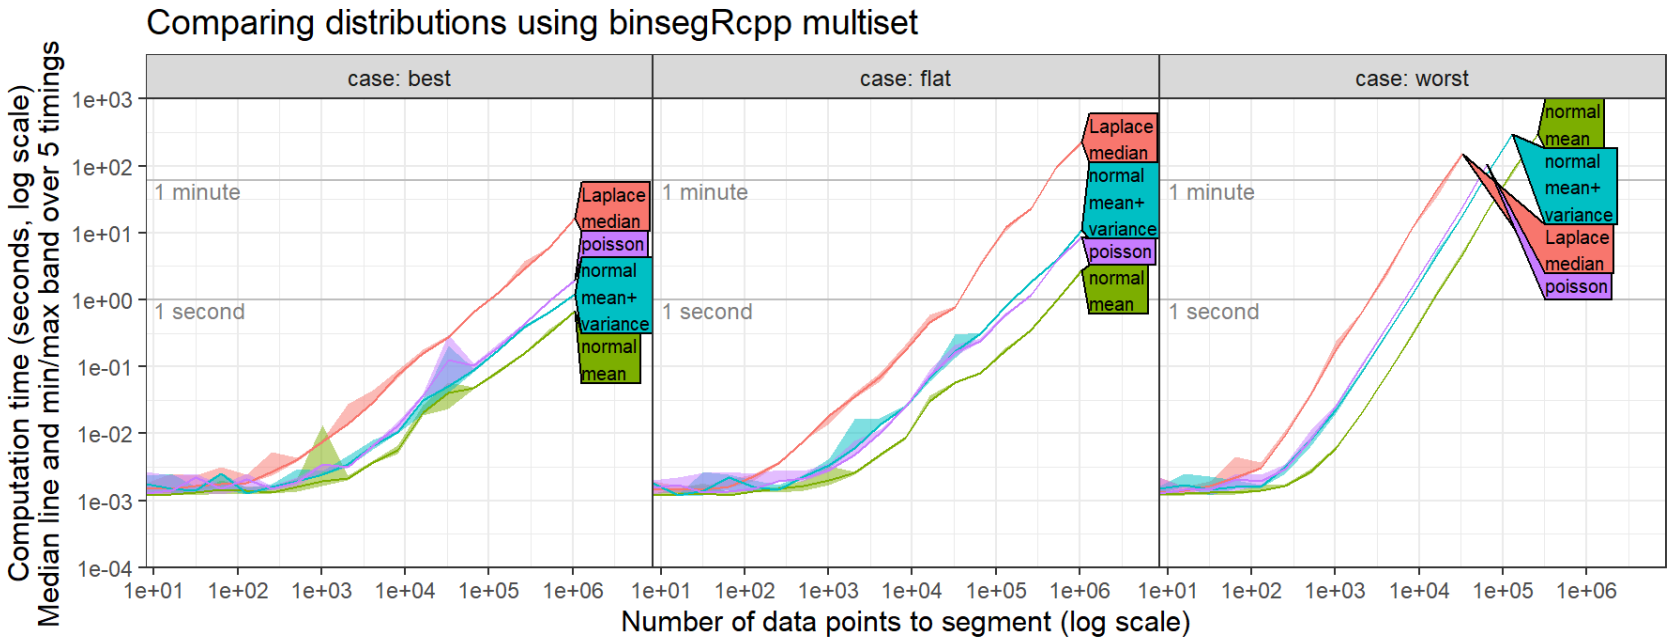

- Laplace (log-linear) is somewhat slower than others (linear).

Accuracy of binsegRcpp and changepoint

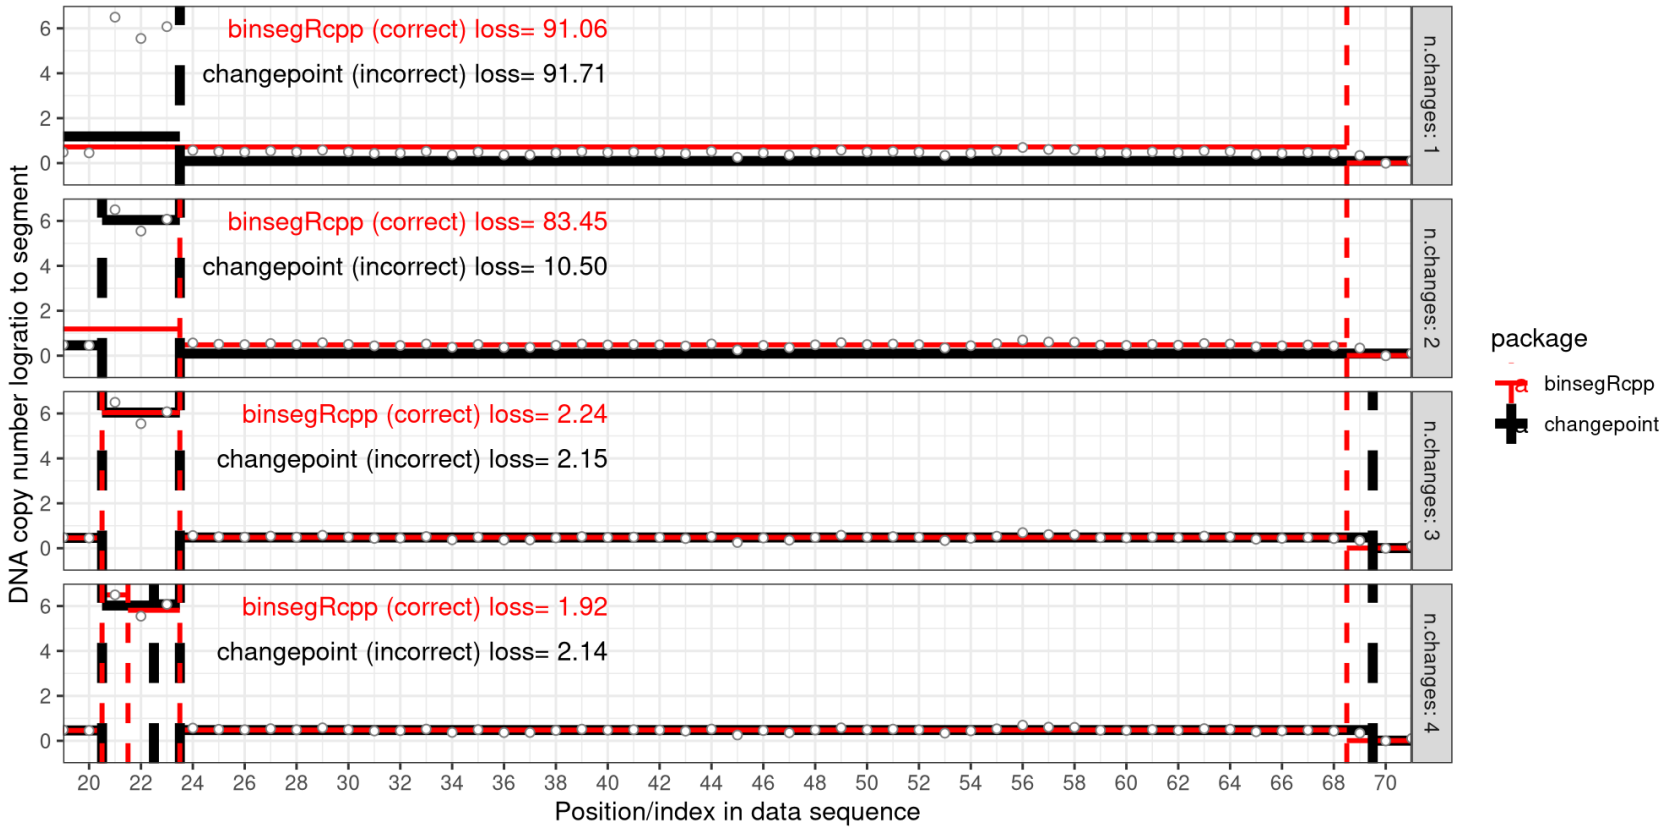

- Real data set for which changepoint does not compute the correct model.
- Loss should be minimized at each iteration.

Conclusions

- Binary segmentation is an important algorithm for change-point detection in sequential data.
- Optimal binary trees used to determine best case number of candidates that the algorithm needs to consider.
- Empirical number of candidates in real data: same as best case, asymptotically.
- Proposed binsegRcpp R package provides C++ code which is asymptotically optimal speed, and correct.
- `install.packages("binsegRcpp")` in R.
- Future work: faster method for computing optimal binary trees?

Comparison with previous algorithms from clustering

- Binary segmentation has segment/cluster-specific mean parameter, as in K-means and Gaussian mixture models. These algorithms attempt optimization of an error function which measures how well the means fit the data (but are not guaranteed to compute the globally optimal/best model).
- Binary segmentation is deterministic (different from K-means/GMM which requires random initialization). It performs a sequence of greedy minimizations (as in hierarchical clustering).
- Binary segmentation defines a sequence of split operations (from 1 segment to N segments), whereas agglomerative hierarchical clustering defines a sequence of join operations (from N clusters to 1 cluster). Data with common segment mean must be adjacent in time/space; hierarchical clustering joins may happen between any pair of data points (no space/time dimension).

Possible exam questions

- Explain in detail one similarity and one difference between binary segmentation and k-means. (gaussian mixture models, hierarchical clustering)
- For a sequence of $n = 10$ data, we need to compute the loss for each of the 9 possible splits in the first iteration of binary segmentation. What is the number of splits for which we must compute the loss in the second/third steps? (best and worst case)